



Red Hat OpenStack Platform 16.1

Instances and Images Guide

Managing Instances and Images

Red Hat OpenStack Platform 16.1 Instances and Images Guide

Managing Instances and Images

OpenStack Team
rhos-docs@redhat.com

Legal Notice

Copyright © 2020 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

The Instances and Images guide provides procedures for the management of instances, images of a Red Hat OpenStack Platform environment.

Table of Contents

PREFACE	6
CHAPTER 1. IMAGE SERVICE	7
1.1. UNDERSTANDING THE IMAGE SERVICE	7
1.1.1. Supported Image service (glance) back ends	7
1.1.2. Image signing and verification	8
1.1.3. Image conversion	9
1.1.4. Image introspection	9
1.1.5. Interoperable image import	10
1.1.6. Improving scalability with Image service caching	10
1.1.7. Image pre-caching	11
1.1.7.1. Configuring the default interval for periodic image pre-caching	11
1.1.7.2. Using a periodic job to pre-cache an image	11
1.2. MANAGE IMAGES	14
1.2.1. Creating an Image	14
1.2.1.1. Use a KVM Guest Image With Red Hat OpenStack Platform	14
1.2.1.2. Create Custom Red Hat Enterprise Linux or Windows Images	15
1.2.1.2.1. Create a Red Hat Enterprise Linux 7 Image	15
1.2.1.2.2. Create a Red Hat Enterprise Linux 6 Image	20
1.2.1.2.3. Create a Windows Image	24
1.2.2. Upload an image	26
1.2.3. Update an image	27
1.2.4. Import an image	28
1.2.4.1. Import from a remote URI	28
1.2.4.2. Import from a local volume	28
1.2.5. Delete an image	29
1.2.6. Hide or unhide an image	29
1.2.7. Show hidden images	29
1.2.8. Enabling image conversion	29
1.2.9. Converting an image to RAW format	30
1.2.9.1. Configuring Image Service to accept RAW and ISO only	30
1.2.10. Storing an image in RAW format	30
CHAPTER 2. IMAGE SERVICE WITH MULTIPLE STORES	32
2.1. REQUIREMENTS OF STORAGE EDGE ARCHITECTURE	32
2.2. IMPORT AN IMAGE TO MULTIPLE STORES	32
2.2.1. Manage image import failures	33
2.2.2. Importing image data to multiple stores	33
2.2.3. Importing image data to multiple stores without failure	33
2.2.4. Importing image data to a single store	34
2.2.5. Checking the progress of the image import operation	35
2.3. COPY AN EXISTING IMAGE TO MULTIPLE STORES	36
2.3.1. Copying an image to all stores	36
2.3.2. Copying an image to specific stores	37
2.4. DELETING AN IMAGE FROM A SPECIFIC STORE	37
2.5. UNDERSTANDING THE LOCATIONS OF IMAGES	38
CHAPTER 3. CONFIGURING THE COMPUTE (NOVA) SERVICE	40
3.1. CONFIGURING MEMORY FOR OVERALLOCATION	41
3.2. CALCULATING RESERVED HOST MEMORY ON COMPUTE NODES	41
3.3. CALCULATING SWAP SIZE	42

CHAPTER 4. SCALING DEPLOYMENTS WITH COMPUTE CELLS	43
4.1. CELL COMPONENTS	43
4.2. CELL DEPLOYMENTS ARCHITECTURE	44
4.3. CONSIDERATIONS FOR MULTI-CELL DEPLOYMENTS	45
4.4. DEPLOYING A MULTI-CELL OVERCLOUD	47
4.5. CREATING AND PROVISIONING A CELL	52
4.6. ADDING COMPUTE NODES TO A CELL	54
4.7. CONFIGURING AN AVAILABILITY ZONE	54
4.8. DELETING A COMPUTE NODE FROM A CELL	55
4.9. DELETING A CELL	55
CHAPTER 5. CREATING AND MANAGING HOST AGGREGATES	57
5.1. ENABLING SCHEDULING ON HOST AGGREGATES	57
5.2. CREATING A HOST AGGREGATE	58
5.3. CREATING AN AVAILABILITY ZONE	59
5.4. DELETING A HOST AGGREGATE	60
5.5. CREATING A TENANT-ISOLATED HOST AGGREGATE	61
CHAPTER 6. CONFIGURE OPENSTACK COMPUTE STORAGE	63
6.1. ARCHITECTURE OVERVIEW	63
6.2. CONFIGURATION	64
CHAPTER 7. VIRTUAL MACHINE INSTANCES	67
7.1. MANAGE INSTANCES	67
7.1.1. Add Components	67
7.1.2. Launch an Instance	67
7.1.2.1. Launch Instance Options	68
7.1.3. Update an Instance (Actions menu)	69
7.1.4. Resize an Instance	71
7.1.5. Connect to an Instance	71
7.1.5.1. Access an Instance Console using the Dashboard	72
7.1.5.2. Directly Connect to a VNC Console	72
7.1.6. View Instance Usage	73
7.1.7. Delete an Instance	73
7.1.8. Manage Multiple Instances at Once	73
7.2. MANAGE INSTANCE SECURITY	74
7.2.1. Manage Key Pairs	74
7.2.1.1. Create a Key Pair	74
7.2.1.2. Import a Key Pair	74
7.2.1.3. Delete a Key Pair	74
7.2.2. Create a Security Group	75
7.2.3. Create, Assign, and Release Floating IP Addresses	75
7.2.3.1. Allocate a Floating IP to the Project	75
7.2.3.2. Assign a Floating IP	75
7.2.3.3. Release a Floating IP	76
7.2.4. Log in to an Instance	76
7.2.5. Inject an admin Password Into an Instance	77
7.3. MANAGE FLAVORS	78
7.3.1. Update Configuration Permissions	79
7.3.2. Create a Flavor	79
7.3.3. Update General Attributes	80
7.3.4. Update Flavor Metadata	80
7.3.4.1. View Metadata	80
7.3.4.2. Add Metadata	80

7.4. SCHEDULE HOSTS	85
7.4.1. Configure Scheduling Filters	86
7.4.2. Configure Scheduling Weights	89
7.4.3. Configure Placement Service Traits	95
7.4.3.1. libvirt virtualization driver capabilities as placement service traits	96
7.4.3.2. Using placement service traits to specify resource provider requirements	96
7.4.4. Configuring a guaranteed minimum bandwidth QoS	97
7.4.4.1. Removing a guaranteed minimum bandwidth QoS from an instance	99
7.4.5. Reserve NUMA Nodes with PCI Devices	99
7.4.6. Configure Emulator Threads to run on Dedicated Physical CPU	99
7.5. MANAGE INSTANCE SNAPSHOTS	100
7.5.1. Create an Instance Snapshot	101
7.5.2. Manage a Snapshot	101
7.5.3. Rebuild an Instance to a State in a Snapshot	102
7.5.4. Consistent Snapshots	102
7.6. USE RESCUE MODE FOR INSTANCES	102
7.6.1. Preparing an Image for a Rescue Mode Instance	102
7.6.1.1. Rescue Image if Using ext4 Filesystem	103
7.6.2. Adding the Rescue Image to the OpenStack Image Service	103
7.6.3. Launching an Instance in Rescue Mode	104
7.6.4. Unrescuing an Instance	104
7.7. SET A CONFIGURATION DRIVE FOR INSTANCES	104
7.7.1. Configuration Drive Options	105
7.7.2. Use a Configuration Drive	105
CHAPTER 8. CONFIGURING COMPUTE NODES FOR PERFORMANCE	106
8.1. CONFIGURING CPU PINNING ON THE COMPUTE NODE	106
8.1.1. Upgrading CPU pinning configuration	108
8.1.2. Launching an instance with CPU pinning	108
8.1.3. Launching a floating instance	109
8.2. CONFIGURING HUGE PAGES ON THE COMPUTE NODE	109
8.2.1. Allocating huge pages to instances	111
8.3. CONFIGURING COMPUTE NODES TO USE FILE-BACKED MEMORY FOR INSTANCES	112
8.3.1. Changing the memory backing directory host disk	113
CHAPTER 9. CONFIGURING SEV-CAPABLE COMPUTE NODES TO PROVIDE MEMORY ENCRYPTION FOR INSTANCES	114
9.1. SECURE ENCRYPTED VIRTUALIZATION (SEV)	114
9.2. CONFIGURING A SEV-CAPABLE COMPUTE NODE	115
9.3. CREATING THE IMAGE AND FLAVOR FOR MEMORY ENCRYPTION	116
9.3.1. Creating a SEV-enabled image for instances	116
9.3.2. Creating a SEV-enabled flavor for instances	117
CHAPTER 10. MIGRATING VIRTUAL MACHINE INSTANCES BETWEEN COMPUTE NODES	118
10.1. MIGRATION TYPES	118
10.2. MIGRATION CONSTRAINTS	120
10.3. PREPARING TO MIGRATE	121
10.4. COLD MIGRATING AN INSTANCE	122
10.5. LIVE MIGRATING AN INSTANCE	123
10.6. CHECKING MIGRATION STATUS	124
10.7. EVACUATING AN INSTANCE	125
10.7.1. Evacuating one instance	125
10.7.2. Evacuating all instances on a host	126
10.7.3. Configuring shared storage	127

10.8. TROUBLESHOOTING MIGRATION	128
10.8.1. Errors during migration	128
10.8.2. Never-ending live migration	129
10.8.3. Instance performance degrades after migration	130
CHAPTER 11. CONFIGURING VIRTUAL GPUS FOR INSTANCES	131
11.1. SUPPORTED CONFIGURATIONS AND LIMITATIONS	131
11.2. CONFIGURING VGPU ON THE COMPUTE NODES	131
11.2.1. Building a custom GPU overcloud image	132
11.2.2. Designating Compute nodes for vGPU	133
11.2.3. Configuring the Compute node for vGPU and deploying the overcloud	136
11.3. CREATING THE VGPU IMAGE AND FLAVOR	137
11.3.1. Creating a custom GPU instance image	137
11.3.2. Creating a vGPU flavor for instances	138
11.3.3. Launching a vGPU instance	139
11.4. ENABLING PCI PASSTHROUGH FOR A GPU DEVICE	139
CHAPTER 12. CONFIGURING REAL-TIME COMPUTE	143
12.1. PREPARING YOUR COMPUTE NODES FOR REAL-TIME	143
12.2. DEPLOYING THE REAL-TIME COMPUTE ROLE	146
12.3. SAMPLE DEPLOYMENT AND TESTING SCENARIO	148
12.4. LAUNCHING AND TUNING REAL-TIME INSTANCES	149
APPENDIX A. IMAGE CONFIGURATION PARAMETERS	152
APPENDIX B. ENABLING THE LAUNCH INSTANCE WIZARD	163

PREFACE

Red Hat OpenStack Platform (RHOSP) provides the foundation to build a private or public Infrastructure-as-a-Service (IaaS) cloud on top of Red Hat Enterprise Linux. It offers a massively scalable, fault-tolerant platform for the development of cloud-enabled workloads.

This guide discusses procedures for creating and managing images, and instances. It also mentions the procedure for configuring the storage for instances for RHOSP.

You can manage the cloud by using either the RHOSP Dashboard or the command-line clients. You can use either method to perform most procedures but some of the more advanced procedures can only be executed on the command line. This guide provides procedures for the Dashboard where possible.



NOTE

For the complete suite of documentation for Red Hat OpenStack Platform, see [Red Hat OpenStack Platform Documentation Suite](#).

CHAPTER 1. IMAGE SERVICE

Manage images and storage in Red Hat OpenStack Platform (RHOSP).

A virtual machine image is a file that contains a virtual disk with a bootable operating system installed. Virtual machine images are supported in different formats. The following formats are available in RHOSP:

- **RAW** - Unstructured disk image format.
- **QCOW2** - Disk format supported by QEMU emulator. This format includes QCOW2v3 (sometimes referred to as QCOW3), which requires QEMU 1.1 or higher.
- **ISO** - Sector-by-sector copy of the data on a disk, stored in a binary file.
- **AKI** - Indicates an Amazon Kernel Image.
- **AMI** - Indicates an Amazon Machine Image.
- **ARI** - Indicates an Amazon RAMDisk Image.
- **VDI** - Disk format supported by VirtualBox virtual machine monitor and the QEMU emulator.
- **VHD** - Common disk format used by virtual machine monitors from VMware, VirtualBox, and others.
- **VMDK** - Disk format supported by many common virtual machine monitors.

Although **ISO** is not normally considered a virtual machine image format, because ISOs contain bootable filesystems with an installed operating system, you use them in the same way as other virtual machine image files.

To download the official Red Hat Enterprise Linux cloud images, your account must have a valid Red Hat Enterprise Linux subscription:

- [Red Hat Enterprise Linux 8 KVM Guest Image](#)
- [Red Hat Enterprise Linux 7 KVM Guest Image](#)
- [Red Hat Enterprise Linux 6 KVM Guest Image](#)

If you are not logged in to the Customer Portal, a prompt opens where you must enter your Red Hat account credentials.

1.1. UNDERSTANDING THE IMAGE SERVICE

Red Hat OpenStack Platform (RHOSP) Image service (glance) features.

1.1.1. Supported Image service (glance) back ends

The following Image service (glance) back end scenarios are supported:

- RBD is the default back end when you use Ceph. For more information, see [Configuring Ceph Storage](#) in the *Advanced OpenCloud Customization* guide.

- RBD multi-store. For more information, see [Deploying the central site](#) in the *Distributed compute node and storage deployment* guide.
- Object Storage (swift). For more information, see [Using an External Object Storage Cluster](#) in the *Advanced Opencloud Customization* guide.
- Block Storage (cinder). For more information, see [Configuring cinder back end for the Image service](#) in the *Advanced Opencloud Customization* guide.
The Image service uses the Block Storage type and back end as the default.
- NFS. For more information, see [Configuring NFS Storage](#) in the *Advanced Opencloud Customization* guide.
Although NFS is a supported Image service deployment option, more robust options are available.

NFS is not native to the Image service. When you mount an NFS share on the Image service, the Image service does not manage the operation. The Image service writes data to the file system but is unaware that the back end is an NFS share.

In this type of deployment, the Image service cannot retry a request if the share fails. This means that when a failure occurs on the back end, the store might enter read-only mode, or it might continue to write data to the local file system, in which case you risk data loss. To recover from this situation, you must ensure that the share is mounted and in sync, and then restart the Image service. For these reasons, Red Hat does not recommend NFS as an Image service back end.

However, if you do choose to use NFS as an Image service back end, some of the following best practices can help to mitigate risks:

- Use a reliable production-grade NFS back end.
- Ensure that you have a strong and reliable connection between Controller nodes and the NFS back end, L2 is recommended.
- Include monitoring and alerts for the mounted share.
- Set underlying FS permissions.
 - Ensure that the user and the group that the glance-api process runs on do not have write permissions on the mount point at the local file system. This means that the process can detect possible mount failure and put the store into read-only mode during a write attempt.
 - The write permissions must be present in the shared file system that you use as a store.

1.1.2. Image signing and verification

Image signing and verification protects image integrity and authenticity by enabling deployers to sign images and save the signatures and public key certificates as image properties.

By taking advantage of this feature, you can:

- Sign an image using your private key and upload the image, the signature, and a reference to your public key certificate (the verification metadata). The Image service then verifies that the signature is valid.

- Create an image in the Compute service, have the Compute service sign the image, and upload the image and its verification metadata. The Image service again verifies that the signature is valid.
- Request a signed image in the Compute service. The Image service provides the image and its verification metadata, allowing the Compute service to validate the image before booting it.

For information on image signing and verification, refer to the [Validate Glance Images](#) chapter of the *Manage Secrets with OpenStack Key Manager Guide*.

1.1.3. Image conversion

Image conversion converts images by calling the task API while importing an image.

As part of the import workflow, a plugin provides the image conversion. This plugin can be activated or deactivated based on the deployer configuration. Therefore, the deployer needs to specify the preferred format of images for the deployment.

Internally, the Image service receives the bits of the image in a particular format. These bits are stored in a temporary location. The plugin is then triggered to convert the image to the target format and moved to a final destination. When the task is finished, the temporary location is deleted. As a result, the format uploaded initially is not retained by the Image service.

For more information about image conversion, see [Enabling image conversion](#).



NOTE

The conversion can be triggered only when importing an image. It does not run when uploading an image. For example:

```
$ glance image-create-via-import \
  --disk-format qcow2 \
  --container-format bare \
  --name NAME \
  --visibility public \
  --import-method web-download \
  --uri http://server/image.qcow2
```

1.1.4. Image introspection

Every image format comes with a set of metadata embedded inside the image itself. For example, a stream optimized **vmdk** would contain the following parameters:

```
$ head -20 so-disk.vmdk

# Disk DescriptorFile
version=1
CID=d5a0bce5
parentCID=ffffff
createType="streamOptimized"

# Extent description
RDONLY 209714 SPARSE "generated-stream.vmdk"

# The Disk Data Base
```

```
#DDB
```

```
ddb.adapterType = "buslogic"
ddb.geometry.cylinders = "102"
ddb.geometry.heads = "64"
ddb.geometry.sectors = "32"
ddb.virtualHWVersion = "4"
```

By introspecting this *vmdk*, you can easily know that the *disk_type* is *streamOptimized*, and the *adapter_type* is *buslogic*. These metadata parameters are useful for the consumer of the image. In Compute, the workflow to instantiate a *streamOptimized* disk is different from the one to instantiate a *flat* disk. This new feature allows metadata extraction. You can achieve image introspection by calling the task API while importing the image. An administrator can override metadata settings.

1.1.5. Interoperable image import

The interoperable image import workflow enables you to import images in two ways:

- Use the **web-download** (default) method to import images from a URI.
- Use the **glance-direct** method to import images from a local file system.

1.1.6. Improving scalability with Image service caching

Use the glance-api caching mechanism to store copies of images on your local machine and retrieve them automatically to improve scalability. With Image service caching, the glance-api can run on multiple hosts. This means that it does not need to retrieve the same image from back-end storage multiple times. Image service caching does not affect any Image service operations.

To configure Image service caching with the Red Hat OpenStack Platform director (tripleo) heat templates, complete the following steps:

Procedure

1. In an environment file, set the value of the **GlanceCacheEnabled** parameter to **true**, which automatically sets the **flavor** value to **keystone+cachemanagement** in the **glance-api.conf** heat template:

```
parameter_defaults:
  GlanceCacheEnabled: true
```

2. Include the environment file in the **openstack overcloud deploy** command when you redeploy the overcloud.
3. Optional: Tune the **glance_cache_pruner** to an alternative frequency when you redeploy the overcloud. The following example shows a frequency of 5 minutes:

```
parameter_defaults:
  ControllerExtraConfig:
    glance::cache::pruner::minute: '*/5'
```

Adjust the frequency according to your needs to avoid file system full scenarios. Include the following elements when you choose an alternative frequency:

- The size of the files that you want to cache in your environment.

- The amount of available file system space.
- The frequency at which the environment caches images.

1.1.7. Image pre-caching

This feature is available in this release as a *Technology Preview*, and therefore is not fully supported by Red Hat. It should only be used for testing, and should not be deployed in a production environment. For more information about Technology Preview features, see [Scope of Coverage Details](#).

1.1.7.1. Configuring the default interval for periodic image pre-caching

Because the Red Hat OpenStack Platform director can now pre-cache images as part of the **glance-api** service, you no longer require **glance-registry** to pre-cache images. The default periodic interval is 300 seconds. You can increase or decrease the default interval based on your requirements.

Procedure

1. Add a new interval with the **ExtraConfig** parameter in an environment file on the undercloud according to your requirements:


```
parameter_defaults:
  ControllerExtraConfig:
    glance::config::glance_api_config:
      DEFAULT/cache_prefetcher_interval:
        value: '<300>'
```

Replace <300> with the number of seconds that you want as an interval to pre-cache images.

2. After you adjust the interval in the environment file in **/home/stack/templates/**, log in as the **stack** user and deploy the configuration:

```
$ openstack overcloud deploy --templates \
-e /home/stack/templates/<ENV_FILE>.yaml
```

Replace <ENV_FILE> with the name of the environment file that contains the **ExtraConfig** settings that you added.



IMPORTANT

If you passed any extra environment files when you created the overcloud, pass them again here using the **-e** option to avoid making undesired changes to the overcloud.

For more information about the **openstack overcloud deploy** command, see [Deployment command](#) in the *Director Installation and Usage* guide.

1.1.7.2. Using a periodic job to pre-cache an image

Prerequisite

To use a periodic job to pre-cache an image, you must use the **glance-cache-manage** command connected directly to the node where the **glance_api** service is running. Do not use a proxy, which hides the node that answers a service request. Because the undercloud might not have access to the network

where the **glance_api** service is running, run commands on the first overcloud node, which is called **controller-0** by default.

Complete the following prerequisite procedure to ensure that you run commands from the correct host, have the necessary credentials, and are also running the **glance-cache-manage** commands from inside the **glance-api** container.

1. Log in to the undercloud as the stack user and identify the provisioning IP address of **controller-0**:

```
(undercloud) [stack@site-undercloud-0 ~]$ openstack server list -f value -c Name -c
Networks | grep controller
overcloud-controller-1 ctlplane=192.168.24.40
overcloud-controller-2 ctlplane=192.168.24.13
overcloud-controller-0 ctlplane=192.168.24.71
(undercloud) [stack@site-undercloud-0 ~]$
```

2. To authenticate to the overcloud, copy the credentials that are stored in **/home/stack/overcloudrc**, by default, to **controller-0**:

```
$ scp ~/overcloudrc heat-admin@192.168.24.71:/home/heat-admin/
```

3. Connect to **controller-0**:

```
$ ssh heat-admin@192.168.24.71
```

4. On **controller-0** as the **heat-admin** user, identify the IP address of the **glance_api** service. In the following example, the IP address is **172.25.1.105**:

```
(overcloud) [root@controller-0 ~]# grep -A 10 '^listen glance_api' 1/var/lib/config-data/puppet-
generated/haproxy/etc/haproxy/haproxy.cfg
listen glance_api
server central-controller0-0.internalapi.redhat.local 172.25.1.105:9292 check fall 5 inter 2000
rise 2
```

5. Because the **glance-cache-manage** command is only available in the **glance_api** container, create a script to exec into that container where the environment variables to authenticate to the overcloud are already set. Create a script called **glance_pod.sh** in **/home/heat-admin** on **controller-0** with the following contents:

```
sudo podman exec -ti \
-e NOVA_VERSION=$NOVA_VERSION \
-e COMPUTE_API_VERSION=$COMPUTE_API_VERSION \
-e OS_USERNAME=$OS_USERNAME \
-e OS_PROJECT_NAME=$OS_PROJECT_NAME \
-e OS_USER_DOMAIN_NAME=$OS_USER_DOMAIN_NAME \
-e OS_PROJECT_DOMAIN_NAME=$OS_PROJECT_DOMAIN_NAME \
-e OS_NO_CACHE=$OS_NO_CACHE \
-e OS_CLOUDNAME=$OS_CLOUDNAME \
-e no_proxy=$no_proxy \
-e OS_AUTH_TYPE=$OS_AUTH_TYPE \
-e OS_PASSWORD=$OS_PASSWORD \
-e OS_AUTH_URL=$OS_AUTH_URL \
-e OS_IDENTITY_API_VERSION=$OS_IDENTITY_API_VERSION \
```



```
-e OS_COMPUTE_API_VERSION=$OS_COMPUTE_API_VERSION \
-e OS_IMAGE_API_VERSION=$OS_IMAGE_API_VERSION \
-e OS_VOLUME_API_VERSION=$OS_VOLUME_API_VERSION \
-e OS_REGION_NAME=$OS_REGION_NAME \
glance_api /bin/bash
```

6. Source the **overcloudrc** file and run the **glance_pod.sh** script to exec into the **glance_api** container with the necessary environment variables to authenticate to the overcloud Controller node.

```
[heat-admin@controller-0 ~]$ source overcloudrc
(overcloudrc) [heat-admin@central-controller-0 ~]$ bash glance_pod.sh
()[glance@controller-0 /]$
```

7. Use a command such as **glance image-list** to verify that the container can run authenticated commands against the overcloud.

```
()[glance@controller-0 /]$ glance image-list
+-----+-----+
| ID                      | Name                               |
+-----+-----+
| ad2f8daf-56f3-4e10-b5dc-d28d3a81f659 | cirros-0.4.0-x86_64-disk.img      |
+-----+-----+
()[glance@controller-0 /]$
```

Procedure

1. As the admin user, queue an image to cache:

```
$ glance-cache-manage --host=<HOST-IP> queue-image <IMAGE-ID>
```

Replace **<HOST-IP>** with the IP address of the Controller node where the **glance-api** container is running, and replace **<IMAGE-ID>** with the ID of the image that you want to queue. When you have queued the images that you want to pre-cache, the **cache_images** periodic job prefetches all queued images concurrently.



NOTE

Because the image cache is local to each node, if your Red Hat OpenStack Platform is deployed with HA (with 3, 5, or 7 Controllers) then you must specify the host address with the **--host** option when you run the **glance-cache-manage** command.

1. Run the following command to view the images in the image cache:

```
$ glance-cache-manage --host=<HOST-IP> list-cached
```

Replace **<HOST-IP>** with the IP address of the host in your environment.

Related information

You can use additional **glance-cache-manage** commands for the following purposes:

- **list-cached** to list all images that are currently cached.

- **list-queued** to list all images that are currently queued for caching.
- **queue-image** to queue an image for caching.
- **delete-cached-image** to purge an image from the cache.
- **delete-all-cached-images** to remove all images from the cache.
- **delete-queued-image** to delete an image from the cache queue.
- **delete-all-queued-images** to delete all images from the cache queue.

1.2. MANAGE IMAGES

The OpenStack Image service (glance) provides discovery, registration, and delivery services for disk and server images. It provides the ability to copy or snapshot a server image, and immediately store it away. Stored images can be used as a template to get new servers up and running quickly and more consistently than installing a server operating system and individually configuring services.

1.2.1. Creating an Image

This section provides you with the steps to manually create OpenStack-compatible images in the QCOW2 format using Red Hat Enterprise Linux 7 ISO files, Red Hat Enterprise Linux 6 ISO files, or Windows ISO files.

1.2.1.1. Use a KVM Guest Image With Red Hat OpenStack Platform

You can use a ready RHEL KVM guest QCOW2 image:

- [Red Hat Enterprise Linux 8 KVM Guest Image](#)
- [Red Hat Enterprise Linux 7 KVM Guest Image](#)
- [Red Hat Enterprise Linux 6 KVM Guest Image](#)

These images are configured with **cloud-init** and must take advantage of ec2-compatible metadata services for provisioning SSH keys in order to function properly.

Ready Windows KVM guest QCOW2 images are not available.



NOTE

For the KVM guest images:

- The **root** account in the image is disabled, but **sudo** access is granted to a special user named **cloud-user**.
- There is no **root** password set for this image.

The **root** password is locked in **/etc/shadow** by placing **!!** in the second field.

For an OpenStack instance, it is recommended that you generate an ssh keypair from the OpenStack dashboard or command line and use that key combination to perform an SSH public authentication to the instance as root.

When the instance is launched, this public key will be injected to it. You can then authenticate using the private key downloaded while creating the keypair.

If you do not want to use keypairs, you can use the **admin** password that has been set using the [Inject an admin Password Into an Instance](#) procedure.

If you want to create custom Red Hat Enterprise Linux or Windows images, see [Create a Red Hat Enterprise Linux 7 Image](#), [Create a Red Hat Enterprise Linux 6 Image](#) , or [Create a Windows Image](#).

1.2.1.2. Create Custom Red Hat Enterprise Linux or Windows Images

Prerequisites:

- Linux host machine to create an image. This can be any machine on which you can install and run the Linux packages.
- libvirt, virt-manager (run command **dnf groupinstall -y @virtualization**). This installs all packages necessary for creating a guest operating system.
- Libguestfs tools (run command **dnf install -y libguestfs-tools-c**). This installs a set of tools for accessing and modifying virtual machine images.
- A Red Hat Enterprise Linux 7 or 6 ISO file (see [RHEL 7.2 Binary DVD](#) or [RHEL 6.8 Binary DVD](#)) or a Windows ISO file. If you do not have a Windows ISO file, visit the [Microsoft TechNet Evaluation Center](#) and download an evaluation image.
- Text editor, if you want to change the **kickstart** files (RHEL only).



NOTE

In the following procedures, all commands with the **[root@host]#** prompt should be run on your host machine.

1.2.1.2.1. Create a Red Hat Enterprise Linux 7 Image

This section provides you with the steps to manually create an OpenStack-compatible image in the QCOW2 format using a Red Hat Enterprise Linux 7 ISO file.

1. Start the installation using **virt-install** as shown below:

```
[root@host]# qemu-img create -f qcow2 rhel7.qcow2 8G
[root@host]# virt-install --virt-type kvm --name rhel7 --ram 2048 \
--cdrom /tmp/rhel-server-7.2-x86_64-dvd.iso \
--disk rhel7.qcow2,format=qcow2 \
--network=bridge:virbr0 --graphics vnc,listen=0.0.0.0 \
--noautoconsole --os-type=linux --os-variant=rhel7
```

This launches an instance and starts the installation process.




NOTE

If the instance does not launch automatically, run the **virt-viewer** command to view the console:

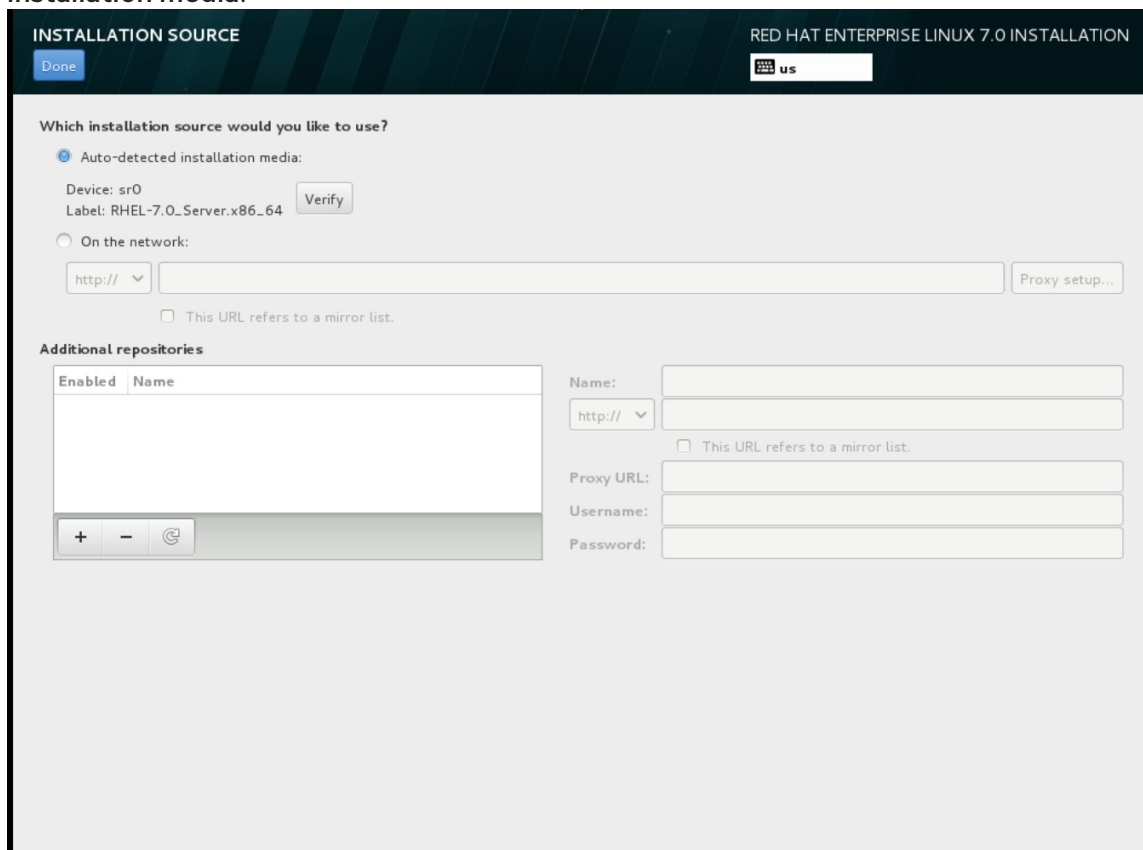
```
[root@host]# virt-viewer rhel7
```

2. Set up the virtual machine as follows:

- a. At the initial Installer boot menu, choose the **Install Red Hat Enterprise Linux 7.X** option.



- b. Choose the appropriate **Language** and **Keyboard** options.
- c. When prompted about which type of devices your installation uses, choose **Auto-detected installation media**.



- d. When prompted about which type of installation destination, choose **Local Standard Disks**.


INSTALLATION DESTINATION RED HAT ENTERPRISE LINUX 7.0 INSTALLATION

[Done](#) us

Device Selection
Select the device(s) you'd like to install to. They will be left untouched until you click on the main menu's "Begin Installation" button.

Local Standard Disks

2.04 GB



Virtio Block Device
vda / 2.04 GB free

Disks left unselected here will not be touched.

Specialized & Network Disks

[Add a disk...](#)

Disks left unselected here will not be touched.

Other Storage Options

Partitioning
☒ Automatically configure partitioning. ☐ I will configure partitioning.
☐ I would like to make additional space available.

Encryption
☐ Encrypt my data. *You'll set a passphrase later.*

[Full disk summary and bootloader...](#) 1 disk selected; 2.04 GB capacity; 2.04 GB free

For other storage options, choose **Automatically configure partitioning**.

- e. For software selection, choose **Minimal Install**.
- f. For network and host name, choose **eth0** for network and choose a **hostname** for your device. The default host name is **localhost.localdomain**.

- a. Choose the **root** password.

g. Choose the root password:

The installation process completes and the **Complete!** screen appears.

- After the installation is complete, reboot the instance and log in as the root user.
- Update the `/etc/sysconfig/network-scripts/ifcfg-eth0` file so it only contains the following values:

```
TYPE=Ethernet
DEVICE=eth0
ONBOOT=yes
BOOTPROTO=dhcp
NM_CONTROLLED=no
```

- Reboot the machine.
- Register the machine with the Content Delivery Network.

```
# sudo subscription-manager register
# sudo subscription-manager attach --pool=Valid-Pool-Number-123456
# sudo subscription-manager repos --enable=rhel-7-server-rpms
```

- Update the system:

```
# dnf -y update
```

- Install the **cloud-init** packages:

```
# dnf install -y cloud-utils-growpart cloud-init
```

- Edit the `/etc/cloud/cloud.cfg` configuration file and under **cloud_init_modules** add:

■

```
- resolv-conf
```

The **resolv-conf** option automatically configures the **resolv.conf** when an instance boots for the first time. This file contains information related to the instance such as **nameservers**, **domain** and other options.

10. Add the following line to **/etc/sysconfig/network** to avoid problems accessing the EC2 metadata service:

```
NOZEROCONF=yes
```

11. To ensure the console messages appear in the **Log** tab on the dashboard and the **nova console-log** output, add the following boot option to the **/etc/default/grub** file:

```
GRUB_CMDLINE_LINUX_DEFAULT="console=tty0 console=ttyS0,115200n8"
```

Run the **grub2-mkconfig** command:

```
# grub2-mkconfig -o /boot/grub2/grub.cfg
```

The output is as follows:

```
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-3.10.0-229.7.2.el7.x86_64
Found initrd image: /boot/initramfs-3.10.0-229.7.2.el7.x86_64.img
Found linux image: /boot/vmlinuz-3.10.0-121.el7.x86_64
Found initrd image: /boot/initramfs-3.10.0-121.el7.x86_64.img
Found linux image: /boot/vmlinuz-0-rescue-b82a3044fb384a3f9aeacf883474428b
Found initrd image: /boot/initramfs-0-rescue-b82a3044fb384a3f9aeacf883474428b.img
done
```

12. Un-register the virtual machine so that the resulting image does not contain the same subscription details for every instance cloned based on it:

```
# subscription-manager repos --disable=*
# subscription-manager unregister
# dnf clean all
```

13. Power off the instance:

```
# poweroff
```

14. Reset and clean the image using the **virt-sysprep** command so it can be used to create instances without issues:

```
[root@host]# virt-sysprep -d rhel7
```

15. Reduce image size using the **virt-sparsify** command. This command converts any free space within the disk image back to free space within the host:

```
[root@host]# virt-sparsify --compress /tmp/rhel7.qcow2 rhel7-cloud.qcow2
```

This creates a new **rhel7-cloud.qcow2** file in the location from where the command is run.

The **rhel7-cloud.qcow2** image file is ready to be uploaded to the Image service. For more information on uploading this image to your OpenStack deployment using the dashboard, see [Upload an Image](#).

1.2.1.2.2. Create a Red Hat Enterprise Linux 6 Image

This section provides you with the steps to manually create an OpenStack-compatible image in the QCOW2 format using a Red Hat Enterprise Linux 6 ISO file.

1. Start the installation using **virt-install**:

```
[root@host]# qemu-img create -f qcow2 rhel6.qcow2 4G
[root@host]# virt-install --connect=qemu:///system --network=bridge:virbr0 \
--name=rhel6 --os-type linux --os-variant rhel6 \
--disk path=rhel6.qcow2,format=qcow2,size=10,cache=none \
--ram 4096 --vcpus=2 --check-cpu --accelerate \
--hvm --cdrom=rhel-server-6.8-x86_64-dvd.iso
```

This launches an instance and starts the installation process.



NOTE

If the instance does not launch automatically, run the **virt-viewer** command to view the console:

```
[root@host]# virt-viewer rhel6
```

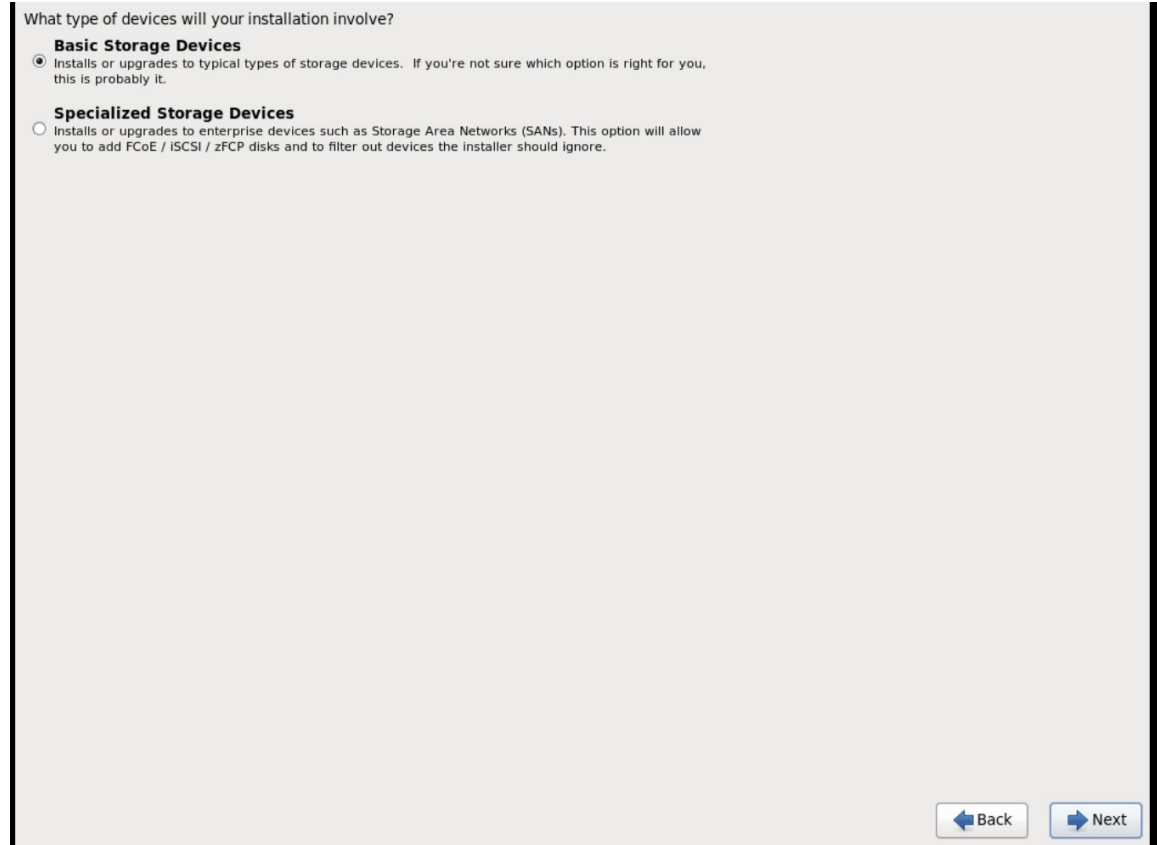
2. Set up the virtual machines as follows:
 - a. At the initial Installer boot menu, choose the **Install or upgrade an existing system** option.



Step through the installation prompts. Accept the defaults.

The installer checks for the disc and lets you decide whether you want to test your installation media before installation. Select **OK** to run the test or **Skip** to proceed without testing.

- b. Choose the appropriate **Language** and **Keyboard** options.
- c. When prompted about which type of devices your installation uses, choose **Basic Storage Devices**.



- d. Choose a **hostname** for your device. The default host name is **localhost.localdomain**.
- e. Set **timezone** and **root** password.

- f. Based on the space on the disk, choose the type of installation.

Which type of installation would you like?

- ☐ **Use All Space**
Removes all partitions on the selected device(s). This includes partitions created by other operating systems.
Tip: This option will remove data from the selected device(s). Make sure you have backups.
- ☐ **Replace Existing Linux System(s)**
Removes only Linux partitions (created from a previous Linux installation). This does not remove other partitions you may have on your storage device(s) (such as VFAT or FAT32).
Tip: This option will remove data from the selected device(s). Make sure you have backups.
- ☒ **Shrink Current System**
Shrinks existing partitions to create free space for the default layout.
- ☐ **Use Free Space**
Retains your current data and partitions and uses only the unpartitioned space on the selected device(s), assuming you have enough free space available.
- ☐ **Create Custom Layout**
Manually create your own custom layout on the selected device(s) using our partitioning tool.

☐ Encrypt system
☐ Review and modify partitioning layout

[< Back](#) [Next >](#)

- g. Choose the **Basic Server** install, which installs an SSH server.

The default installation of Red Hat Enterprise Linux is a basic server install. You can optionally select a different set of software now.

- ☒ Basic Server
- ☐ Database Server
- ☐ Web Server
- ☐ Identity Management Server
- ☐ Virtualization Host
- ☐ Desktop
- ☐ Software Development Workstation
- ☐ Minimal

Please select any additional repositories that you want to use for software installation.

- ☐ Hight Availability
- ☐ Load Balancer
- ☒ Red Hat Enterprise Linux

[Add additional software repositories](#) [Modify repository](#)

You can further customize the software selection now, or after install via the software management application.

☒ Customize later ☐ Customize now

[< Back](#) [Next >](#)

- h. The installation process completes and **Congratulations, your Red Hat Enterprise Linux installation is complete** screen appears.

3. Reboot the instance and log in as the **root** user.

4. Update the **/etc/sysconfig/network-scripts/ifcfg-eth0** file so it only contains the following values:

```
TYPE=Ethernet
DEVICE=eth0
ONBOOT=yes
BOOTPROTO=dhcp
NM_CONTROLLED=no
```

5. Reboot the machine.
6. Register the machine with the Content Delivery Network:

```
# sudo subscription-manager register
# sudo subscription-manager attach --pool=Valid-Pool-Number-123456
# sudo subscription-manager repos --enable=rhel-6-server-rpms
```

7. Update the system:

```
# dnf -y update
```

8. Install the **cloud-init** packages:

```
# dnf install -y cloud-utils-growpart cloud-init
```

9. Edit the **/etc/cloud/cloud.cfg** configuration file and under **cloud_init_modules** add:

```
- resolv-conf
```

The **resolv-conf** option automatically configures the **resolv.conf** configuration file when an instance boots for the first time. This file contains information related to the instance such as **nameservers**, **domain**, and other options.

10. To prevent network issues, create the **/etc/udev/rules.d/75-persistent-net-generator.rules** file as follows:

```
# echo "#" > /etc/udev/rules.d/75-persistent-net-generator.rules
```

This prevents **/etc/udev/rules.d/70-persistent-net.rules** file from being created. If **/etc/udev/rules.d/70-persistent-net.rules** is created, networking may not function properly when booting from snapshots (the network interface is created as "eth1" rather than "eth0" and IP address is not assigned).

11. Add the following line to **/etc/sysconfig/network** to avoid problems accessing the EC2 metadata service:

```
NOZEROCONF=yes
```

12. To ensure the console messages appear in the **Log** tab on the dashboard and the **nova console-log** output, add the following boot option to the **/etc/grub.conf**:

```
console=tty0 console=ttyS0,115200n8
```

- Un-register the virtual machine so that the resulting image does not contain the same subscription details for every instance cloned based on it:

```
# subscription-manager repos --disable=*
# subscription-manager unregister
# dnf clean all
```

- Power off the instance:

```
# poweroff
```

- Reset and clean the image using the **virt-sysprep** command so it can be used to create instances without issues:

```
[root@host]# virt-sysprep -d rhel6
```

- Reduce image size using the **virt-sparsify** command. This command converts any free space within the disk image back to free space within the host:

```
[root@host]# virt-sparsify --compress rhel6.qcow2 rhel6-cloud.qcow2
```

This creates a new **rhel6-cloud.qcow2** file in the location from where the command is run.



NOTE

You will need to manually resize the partitions of instances based on the image in accordance with the disk space in the flavor that is applied to the instance.

The **rhel6-cloud.qcow2** image file is ready to be uploaded to the Image service. For more information on uploading this image to your OpenStack deployment using the dashboard, see [Upload an Image](#)

1.2.1.2.3. Create a Windows Image

This section provides you with the steps to manually create an OpenStack-compatible image in the QCOW2 format using a Windows ISO file.

- Start the installation using **virt-install** as shown below:

```
[root@host]# virt-install --name=name \
--disk size=size \
--cdrom=path \
--os-type=windows \
--network=bridge:virbr0 \
--graphics spice \
--ram=RAM
```

Replace the values of the **virt-install** parameters as follows:

- name* – the name that the Windows guest should have.
- size* – disk size in GB.
- path* – the path to the Windows installation ISO file.

- *RAM* – the requested amount of RAM in MB.



NOTE

The **--os-type=windows** parameter ensures that the clock is set up correctly for the Windows guest, and enables its Hyper-V enlightenment features.

Note that **virt-install** saves the guest image as `/var/lib/libvirt/images/name.qcow2` by default. If you want to keep the guest image elsewhere, change the parameter of the **--disk** option as follows:

```
--disk path=filename,size=size
```


Replace *filename* with the name of the file which should store the guest image (and optionally its path); for example **path=win8.qcow2,size=8** creates an 8 GB file named **win8.qcow2** in the current working directory.

TIP

If the guest does not launch automatically, run the **virt-viewer** command to view the console:

```
[root@host]# virt-viewer name
```

2. Installation of Windows systems is beyond the scope of this document. For instructions on how to install Windows, see the relevant Microsoft documentation.
3. To allow the newly installed Windows system to use the virtualized hardware, you might need to install *virtio drivers*. To do so, first install the image, which you must attach as a CD-ROM drive to the Windows guest. To install the **virtio-win** package you must add the virtio ISO image to the guest, and install the virtio drivers. See [Installing KVM paravirtualized drivers for Windows virtual machines](#) in the *Configuring and managing virtualization* guide.
4. To complete the setup, download and execute [Cloudbase-Init](#) on the Windows system. At the end of the installation of Cloudbase-Init, select the **Run Sysprep** and **Shutdown** check boxes. The **Sysprep** tool makes the guest unique by generating an OS ID, which is used by certain Microsoft services.



IMPORTANT

Red Hat does not provide technical support for Cloudbase-Init. If you encounter an issue, [contact Cloudbase Solutions](#).

When the Windows system shuts down, the *name.qcow2* image file is ready to be uploaded to the Image service. For more information on uploading this image to your OpenStack deployment using the dashboard or the command line, see [Upload an Image](#).



NOTE

libosinfo data

The Compute Service has deprecated support for using libosinfo data to set default device models. Instead, use the following image metadata properties to configure the optimal virtual hardware for an instance:

- **os_distro**
- **os_version**
- **hw_cdrom_bus**
- **hw_disk_bus**
- **hw_scsi_model**
- **hw_vif_model**
- **hw_video_model**
- **hypervisor_type**

For more information on these metadata properties, see [Appendix A, Image Configuration Parameters](#).

1.2.2. Upload an image

1. In the dashboard, select **Project > Compute > Images**
2. Click **Create Image**.
3. Fill out the values, and click **Create Image** when finished.

Table 1.1. Image Options

Field	Notes
Name	Name for the image. The name must be unique within the project.
Description	Brief description to identify the image.
Image Source	Image source: Image Location or Image File . Based on your selection, the next field is displayed.
Image Location or Image File	<ul style="list-style-type: none"> • Select Image Location option to specify the image location URL. • Select Image File option to upload an image from the local disk.
Format	Image format (for example, qcow2).

Field	Notes
Architecture	Image architecture. For example, use i686 for a 32-bit architecture or x86_64 for a 64-bit architecture.
Minimum Disk (GB)	Minimum disk size required to boot the image. If this field is not specified, the default value is 0 (no minimum).
Minimum RAM (MB)	Minimum memory size required to boot the image. If this field is not specified, the default value is 0 (no minimum).
Public	If selected, makes the image public to all users with access to the project.
Protected	If selected, ensures only users with specific permissions can delete this image.

When the image has been successfully uploaded, its status is changed to **active**, which indicates that the image is available for use. Note that the Image service can handle even large images that take a long time to upload – longer than the lifetime of the Identity service token which was used when the upload was initiated. This is due to the fact that the Image service first creates a trust with the Identity service so that a new token can be obtained and used when the upload is complete and the status of the image is to be updated.



NOTE

You can also use the **glance image-create** command with the **property** option to upload an image. More values are available on the command line. For a complete listing, see [Image Configuration Parameters](#).

1.2.3. Update an image

1. In the dashboard, select **Project > Compute > Images**
2. Click **Edit Image** from the dropdown list.



NOTE

The **Edit Image** option is available only when you log in as an **admin** user. When you log in as a **demo** user, you have the option to **Launch an instance** or **Create Volume**.

3. Update the fields and click **Update Image** when finished. You can update the following values – name, description, kernel ID, ramdisk ID, architecture, format, minimum disk, minimum RAM, public, protected.
4. Click the drop-down menu and select **Update Metadata** option.
5. Specify metadata by adding items from the left column to the right one. In the left column, there are metadata definitions from the Image Service Metadata Catalog. Select **Other** to add metadata with the key of your choice and click **Save** when finished.

**NOTE**

You can also use the **glance image-update** command with the **property** option to update an image. More values are available on the command line; for a complete listing, see [Image Configuration Parameters](#).

1.2.4. Import an image

You can import images into the Image service (glance) using **web-download** to import an image from a URI and **glance-direct** to import an image from a local file system. The **web-download** method is enabled by default.

Import methods are configured by the cloud administrator. Run the **glance import-info** command to list available import options.

1.2.4.1. Import from a remote URI

You can use the **web-download** method to copy an image from a remote URI.

1. Create an image and specify the URI of the image to import.

```
glance image-create --uri <URI>
```

2. You can monitor the image's availability using the **glance image-show <image-ID>** command where the ID is the one provided during image creation.

The Image service web-download method uses a two-stage process to perform the import. First, it creates an image record. Second, it retrieves the image the specified URI. This method provides a more secure way to import images than the deprecated **copy-from** method used in Image API v1.

The URI is subject to optional blacklist and whitelist filtering as described in the Advanced OpenStack Customization Guide.

The Image Property Injection plugin may inject metadata properties to the image as described in the Advanced OpenStack Customization Guide. These injected properties determine which compute nodes the image instances are launched on.

1.2.4.2. Import from a local volume

The **glance-direct** method creates an image record, which generates an image ID. Once the image is uploaded to the service from a local volume, it is stored in a staging area and is made active after it passes any configured checks. The **glance-direct** method requires a shared staging area when used in a highly available (HA) configuration.

**NOTE**

Image uploads using the **glance-direct** method fail in an HA environment if a common staging area is not present. In an HA active-active environment, API calls are distributed to the glance controllers. The download API call could be sent to a different controller than the API call to upload the image. For more information about configuring the staging area, refer to the [Storage Configuration](#) section in the *Advanced OpenStack Customization Guide*.

The glance-direct method uses three different calls to import an image:

- **glance image-create**
- **glance image-stage**
- **glance image-import**

You can use the **glance image-create-via-import** command to perform all three of these calls in one command. In the example below, uppercase words should be replaced with the appropriate options.

```
glance image-create-via-import --container-format FORMAT --disk-format DISKFORMAT --name
NAME --file /PATH/TO/IMAGE
```

Once the image moves from the staging area to the back end location, the image is listed. However, it may take some time for the image to become active.

You can monitor the image's availability using the **glance image-show <image-ID>** command where the ID is the one provided during image creation.

1.2.5. Delete an image

1. In the dashboard, select **Project > Compute > Images**
2. Select the image you want to delete and click **Delete Images**.

1.2.6. Hide or unhide an image

You can hide public images from normal listings presented to users. For instance, you can hide obsolete CentOS 7 images and show only the latest version to simplify the user experience. Users can discover and use hidden images.

To hide an image:

```
glance image-update <image-id> --hidden 'true'
```

To create a hidden image, add the **--hidden** argument to the **glance image-create** command.

To unhide an image:

```
glance image-update <image-id> --hidden 'false'
```

1.2.7. Show hidden images

To list hidden images:

```
glance image-list --hidden 'true'
```

1.2.8. Enabling image conversion

With the **GlanceImageImportPlugins** parameter enabled, you can upload a QCOW2 image, and the Image service will convert it to RAW.

**NOTE**

Image conversion is automatically enabled when you use Red Hat Ceph Storage RBD to store images and boot Nova instances.

To enable image conversion, create an environment file that contains the following parameter value and include the new environment file with the **-e** option in the **openstack overcloud deploy** command:

```
parameter_defaults:
  GlanceImageImportPlugins:'image_conversion'
```

1.2.9. Converting an image to RAW format

Red Hat Ceph Storage can store, but does not support using, QCOW2 images to host virtual machine (VM) disks.

When you upload a QCOW2 image and create a VM from it, the compute node downloads the image, converts the image to RAW, and uploads it back into Ceph, which can then use it. This process affects the time it takes to create VMs, especially during parallel VM creation.

For example, when you create multiple VMs simultaneously, uploading the converted image to the Ceph cluster may impact already running workloads. The upload process can starve those workloads of IOPS and impede storage responsiveness.

To boot VMs in Ceph more efficiently (ephemeral back end or boot from volume), the glance image format must be RAW.

Converting an image to RAW may yield an image that is larger in size than the original QCOW2 image file. Run the following command before the conversion to determine the final RAW image size:

```
qemu-img info <image>.qcow2
```

To convert an image from QCOW2 to RAW format, do the following:

```
qemu-img convert -p -f qcow2 -O raw <original qcow2 image>.qcow2 <new raw image>.raw
```

1.2.9.1. Configuring Image Service to accept RAW and ISO only

Optionally, to configure the Image Service to accept only RAW and ISO image formats, deploy using an additional environment file that contains the following:

```
parameter_defaults:
  ExtraConfig:
    glance::config::api_config:
      image_format/disk_formats:
        value: "raw,iso"
```

1.2.10. Storing an image in RAW format

With the **GlanceImageImportPlugins** parameter enabled, run the following command to store a previously created image in RAW format:

```
$ glance image-create-via-import \  
  --disk-format qcow2 \  
  --container-format bare \  
  --name NAME \  
  --visibility public \  
  --import-method web-download \  
  --uri http://server/image.qcow2
```

- For **--name**, replace **NAME** with the name of the image; this is the name that will appear in **glance image-list**.
- For **--uri**, replace **http://server/image.qcow2** with the location and file name of the QCOW2 image.



NOTE

This command example creates the image record and imports it by using the **web-download** method. The glance-api downloads the image from the **--uri** location during the import process. If **web-download** is not available, **glanceclient** cannot automatically download the image data. Run the **glance import-info** command to list the available image import methods.

CHAPTER 2. IMAGE SERVICE WITH MULTIPLE STORES

The Red Hat OpenStack Platform Image service (glance) supports using multiple stores with distributed edge architecture so that you can have an image pool at every edge site. You can copy images between the central site, which is also known as the hub site, and the edge sites.

The image metadata contains the location of each copy. For example, an image present on two edge sites is exposed as a single UUID with three locations: the central site plus the two edge sites. This means you can have copies of image data that share a single UUID on many stores. For more information about locations, see [Understanding the location of images](#).

With an RBD image pool at every edge site, you can boot VMs quickly by using Ceph RBD copy-on-write (COW) and snapshot layering technology. This means that you can boot VMs from volumes and have live migration. For more information about layering with Ceph RBD, see [Ceph block device layering](#) in the *Block Device Guide*.


2.1. REQUIREMENTS OF STORAGE EDGE ARCHITECTURE

- A copy of each image must exist in the Image service at the central location.
- Prior to creating an instance at an edge site, you must have a local copy of the image at that edge site.
- Source the **centralrc** authentication file to schedule workloads at edge sites as well as at the central location. Authentication files that are automatically generated for edge sites are not needed.
- Images uploaded to an edge site must be copied to the central location before they can be copied to other edge sites.
- Use the Image service RBD driver for all edge sites. Mixed architecture is not supported.
- Multistack must be used with a single stack at each site.
- RBD must be the storage driver for the Image, Compute and Block Storage services.
- For each site, you must assign the same value to the **NovaComputeAvailabilityZone** and **CinderStorageAvailabilityZone** parameters.

2.2. IMPORT AN IMAGE TO MULTIPLE STORES

Use the interoperable image import workflow to import image data into multiple Ceph Storage clusters. You can import images into the Image service that are available on the local file system or through a web server.

If you import an image from a web server, the image can be imported into multiple stores at once. If the image is not available on a web server, you can import the image from a local file system into the central store and then copy it to additional stores. For more information, see [Copy an existing image to multiple stores](#).



IMPORTANT

Always store an image copy on the central site, even if there are no instances using the image at the central location. For more information about importing images into the Image service, see the [Distributed compute node and storage deployment](#) guide.


2.2.1. Manage image import failures

You can manage failures of the image import operation by using the **--allow-failure** parameter:

- If the value of the **--allow-failure** parameter to **true**, the image status becomes **active** after the first store successfully imports the data. This is the default setting. You can view a list of stores that failed to import the image data by using the **os_glance_failed_import** image property.
- If you set the value of the **--allow-failure** parameter to **false**, the image status only becomes **active** after all specified stores successfully import the data. Failure of any store to import the image data results in an image status of **failed**. The image is not imported into any of the specified stores.

2.2.2. Importing image data to multiple stores

Because the default setting of the **--allow-failure** parameter is **true**, you do not need to include the parameter in the command if it is acceptable for some stores to fail to import the image data.



NOTE

This procedure does not require all stores to successfully import the image data.

Procedure

1. Import image data to multiple, specified stores:

```
$ glance image-create-via-import \
--container-format bare \
--name IMAGE-NAME \
--import-method web-download \
--uri URI \
--stores STORE1,STORE2,STORE3
```

- Replace *IMAGE-NAME* with the name of the image you want to import.
- Replace *URI* with the URI of the image.
- Replace *STORE1*, *STORE2*, and *STORE3* with the names of the stores to which you want to import the image data.
- Alternatively, replace **--stores** with **--all-stores true** to upload the image to all the stores.



NOTE

The **glance image-create-via-import** command, which automatically converts the QCOW2 image to RAW format, works only with the **web-download** method. The **glance-direct** method is available, but it works only in deployments with a configured shared file system. For more information, see [Storing an image in RAW format](#).

2.2.3. Importing image data to multiple stores without failure

This procedure requires all stores to successfully import the image data.

Procedure

1. Import image data to multiple, specified stores:

```
$ glance image-create-via-import \
--container-format bare \
--name IMAGE-NAME \
--import-method web-download \
--uri URI \
--stores STORE1,STORE2
```

- Replace *IMAGE-NAME* with the name of the image you want to import.
- Replace *URI* with the URI of the image.
- Replace *STORE1*, *STORE2*, and *STORE3* with the names of stores to which you want to copy the image data.
- Alternatively, replace **--stores** with **--all-stores true** to upload the image to all the stores.



NOTE

With the **--allow-failure** parameter set to **false**, the Image service does not ignore stores that fail to import the image data. You can view the list of failed stores with the image property **os_glance_failed_import**. For more information see [Checking the progress of image import operation](#).

2. Verify that the image data was added to specific stores:

```
$ glance image-show IMAGE-ID | grep stores
```

Replace *IMAGE-ID* with the ID of the original existing image.

The output displays a comma-delimited list of stores.

2.2.4. Importing image data to a single store

You can import image data to a single store.

Procedure

1. Import image data to a single store:

```
$ glance image-create-via-import \
--container-format bare \
--name IMAGE-NAME \
--import-method web-download \
--uri URI \
--store STORE
```

- Replace *IMAGE-NAME* with the name of the image you want to import.
- Replace *URI* with the URI of the image.
- Replace *STORE* with the name of the store to which you want to copy the image data.

**NOTE**

If you do not include the options of **--stores**, **--all-stores**, or **--store** in the command, the Image service creates the image in the central store.

2. Verify that the image data was added to specific store:

```
$ glance image-show IMAGE-ID | grep stores
```

Replace *IMAGE-ID* with the ID of the original existing image.

The output displays a comma-delimited list of stores.

2.2.5. Checking the progress of the image import operation

The interoperable image import workflow sequentially imports image data into stores. The size of the image, the number of stores, and the network speed between the central site and the edge sites impact how long it takes for the image import operation to complete.

You can follow the progress of the image import by looking at two image properties, which appear in notifications sent during the image import operation:

- The **os_glance_importing_to_stores** property lists the stores that have not imported the image data. At the beginning of the import, all requested stores show up in the list. Each time a store successfully imports the image data, the Image service removes the store from the list.
- The **os_glance_failed_import** property lists the stores that fail to import the image data. This list is empty at the beginning of the image import operation.

**NOTE**

In the following procedure, the environment has three Ceph Storage clusters: the **central** store and two stores at the edge, **dcn0** and **dcn1**.

Procedure

1. Verify that the image data was added to specific stores:

```
$ glance image-show IMAGE-ID
```

Replace *IMAGE-ID* with the ID of the original existing image.

The output displays a comma-delimited list of stores similar to the following example snippet:

```
| os_glance_failed_import      |
| os_glance_importing_to_stores | central,dcn0,dcn1
| status                      | importing
```

2. Monitor the status of the image import operation. When you precede a command with **watch**, the command output refreshes every two seconds.

```
$ watch glance image-show IMAGE-ID
```

Replace *IMAGE-ID* with the ID of the original existing image.

The status of the operation changes as the image import operation progresses:

```
| os_glance_failed_import      |
| os_glance_importing_to_stores | dcn0,dcn1
| status                       | importing
```

Output that shows that an image failed to import resembles the following example:

```
| os_glance_failed_import      | dcn0
| os_glance_importing_to_stores | dcn1
| status                       | importing
```

After the operation completes, the status changes to active:

```
| os_glance_failed_import      | dcn0
| os_glance_importing_to_stores |
| status                       | active
```

2.3. COPY AN EXISTING IMAGE TO MULTIPLE STORES

This feature enables you to copy existing images using Red Hat OpenStack Image service (glance) image data into multiple Ceph Storage stores at the edge by using the interoperable image import workflow.



NOTE

The image must be present at the central site before you copy it to any edge sites. Only the image owner or administrator can copy existing images to newly added stores.

You can copy existing image data either by setting **--all-stores** to **true** or by specifying specific stores to receive the image data.

- The default setting for the **--all-stores** option is **false**. If **--all-stores** is **false**, you must specify which stores receive the image data by using **--stores STORE1,STORE2**. If the image data is already present in any of the specified stores, the request fails.
- If you set **all-stores** to **true**, and the image data already exists in some of the stores, then those stores are excluded from the list.

After you specify which stores receive the image data, the Image service copies data from the central site to a staging area. Then the Image service imports the image data by using the interoperable image import workflow. For more information, see [Importing an image to multiple stores](#).



IMPORTANT

Red Hat recommends that administrators carefully avoid closely timed image copy requests. Two closely timed copy-image operations for the same image causes race conditions and unexpected results. Existing image data remains as it is, but copying data to new stores fails.

2.3.1. Copying an image to all stores

Use the following procedure to copy image data to all available stores.

Procedure

1. Copy image data to all available stores:

```
$ glance image-import IMAGE-ID \
--all-stores true \
--import-method copy-image
```

Replace *IMAGE-ID* with the name of the image you want to copy.

2. Confirm that the image data successfully replicated to all available stores:

```
$ glance image-list --include-stores
```

For information about how to check the status of the image import operation, see [Checking the progress of the image import operation](#).

2.3.2. Copying an image to specific stores

Use the following procedure to copy image data to specific stores.

Procedure

1. Copy image data to specific stores:

```
$ glance image-import IMAGE-ID \
--stores STORE1,STORE2 \
--import-method copy-image
```

- Replace *IMAGE-ID* with the name of the image you want to copy.
- Replace *STORE1* and *STORE2* with the names of the stores to which you want to copy the image data.


2. Confirm that the image data successfully replicated to the specified stores:

```
$ glance image-list --include-stores
```

For information about how to check the status of the image import operation, see [Checking the progress of the image import operation](#).

2.4. DELETING AN IMAGE FROM A SPECIFIC STORE

This feature enables you to delete an existing image copy on a specific store using Red Hat OpenStack Image service (glance).



IMPORTANT

This feature is available in this release as a *Technology Preview*, and therefore is not fully supported by Red Hat. It should only be used for testing, and should not be deployed in a production environment. For more information about Technology Preview features, see [Scope of Coverage Details](#).

Procedure

Delete an image from a specific store:

```
$ glance stores-delete --store _STORE_ID_ _IMAGE_ID_
```

- Replace `_STORE_ID_` with the name of the store on which the image copy should be deleted.
- Replace `_IMAGE_ID_` with the ID of the image you want to delete.



WARNING

Using **glance image-delete** will permanently delete the image across all the sites. All image copies will be deleted, as well as the image instance and metadata.

2.5. UNDERSTANDING THE LOCATIONS OF IMAGES

Although an image can be present on multiple sites, there is only a single UUID for a given image. The image metadata contains the locations of each copy. For example, an image present on two edge sites is exposed as a single UUID with three locations: the central site plus the two edge sites.

Procedure

1. Show the sites on which a copy of the image exists:

```
$ glance image-show ID | grep "stores"

| stores | default_backend,dcn1,dcn2
```

In the example, the image is present on the central site, the **default_backend**, and on the two edge sites **dcn1** and **dcn2**.

2. Alternatively, you can run the **glance image-list** command with the **--include-stores** option to see the sites where the images exist:

```
$ glance image-list --include-stores

| ID | Name | Stores |
| 2bd882e7-1da0-4078-97fe-f1bb81f61b00 | cirros | default_backend,dcn1,dcn2
```

3. List the image locations properties to show the details of each location:

```
$ openstack image show ID -c properties

| properties |
(--- cut ---)
locations='[{"url": "rbd://79b70c32-df46-4741-93c0-8118ae2ae284/images/2bd882e7-1da0-4078-97fe-f1bb81f61b00/snap", "metadata": {"store": "default_backend"}}, {"url": "rbd://63df2767-
```

```
8ddb-4e06-8186-8c155334f487/images/2bd882e7-1da0-4078-97fe-f1bb81f61b00/snap',
'metadata': {'store': 'dcn1'}}, {'url': 'rbd://1b324138-2ef9-4ef9-bd9e-
aa7e6d6ead78/images/2bd882e7-1da0-4078-97fe-f1bb81f61b00/snap', 'metadata': {'store':
'dcn2'}}}],
(--- cut ---)
```

The image properties show the different Ceph RBD URIs for the location of each image.

In the example, the central image location URI is:

```
rbd://79b70c32-df46-4741-93c0-8118ae2ae284/images/2bd882e7-1da0-4078-97fe-
f1bb81f61b00/snap', 'metadata': {'store': 'default_backend'}}
```

The URI is composed of the following data:

- **79b70c32-df46-4741-93c0-8118ae2ae284** corresponds to the central Ceph FSID. Each Ceph cluster has a unique FSID.
- The default value for all sites is **images**, which corresponds to the Ceph pool on which the images are stored.
- **2bd882e7-1da0-4078-97fe-f1bb81f61b00** corresponds to the image UUID. The UUID is the same for a given image regardless of its location.
- The metadata shows the glance store to which this location maps. In this example, it maps to the **default_backend**, which is the central hub site.

CHAPTER 3. CONFIGURING THE COMPUTE (NOVA) SERVICE

Use environment files to customize the Compute (nova) service. Puppet generates and stores this configuration in the `/var/lib/config-data/puppet-generated/<nova_container>/etc/nova/nova.conf` file. Use the following configuration methods to customize the Compute service configuration:

- **Heat parameters** – as detailed in the [Compute \(nova\) Parameters](#) section in the *Overcloud Parameters* guide. For example:

```
parameter_defaults:
  NovaSchedulerDefaultFilters:
  AggregateInstanceExtraSpecsFilter,RetryFilter,ComputeFilter,ComputeCapabilitiesFilter,Image
  PropertiesFilter
  NovaNfsEnabled: true
  NovaNfsShare: '192.0.2.254:/export/nova'
  NovaNfsOptions: 'context=system_u:object_r:nfs_t:s0'
  NovaNfsVersion: '4.2'
```

- **Puppet parameters** – as defined in `/etc/puppet/modules/nova/manifests/*`:

```
parameter_defaults:
  ComputeExtraConfig:
    nova::compute::force_raw_images: True
```



NOTE

Only use this method if an equivalent heat parameter does not exist.

- **Manual hieradata overrides** – for customizing parameters when no heat or Puppet parameter exists. For example, the following sets the `timeout_nbd` in the **[DEFAULT]** section on the Compute role:

```
parameter_defaults:
  ComputeExtraConfig:
    nova::config::nova_config:
      DEFAULT/timeout_nbd:
        value: '20'
```



WARNING

If a heat parameter exists, it must be used instead of the Puppet parameter; if a Puppet parameter exists, but not a heat parameter, then the Puppet parameter must be used instead of the manual override method. The manual override method must only be used if there is no equivalent heat or Puppet parameter.

TIP

Follow the guidance in [Identifying Parameters to Modify](#) to determine if a heat or Puppet parameter is available for customizing a particular configuration.

See [Parameters](#) in the *Advanced Openstack Customization* guide for further details on configuring overcloud services.

3.1. CONFIGURING MEMORY FOR OVERALLOCATION

When you use memory overcommit (**NovaRAMAllocationRatio** ≥ 1.0), you need to deploy your overcloud with enough swap space to support the allocation ratio.

**NOTE**

If your **NovaRAMAllocationRatio** parameter is set to < 1 , follow the RHEL recommendations for swap size. For more information, see [Recommended system swap space](#) in the RHEL *Managing Storage Devices* guide.

Prerequisites

- You have calculated the swap size your node requires. For more information, see [Section 3.3, “Calculating swap size”](#).

Procedure

1. Copy the **/usr/share/openstack-tripleo-heat-templates/environments/enable-swap.yaml** file to your environment file directory:

```
$ cp /usr/share/openstack-tripleo-heat-templates/environments/enable-swap.yaml
/home/stack/templates/enable-swap.yaml
```

2. Configure the swap size by adding the following parameters to your **enable-swap.yaml** file:

```
parameter_defaults:
  swap_size_megabytes: <swap size in MB>
  swap_path: <full path to location of swap, default: /swap>
```

3. To apply this configuration, add the **enable_swap.yaml** environment file to the stack with your other environment files and deploy the overcloud:

```
(undercloud) $ openstack overcloud deploy --templates \
-e [your environment files] \
-e /home/stack/templates/enable-swap.yaml \
```

3.2. CALCULATING RESERVED HOST MEMORY ON COMPUTE NODES

To determine the total amount of RAM to reserve for host processes, you need to allocate enough memory for each of the following:

- The resources that run on the node, for instance, OSD consumes 3 GB of memory.
- The emulator overhead required to visualize instances on a host.

- The hypervisor for each instance.

After you calculate the additional demands on memory, use the following formula to help you determine the amount of memory to reserve for host processes on each node:

```
NovaReservedHostMemory = total_RAM - ( vm_no * (avg_instance_size + overhead)) + (resource1 * resource_ram) + (resource_n * resource_ram)
```

- Replace **vm_no** with the number of instances.
- Replace **avg_instance_size** with the average amount of memory each instance can use.
- Replace **overhead** with the hypervisor overhead required for each instance.
- Replace **resource1** with the number of a resource type on the node.
- Replace **resource_ram** with the amount of RAM each resource of this type requires.

3.3. CALCULATING SWAP SIZE

The allocated swap size must be large enough to handle any memory overcommit. You can use the following formulas to calculate the swap size your node requires:

- $\text{overcommit_ratio} = \text{NovaRAMAllocationRatio} - 1$
- Minimum swap size (MB) = $(\text{total_RAM} * \text{overcommit_ratio}) + \text{RHEL_min_swap}$
- Recommended (maximum) swap size (MB) = $\text{total_RAM} * (\text{overcommit_ratio} + \text{percentage_of_RAM_to_use_for_swap})$

The **percentage_of_RAM_to_use_for_swap** variable creates a buffer to account for QEMU overhead and any other resources consumed by the operating system or host services.

For instance, to use 25% of the available RAM for swap, with 64GB total RAM, and **NovaRAMAllocationRatio** set to **1**:

- Recommended (maximum) swap size = $64000 \text{ MB} * (0 + 0.25) = 16000 \text{ MB}$

For information on how to calculate the **NovaReservedHostMemory** value, see [Section 3.2, “Calculating reserved host memory on Compute nodes”](#).

For information on how to determine the **RHEL_min_swap** value, see [Recommended system swap space](#) in the RHEL *Managing Storage Devices* guide.

CHAPTER 4. SCALING DEPLOYMENTS WITH COMPUTE CELLS

You can use cells to divide Compute nodes in large deployments into groups, each with a message queue and dedicated database that contains instance information.

By default, the director installs the overcloud with a single cell for all Compute nodes. This single-cell deployment contains all instances and instance metadata. For larger deployments, you can deploy the overcloud with multiple cells to accommodate a larger number of Compute nodes.

In multi-cell deployments, each cell runs standalone copies of the cell-specific components and stores instance metadata only for instances in that cell. Global information and cell mappings are stored in the global Controller cell, which helps with security and recovery in case one of the cells fails.

You can add cells to your environment when you install a new overcloud or at any time afterwards.

4.1. CELL COMPONENTS

In single-cell deployments, all components are contained in the same cell. In multi-cell deployments, the global services run on the main Controller cell, and each Compute cell runs standalone copies of the cell-specific components and contains the database and message queue for the Compute nodes in that cell.

Global components

The following components are deployed in a Controller cell once for each overcloud, regardless of the number of Compute cells.

Compute API

Provides the external REST API to users.

Scheduler

Determines to which Compute node to assign the instances.

Placement service

Monitors and allocates Compute resources to the instances.

API database

Used by the Compute API and the Compute scheduler services to track location information about instances, and provides a temporary location for instances that are built but not scheduled.

In multi-cell deployments, this database also contains *cell mappings* that specify the database connection for each cell.

cell0 database

Dedicated database for information about instances that failed to be scheduled.

Super conductor

In multi-cell deployments, this service coordinates between the global services and each Compute cell, and also sends failed instance information to the **cell0** database.



NOTE

This component exists only in multi-cell deployments.

Cell-specific components

The following components are deployed in each Compute cell.

Cell database

Contains most of the information about instances. Used by the global API, the conductor, and the Compute services.

Conductor

Coordinates database queries and long-running tasks from the global services, and insulates Compute nodes from direct database access.

Message queue

Messaging service used by all services to communicate with each other within the cell and with the global services.

Configuration files

The overcloud includes configuration files that define the following information for the Compute cells:

- **[DEFAULT]/transport_url**: Message queue endpoint for each cell.
- **[DATABASE]/connection**: Database connection for each cell.
- **[API_DATABASE]/connection**: Routing and placement information for the global components.
- (Multi-cell deployments only) Cell mapping records to be stored in the global API database.

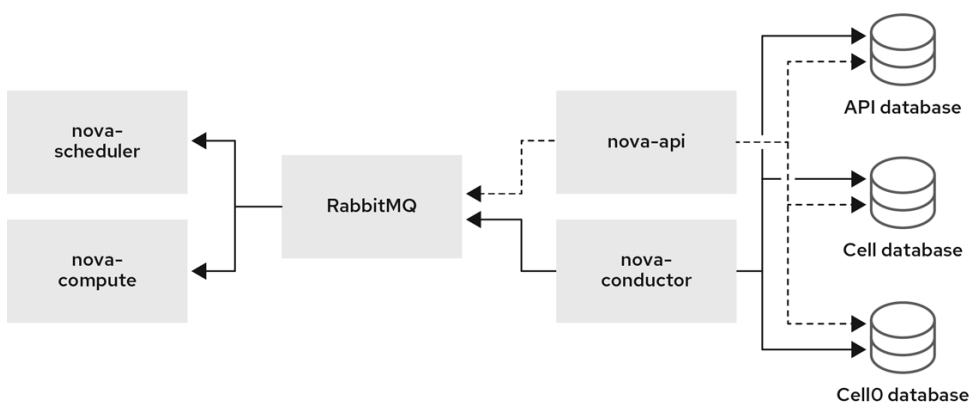
This information is extracted from the overcloud when you deploy the multi-cell environment, as described in [Section 4.4, "Deploying a multi-cell overcloud"](#).

4.2. CELL DEPLOYMENTS ARCHITECTURE

Each deployment type allows you to optimize your overcloud for different use-cases.

Single-cell deployment architecture (default)

The following diagram shows an example of the basic structure and interaction in a default single-cell overcloud.



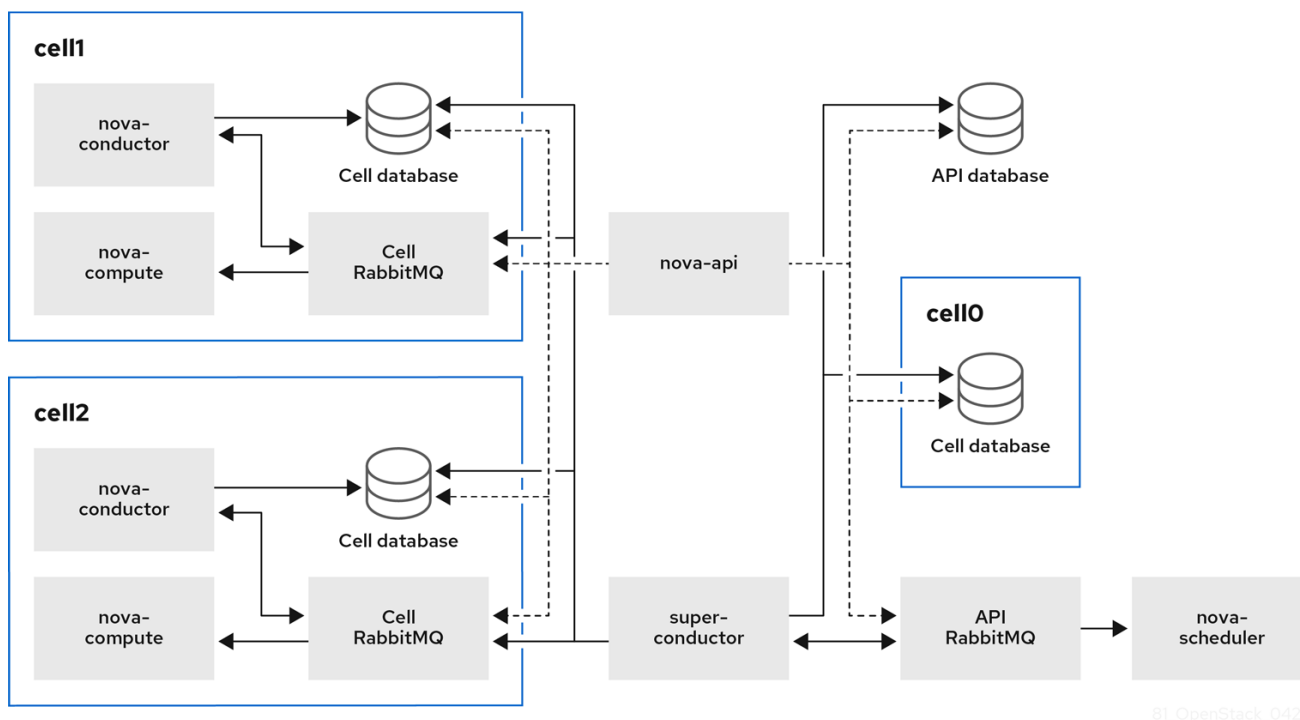
81_OpenStack_0420

In this deployment, all services are configured to use a single conductor to communicate between the Compute API and the Compute nodes, and a single database stores all live instance data.

In smaller deployments this configuration might be sufficient, but if any API-level (global) service or the database fails, the entire Compute deployment cannot send or receive information, regardless of high availability configurations.

Multi-cell deployment architecture (custom)

The following diagram shows an example of the basic structure and interaction in a custom multi-cell overcloud.



In this deployment, the Compute nodes are divided to multiple cells, each with their own conductor, database, and message queue. The global services use the super conductor to communicate with each cell, and the global database contains only information required for the whole overcloud.

The cell-level services cannot access global services directly. This isolation provides additional security and fail-safe capabilities in case of cell failure.



IMPORTANT

In Edge deployments, you must deploy the first cell on the central site, therefore, do not deploy the first cell on any of the edge sites. Do not run any Compute services on the first cell. Instead, deploy each new cell containing the Compute nodes separately on the edge sites.

4.3. CONSIDERATIONS FOR MULTI-CELL DEPLOYMENTS

Maximum number of Compute nodes in a multi-cell deployment

The maximum number of Compute nodes is 500 across all cells.

SSL/TLS

You cannot enable SSL/TLS on the overcloud.

Cross-cell instance migrations

Migrating an instance from a host in one cell to a host in another cell is not supported. This limitation affects the following operations:

- cold migration
- live migration
- unshelve
- resize
- evacuation

Service quotas

Compute service quotas are calculated dynamically at each resource consumption point, instead of statically in the database. In multi-cell deployments, unreachable cells cannot provide usage information in real-time, which might cause the quotas to be exceeded when the cell is reachable again.

You can use the Placement service and API database to configure the quota calculation to withstand failed or unreachable cells.

API database

The Compute API database is always global for all cells and cannot be duplicated for each cell.

Console proxies

You must configure console proxies for each cell, because console token authorizations are stored in cell databases. Each console proxy server needs to access the **database.connection** information of the corresponding cell database.

Template URLs in cell mappings

You can create templates for the **--database_connection** and **--transport-url** in cell mappings with variables that are dynamically updated each time you query the global database. The values are taken from the configuration files of the Compute nodes.

The format of a template URL is as follows:

```
{scheme}://{username}:{password}@{hostname}/{path}
```

The following table shows the variables that you can use in cell mapping URLs:

Variable	Description
scheme	Prefix before <code>://</code>
username	User name
password	Password
hostname	Host name or IP address
port	Port number (must be specified)
path	Path to the directory in the host (without leading slash)

Variable	Description
query	Full query with string arguments (without leading question mark)
fragment	Path after the first hash # sign

Compute metadata API

You can run the Compute metadata API globally or in each cell. Choose one of the following:

- If you have networks that cover multiple cells, you need to run the metadata API globally so that it can bridge between the cells. In this case, the metadata API needs to access the **api_database.connection** information.
- If you have networks in separate segments for each cell, you can run the metadata API separately in each cell. This configuration can improve performance and data isolation. In this case, **neutron-metadata-agent** service point to the corresponding **nova-api-metadata** service.

You use the **api.local_metadata_per_cell** configuration option to set which method to implement. For details on configuring this option, see the *Create environment files with cell parameters* section in [Section 4.4, “Deploying a multi-cell overcloud”](#).

4.4. DEPLOYING A MULTI-CELL OVERCLOUD

Deploying a multi-cell overcloud includes the following stages:

1. Extracting parameter information from the default first cell in the basic overcloud. This cell becomes the global Controller after you redeploy the overcloud.
2. Configuring a custom role and flavor for the cell.
3. Creating an environment file with cell-specific parameters.
4. Redeploying the overcloud with the new cell stack.



NOTE

- This process adds one cell to the overcloud. Repeat these steps for each additional cell you want to deploy in the overcloud.
- In this procedure, the name of the new cell is **cell1**. Replace the name in all commands with the actual cell name.

Prerequisites

- Deploy a basic overcloud with the required number of Controller and Compute nodes.
- Review the requirements and limitations for a multi-cell overcloud as described in [Section 4.3, “Considerations for multi-cell deployments”](#).

Extract parameter information from the overcloud

1. Create a new directory for the new cell and export the contents to the new directory. For example:

```
$ source ~/stackrc
(undercloud) $ mkdir cell1
(undercloud) $ export DIR=cell1
```

2. Export the **EndpointMap**, **HostsEntry**, **AllNodesConfig**, **GlobalConfig** parameters, and the password information from the overcloud to a new environment file for the cell. For example:

```
(undercloud) $ openstack overcloud cell export cell1 -o cell1/cell1-ctrl-input.yaml
```



NOTE

If the environment file already exists, run the command with the **--force-overwrite** or **-f** option.

Configure a custom role for a cell

1. Add the **CellController** role to your roles data file and regenerate the file. For example:

```
(undercloud) $ openstack overcloud roles generate --roles-path \
    /usr/share/openstack-tripleo-heat-templates/roles \
    -o $DIR/cell_roles_data.yaml Compute CellController
```

The **CellController** custom role includes the services from the default **Compute** role and additional configuration for the following services:

- Galera database
 - RabbitMQ
 - **nova-conductor**
 - **nova novnc proxy**
 - **nova metadata** (only in case you set the **NovaLocalMetadataPerCell** parameter)
2. In case you want to divide your network between the global Controller and the cells, configure network access in the roles file that you created. For example:

```
name: Compute
description: |
  Basic Compute Node role
CountDefault: 1
# Create external Neutron bridge (unset if using ML2/OVS without DVR)
tags:
  - external_bridge
networks:
  InternalApi:
    subnet: internal_api_cell1
  Tenant:
    subnet: tenant_subnet
  Storage:
    subnet: storage_cell1
```

```

...
- name: CellController
  description: |
    CellController role for the nova cell_v2 controller services
  CountDefault: 1
  tags:
    - primary
    - controller
  networks:
    External:
      subnet: external_cell1
    InternalApi:
      subnet: internal_api_cell1
    Storage:
      subnet: storage_cell1
    StorageMgmt:
      subnet: storage_mgmt_cell1
    Tenant:
      subnet: tenant_subnet

```

Configure a flavor and tag nodes to a cell

1. Create the **cellcontroller** flavor to tag nodes that you want to allocate to the cell. For example:

```

(undercloud) $ openstack flavor create --id auto --ram 4096 --disk 40 --vcpus 1 cellcontroller
(undercloud) $ openstack flavor set --property "cpu_arch"="x86_64" \
    --property "capabilities:boot_option"="local" \
    --property "capabilities:profile"="cellcontroller" \
    --property "resources:CUSTOM_BAREMETAL=1" \
    --property "resources:DISK_GB=0" \
    --property "resources:MEMORY_MB=0" \
    --property "resources:VCPU=0" \
    cellcontroller

```

2. Tag each node that you want to assign to the cell with the **cellcontroller** profile.

```

(undercloud) $ openstack baremetal node set --property \
    capabilities='profile:cellcontroller,boot_option:local' <NODE_UUID>

```

Replace **<NODE_UUID>** with the actual ID of the Compute node that you want to assign to the cell.

Create environment files with cell parameters

1. Create a new environment file in the directory for the cell, such as */cell1/cell1.yaml*, and add the following parameters:

```

resource_registry:
  # since the same networks are used in this example, the
  # creation of the different networks is omitted
  OS::TripleO::Network::External: OS::Heat::None
  OS::TripleO::Network::InternalApi: OS::Heat::None
  OS::TripleO::Network::Storage: OS::Heat::None
  OS::TripleO::Network::StorageMgmt: OS::Heat::None
  OS::TripleO::Network::Tenant: OS::Heat::None

```

```

OS::TripleO::Network::Management: OS::Heat::None
OS::TripleO::Network::Ports::OVNDBsVipPort: /usr/share/openstack-tripleo-heat-
templates/network/ports/noop.yaml
OS::TripleO::Network::Ports::RedisVipPort: /usr/share/openstack-tripleo-heat-
templates/network/ports/noop.yaml

parameter_defaults:
  # CELL Parameter to reflect that this is an additional CELL
  NovaAdditionalCell: True

  # mapping of the CellController flavor to the CellController role
  CellControllerFlavor: cellcontroller

  # The DNS names for the VIPs for the cell
  CloudName: cell1.ooo.test
  CloudNameInternal: cell1.internalapi.ooo.test
  CloudNameStorage: cell1.storage.ooo.test
  CloudNameStorageManagement: cell1.storagemgmt.ooo.test
  CloudNameCtlplane: cell1.ctlplane.ooo.test

  # Flavors used for the cell controller and computes
  OvercloudCellControllerFlavor: cellcontroller
  OvercloudComputeFlavor: compute

  # Number of controllers/computes in the cell
  CellControllerCount: 3
  ComputeCount: 1

  # Compute node name (must be unique)
  ComputeHostnameFormat: 'cell1-compute-%index%'

  # default gateway
  ControlPlaneStaticRoutes:
    - ip_netmask: 0.0.0.0/0
      next_hop: 192.168.24.1
      default: true
  DnsServers:
    - x.x.x.x

```

Change the parameter values in this example according to your deployment needs.

2. Depending on your network configuration, you might need to allocate a network resource to the cell. Add the following parameter if you need to register cells to the network:

```

resource_registry:
  OS::TripleO::CellController::Net::SoftwareConfig: single-nic-vlans/controller.yaml
  OS::TripleO::Compute::Net::SoftwareConfig: single-nic-vlans/compute.yaml

```

3. If you divide your network between the global Controller and the cells and want to run the Compute metadata API in each cell instead of in the global Controller, add the following parameter:

```

parameter_defaults:
  NovaLocalMetadataPerCell: True

```

**NOTE**

- The parameters in this file restrict the overcloud to use a single network for all cells.
- The Compute host names must be unique across all cells.

4. Copy the **network_data.yaml** file and name it according to the cell name. For example:

```
(undercloud) $ cp /usr/share/openstack-tripleo-heat-templates/network_data.yaml
cell1/network_data-ctrl.yaml
```

5. Add the UUIDs for the network components you want to reuse for the cells to the new network data file.

```
external_resource_network_id: [EXISTING_NETWORK_UUID]
external_resource_subnet_id: [EXISTING_SUBNET_UUID]
external_resource_segment_id: [EXISTING_SEGMENT_UUID]
external_resource_vip_id: [EXISTING_VIP_UUID]
```

(Optional) Configure networking for segmented networks

If you want to divide your network between the global Controller and the Compute cells, create an environment file such as **routes.yaml** and add the routing information and virtual IP address (VIP) information for the cell. For example:

```
parameter_defaults:
  InternalApiInterfaceRoutes:
    - destination: 172.17.2.0/24
      nexthop: 172.16.2.254
  StorageInterfaceRoutes:
    - destination: 172.17.1.0/24
      nexthop: 172.16.1.254
  StorageMgmtInterfaceRoutes:
    - destination: 172.17.3.0/24
      nexthop: 172.16.3.254

parameter_defaults:
  VipSubnetMap:
    InternalApi: internal_api_cell1
    Storage: storage_cell1
    StorageMgmt: storage_mgmt_cell1
    External: external_cell1
```

(Optional) Configure networking for Edge sites

To distribute Compute nodes across Edge sites, create one environment file for the main Controller cell and separate environment files for each Compute cell in that Edge site.

- In the primary environment file, set the **ComputeCount** parameter to **0** in the Controller cell. This cell is separate from the Edge site Compute cells, which will contain the actual Compute nodes.
- In the Compute cell environment files, add the following parameter to disable external VIP ports:

—

```
resource_registry:
  # Since the compute stack deploys only compute nodes ExternalVIPPorts are not required.
  OS::TripleO::Network::Ports::ExternalVipPort: /usr/share/openstack-tripleo-heat-
  templates/network/ports/noop.yaml
```

Deploy the overcloud

Choose one of the following:

Multi-cell deployment with a single network

Run the **overcloud deploy** command and add the environment files that you created to configure the new cell stack. For example:


```
$ openstack overcloud deploy \
  --templates /usr/share/openstack-tripleo-heat-templates \
  --stack cell1 \
  -r $HOME/$DIR/cell_roles_data.yaml \
  -e $HOME/$DIR/cell1-ctrl_input.yaml \
  -e $HOME/$DIR/cell1.yaml
```

Multi-cell deployment with segmented networks

Run the **overcloud deploy** command with the additional network data environment file that you created in the previous steps.

The following example shows the **overcloud deploy** command with the environment files that you created to designate a network segment for the cell. Edit the command according to the actual number and names of the cells that you want to deploy.

```
openstack overcloud deploy \
  --templates /usr/share/openstack-tripleo-heat-templates \
  --stack cell1-ctrl \
  -r $HOME/$DIR/cell_roles_data.yaml \
  -n $HOME/$DIR/cell1_routes.yaml \
  -n $HOME/$DIR/network_data-ctrl.yaml \
  -e $HOME/$DIR/cell1-ctrl-input.yaml \
  -e $HOME/$DIR/cell1.yaml
```



NOTE

If you deploy Compute cells in Edge sites, run the **overcloud deploy** command in each site with the environment files and configuration for each Compute cell in that site.

4.5. CREATING AND PROVISIONING A CELL

After you deploy the overcloud with a new cell stack as described in [Section 4.4, "Deploying a multi-cell overcloud"](#), you create and provision the Compute cell.



NOTE

This process must be repeated for each cell that you create and launch. You can automate the steps in an Ansible playbook. For an example of an Ansible playbook, see the [Create the cell and discover Compute nodes](#) section of the OpenStack community documentation. Community documentation is provided as-is and is not officially supported.

1. Get the IP addresses of the control plane and cell controller.

```
$ CTRL_IP=$(openstack server list -f value -c Networks --name overcloud-controller-0 | sed
's/ctlplane=//')
$ CELL_CTRL_IP=$(openstack server list -f value -c Networks --name cellcontroller-0 | sed
's/ctlplane=//')
```

2. Add the cell information to all Controller nodes. This information is used to connect to the cell endpoint from the undercloud.

```
(undercloud) [stack@undercloud ~]$ CELL_INTERNALAPI_INFO=$(ssh heat-
admin@${CELL_CTRL_IP} egrep \
    cellcontrol.*\.internalapi /etc/hosts)
(undercloud) [stack@undercloud ~]$ ansible -i /usr/bin/tripleo-ansible-inventory Controller -b \
    -m lineinfile -a "dest=/etc/hosts line=\"${CELL_INTERNALAPI_INFO}\""
```

3. Get the **transport_url** and **database.connection** endpoint information from the controller cell.

```
(undercloud) [stack@undercloud ~]$ CELL_TRANSPORT_URL=$(ssh heat-
admin@${CELL_CTRL_IP} sudo \
    crudini --get /var/lib/config-data/nova/etc/nova/nova.conf DEFAULT
transport_url)
(undercloud) [stack@undercloud ~]$ CELL_MYSQL_VIP=$(ssh heat-
admin@${CELL_CTRL_IP} sudo \
    crudini --get /var/lib/config-data/nova/etc/nova/nova.conf database
connection \
    | perl -nle'/(\d+\.\d+\.\d+\.\d+)/ && print $1')
```

4. Log in to one of the global Controller nodes to create the cell based on the information that you retrieved in the previous steps. For example:

```
$ export CONTAINERCLI='podman'

$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \
    nova-manage cell_v2 create_cell --name computecell1 \
    --database_connection "{scheme}://{username}:{password}@${CELL_MYSQL_VIP}/nova?
{query}" \
    --transport-url "${CELL_TRANSPORT_URL}"
```

5. Check that the cell is created and appears in the cell list.

```
$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \
    nova-manage cell_v2 list_cells --verbose
```

6. Restart the Compute services on the Controller nodes.

```
$ ansible -i /usr/bin/tripleo-ansible-inventory Controller -b -a \
"systemctl restart tripleo_nova_api tripleo_nova_conductor tripleo_nova_scheduler"
```

7. Check that the cell controller services are provisioned.

```
(overcloud) [stack@undercloud ~]$ nova service-list
```

4.6. ADDING COMPUTE NODES TO A CELL

1. Log into one of the Controller nodes.
2. Get the IP address of the control plane for the cell and run the host discovery command to expose and assign Compute hosts to the cell.

```
$ CTRL=overcloud-controller-0
$ CTRL_IP=$(openstack server list -f value -c Networks --name $CTRL | sed 's/ctlplane=//')

$ export CONTAINERCLI='podman'

$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \
nova-manage cell_v2 discover_hosts --by-service --verbose
```

3. Verify that the Compute hosts were assigned to the cell.

```
$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \
nova-manage cell_v2 list_hosts
```

4.7. CONFIGURING AN AVAILABILITY ZONE

You must assign each cell to an availability zone (AZ) to keep the Compute nodes in that cell during instance creation and migration. The Controller cell must be in a different AZ from the Compute cells.

You can use host aggregates to configure the AZ for the Compute cell. The following example shows the command to create a host aggregate for the cell **cell1**, define the AZ for the host aggregate, and add the hosts within the cell to the AZ:

```
(undercloud)$ source ~/overcloudrc
(overcloud)$ openstack aggregate create cell1 --zone cell1
(overcloud)$ openstack aggregate add host cell1 hostA
(overcloud)$ openstack aggregate add host cell1 hostB
```



NOTE

- You cannot use the **OS::TripleO::Services::NovaAZConfig** parameter to automatically create the AZ during deployment, because the cell is not created at this stage.
- Migrating instances between cells is not supported. To move an instance to a different cell, you must delete it from the old cell and re-create it in the new cell.

For more information on host aggregates and availability zones, see [Creating and managing host aggregates](#).

4.8. DELETING A COMPUTE NODE FROM A CELL

To delete a Compute node from a cell, you must delete all instances from the cell and delete the host names from the Placement database.

1. Delete all instances from the Compute nodes in the cell.



NOTE

Migrating instances between cells is not supported. You must delete the instances and re-create them in another cell.

2. On one of the global Controllers, delete all Compute nodes from the cell.

```
$ CTRL=overcloud-controller-0
$ CTRL_IP=$(openstack server list -f value -c Networks --name $CTRL | sed 's/ctlplane=//')

$ export CONTAINERCLI='podman'

$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \
  nova-manage cell_v2 list_hosts

$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \
  nova-manage cell_v2 delete_host --cell_uuid <uuid> --host <compute>
```

3. Delete the resource providers for the cell from the Placement service, to ensure that the host name is available in case you want to add Compute nodes with the same host name to another cell later. For example:

```
(undercloud) $ source ~/overcloudrc

(overcloud) $ openstack resource provider list
+-----+-----+-----+
| uuid                               | name                               | generation |
+-----+-----+-----+
| 9cd04a8b-5e6c-428e-a643-397c9bebcc16 | computecell1-novacompute-0.site1.test | 11         |
+-----+-----+-----+

(overcloud) $ openstack resource provider delete 9cd04a8b-5e6c-428e-a643-397c9bebcc16
```

4.9. DELETING A CELL

To delete a cell, you must first delete all instances and Compute nodes from the cell, as described in [Section 4.8, "Deleting a Compute node from a cell"](#). Then, you delete the cell itself and the cell stack.

1. On one of the global Controllers, delete the cell.

```
$ CTRL=overcloud-controller-0
$ CTRL_IP=$(openstack server list -f value -c Networks --name $CTRL | sed 's/ctlplane=//')

$ export CONTAINERCLI='podman'

$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \
```

```
nova-manage cell_v2 list_cells
```

```
$ ssh heat-admin@${CTRL_IP} sudo ${CONTAINERCLI} exec -i -u root nova_api \  
nova-manage cell_v2 delete_cell --cell_uuid <uuid>
```

2. Delete the cell stack from the overcloud.

```
$ openstack stack delete <stack name> --wait --yes && openstack overcloud plan delete  
<STACK_NAME>
```

**NOTE**

If you deployed separate cell stacks for a Controller and Compute cell, delete the Compute cell stack first and then the Controller cell stack.

CHAPTER 5. CREATING AND MANAGING HOST AGGREGATES

You can partition a Compute deployment into logical groups for performance or administrative purposes. Red Hat OpenStack Platform (RHOSP) provides the following mechanisms for partitioning logical groups:

Host aggregate

A host aggregate is a grouping of Compute nodes into a logical unit based on attributes such as the hardware or performance characteristics. You can assign a Compute node to one or more host aggregates.

You can map flavors and images to host aggregates by setting metadata on the host aggregate, and then matching flavor extra specs or image metadata properties to the host aggregate metadata. The Compute scheduler can use this metadata to schedule instances when the required filters are enabled. Metadata that you specify in a host aggregate limits the use of that host to any instance that has the same metadata specified in its flavor or image.

You can configure weight multipliers for each host aggregate by setting the **xxx_weight_multiplier** configuration option in the host aggregate metadata.

You can use host aggregates to handle load balancing, enforce physical isolation or redundancy, group servers with common attributes, or separate classes of hardware.

When you create a host aggregate, you can specify a zone name. This name is presented to cloud users as an availability zone that they can select.

Availability zones

An availability zone is the cloud user view of a host aggregate. A cloud user cannot view the Compute nodes in the availability zone, or view the metadata of the availability zone. The cloud user can only see the name of the availability zone.

You can assign each Compute node to only one availability zone. You can configure a default availability zone where instances will be scheduled when the cloud user does not specify a zone. You can direct cloud users to use availability zones that have specific capabilities.

5.1. ENABLING SCHEDULING ON HOST AGGREGATES

To schedule instances on host aggregates that have specific attributes, update the configuration of the Compute scheduler to enable filtering based on the host aggregate metadata.

Procedure

1. Open your Compute environment file.
2. Add the following values to the **NovaSchedulerDefaultFilters** parameter, if they are not already present:
 - **AggregateInstanceExtraSpecsFilter**: Add this value to filter Compute nodes by host aggregate metadata that match flavor extra specs.



NOTE

For this filter to perform as expected, you must scope the flavor extra specs by prefixing the **extra_specs** key with the **aggregate_instance_extra_specs** namespace.

- **AggregateImagePropertiesIsolation:** Add this value to filter Compute nodes by host aggregate metadata that match image metadata properties.

**NOTE**

To filter host aggregate metadata using image metadata properties, the host aggregate metadata key must match a valid image metadata property. For details on valid image metadata properties, see [Appendix A. Image Configuration Parameters](#).

- **AvailabilityZoneFilter:** Add this value to filter by availability zone when launching an instance.

**NOTE**

Instead of using the **AvailabilityZoneFilter** Compute scheduler service filter, you can use the Placement service to process availability zone requests.

3. Save the updates to your Compute environment file.
4. Deploy the overcloud, adding your Compute environment file to the stack along with your other environment files:

```
(undercloud)$ openstack overcloud deploy --templates \
-e [your environment files] \
-e /home/stack/templates/<compute_environment_file>.yaml
```

5.2. CREATING A HOST AGGREGATE

You can create as many host aggregates as you require.

Procedure

1. To create a host aggregate, enter the following command:

```
(overcloud)# openstack aggregate create <aggregate_name>
```

Replace **<aggregate_name>** with the name you want to assign to the host aggregate.

2. Add metadata to the host aggregate:

```
(overcloud)# openstack aggregate set --property <key=value> \
--property <key=value> \
<aggregate_name>
```

- Replace **<key=value>** with the metadata key-value pair. If you are using the **AggregateInstanceExtraSpecsFilter** filter, the key can be any arbitrary string, for example, **ssd=true**. If you are using the **AggregateImagePropertiesIsolation** filter, the key must match a valid image metadata property. For more information on valid image metadata properties, see [Appendix A. Image Configuration Parameters](#).
- Replace **<aggregate_name>** with the name of the host aggregate.

3. Add the Compute nodes to the host aggregate:

```
(overcloud)# openstack aggregate add host <aggregate_name> \
    <host_name>
```

- Replace **<aggregate_name>** with the name of the host aggregate to add the Compute node to.
- Replace **<host_name>** with the name of the Compute node to add to the host aggregate.

4. Create a flavor or image for the host aggregate:

- Create a flavor:

```
(overcloud)$ openstack flavor create \
    --ram <size-mb> \
    --disk <size-gb> \
    --vcpus <no_reserved_vcpus> \
    host-agg-flavor
```

- Create an image:

```
(overcloud)$ openstack image create host-agg-image
```

For information on how to create an image, see [Creating an image](#).

5. Set one or more key-value pairs on the flavor or image that match the key-value pairs on the host aggregate.

- To set the key-value pairs on a flavor, use the scope **aggregate_instance_extra_specs**:

```
(overcloud)# openstack flavor set \
    --property aggregate_instance_extra_specs:ssd=true \
    host-agg-flavor
```

- To set the key-value pairs on an image, use valid image metadata properties as the key:

```
(overcloud)# openstack image set --property os_type=linux \
    host-agg-image
```

5.3. CREATING AN AVAILABILITY ZONE

You can create an availability zone that cloud users can select when they create an instance.

Procedure

1. To create an availability zone, you can create a new availability zone host aggregate, or make an existing host aggregate an availability zone:
 - a. To create a new availability zone host aggregate, enter the following command:

```
(overcloud)# openstack aggregate create --zone <availability_zone> \
    <aggregate_name>
```

- Replace **<availability_zone>** with the name you want to assign to the availability zone.
- Replace **<aggregate_name>** with the name you want to assign to the host aggregate.

b. To make an existing host aggregate an availability zone, enter the following command:

```
(overcloud)# openstack aggregate set --zone <availability_zone> \
    <aggregate_name>
```

- Replace **<availability_zone>** with the name you want to assign to the availability zone.
- Replace **<aggregate_name>** with the name of the host aggregate.

2. Optional: Add metadata to the availability zone:

```
(overcloud)# openstack aggregate set --property <key=value> \
    <aggregate_name>
```

- Replace **<key=value>** with your metadata key-value pair. You can add as many key-value properties as required.
- Replace **<aggregate_name>** with the name of the availability zone host aggregate.

3. Add Compute nodes to the availability zone host aggregate:

```
(overcloud)# openstack aggregate add host <aggregate_name> \
    <host_name>
```

- Replace **<aggregate_name>** with the name of the availability zone host aggregate to add the Compute node to.
- Replace **<host_name>** with the name of the Compute node to add to the availability zone.

5.4. DELETING A HOST AGGREGATE

To delete a host aggregate, you first remove all the Compute nodes from the host aggregate.

Procedure

1. To view a list of all the Compute nodes assigned to the host aggregate, enter the following command:

```
(overcloud)# openstack aggregate show <aggregate_name>
```

2. To remove all assigned Compute nodes from the host aggregate, run the following command for each Compute node:

```
(overcloud)# openstack aggregate remove host <aggregate_name> \
    <host_name>
```

- Replace **<aggregate_name>** with the name of the host aggregate to remove the Compute node from.

- Replace **<host_name>** with the name of the Compute node to remove from the host aggregate.
3. After you remove all the Compute nodes from the host aggregate, enter the following command to delete the host aggregate:

```
(overcloud)# openstack aggregate delete <aggregate_name>
```

5.5. CREATING A TENANT-ISOLATED HOST AGGREGATE

You can create a host aggregate that is available only to specific tenants. Only the tenants that you assign to the host aggregate can launch instances on the host aggregate.



NOTE

Tenant isolation uses the Placement service to filter host aggregates for each tenant. This process supersedes the functionality of the **AggregateMultiTenancyIsolation** filter. You therefore do not need to use the **AggregateMultiTenancyIsolation** filter.

Procedure

1. Open your Compute environment file.
2. To schedule the tenant instances on the tenant-isolated host aggregate, set the **NovaSchedulerLimitTenantsToPlacementAggregate** parameter to **True** in the Compute environment file.
3. Optional: To ensure that only the tenants that you assign to a host aggregate can create instances on your cloud, set the **NovaSchedulerPlacementAggregateRequiredForTenants** parameter to **True**.



NOTE

NovaSchedulerPlacementAggregateRequiredForTenants is **False** by default. When this parameter is **False**, tenants that are not assigned to a host aggregate can create instances on any host aggregate.

4. Save the updates to your Compute environment file.
5. Deploy the overcloud, adding your Compute environment file to the stack along with your other environment files:

```
(undercloud)$ openstack overcloud deploy --templates \
-e [your environment files] \
-e /home/stack/templates/<compute_environment_file>.yaml \
```

6. Create the host aggregate. For more information, see [Creating a host aggregate](#).
7. Retrieve the list of tenant IDs:

```
(overcloud)# openstack project list
```

8. Use the **filter_tenant_id<suffix>** metadata key to assign tenants to the host aggregate:

—

```
(overcloud)# openstack aggregate set \
  --property filter_tenant_id<ID0>=<tenant_id0> \
  --property filter_tenant_id<ID1>=<tenant_id1> \
  ...
  --property filter_tenant_id<IDn>=<tenant_idn> \
  <aggregate_name>
```

- Replace **<ID0>**, **<ID1>**, and all IDs up to **<IDn>** with unique values for each tenant filter that you want to create.
- Replace **<tenant_id0>**, **<tenant_id1>**, and all tenant IDs up to **<tenant_idn>** with the ID of each tenant that you want to assign to the host aggregate.
- Replace **<aggregate_name>** with the name of the tenant-isolated host aggregate.
For example, use the following syntax to assign tenants **78f1**, **9d3t**, and **aa29** to the host aggregate **tenant-isolated-aggregate**:

```
(overcloud) # openstack aggregate set \
  --property filter_tenant_id0=78f1 \
  --property filter_tenant_id1=9d3t \
  --property filter_tenant_id2=aa29 \
  tenant-isolated-aggregate
```

TIP

You can create a host aggregate that is available only to a single specific tenant by omitting the suffix from the **filter_tenant_id** metadata key:

```
(overcloud) # openstack aggregate set \
  --property filter_tenant_id=78f1 \
  single-tenant-isolated-aggregate
```

CHAPTER 6. CONFIGURE OPENSTACK COMPUTE STORAGE

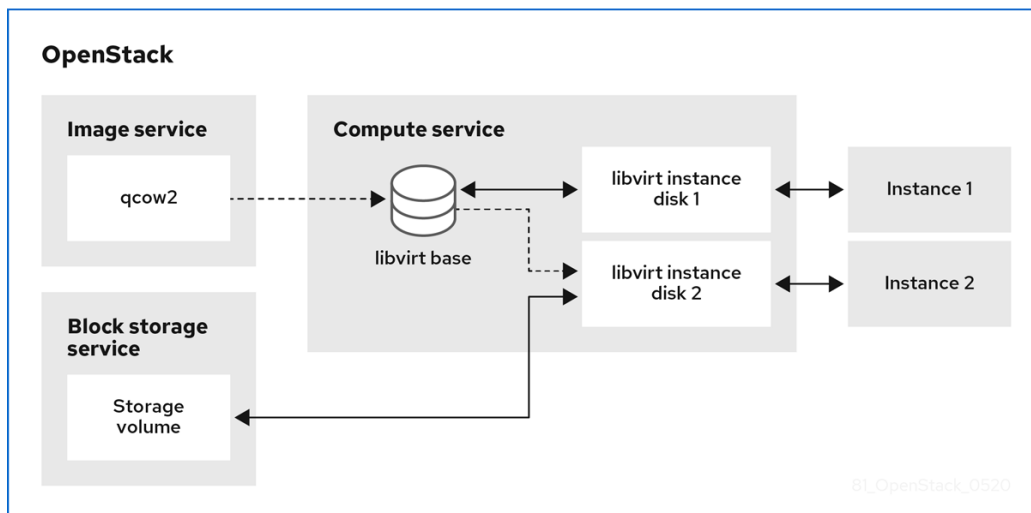
This chapter describes the architecture for the back-end storage of images in OpenStack Compute (nova), and provides basic configuration options.

6.1. ARCHITECTURE OVERVIEW

In Red Hat OpenStack Platform, the OpenStack Compute service uses the KVM hypervisor to execute compute workloads. The **libvirt** driver handles all interactions with KVM, and enables the creation of virtual machines.

Two types of **libvirt** storage must be considered for Compute:

- Base image, which is a cached and formatted copy of the Image service image.
- Instance disk, which is created using the **libvirt** base and is the back end for the virtual machine instance. Instance disk data can be stored either in Compute's ephemeral storage (using the **libvirt** base) or in persistent storage (for example, using Block Storage).



The steps that Compute takes to create a virtual machine instance are:

1. Cache the Image service's backing image as the **libvirt** base.
2. Convert the base image to the raw format (if configured).
3. Resize the base image to match the VM's flavor specifications.
4. Use the base image to create the libvirt instance disk.

In the diagram above, the #1 instance disk uses ephemeral storage; the #2 disk uses a block-storage volume.

Ephemeral storage is an empty, unformatted, additional disk available to an instance. This storage value is defined by the instance flavor. The value provided by the user must be less than or equal to the ephemeral value defined for the flavor. The default value is **0**, meaning no ephemeral storage is created.

The ephemeral disk appears in the same way as a plugged-in hard drive or thumb drive. It is available as a block device which you can check using the **lsblk** command. You can format it, mount it, and use it however you normally would a block device. There is no way to preserve or reference that disk beyond

the instance it is attached to.

Block storage volume is persistent storage available to an instance regardless of the state of the running instance.

6.2. CONFIGURATION

You can configure performance tuning and security for your virtual disks by customizing the Compute (nova) configuration files. Compute is configured in custom environment files and heat templates using the parameters detailed in the [Compute \(nova\) Parameters](#) section in the *Overcloud Parameters* guide. This configuration is generated and stored in the `/var/lib/config-data/puppet-generated/<nova_container>/etc/nova/nova.conf` file, as detailed in the following table.

Table 6.1. Compute Image Parameters

Section	Parameter	Description	Default
[DEFAULT]	force_raw_images	<p>Whether to convert a non-raw cached base image to be raw (boolean). If a non-raw image is converted to raw, Compute:</p> <ul style="list-style-type: none"> Disallows backing files (which might be a security issue). Removes existing compression (to avoid CPU bottlenecks). <p>Converting the base to raw uses more space for any image that could have been used directly by the hypervisor (for example, a qcow2 image). If you have a system with slower I/O or less available space, you might want to specify <i>false</i>, trading the higher CPU requirements of compression for that of minimized input bandwidth.</p> <p>Raw base images are always used with libvirt_images_type=lvm.</p>	true
[DEFAULT]	use_cow_images	<p>Whether to use CoW (Copy on Write) images for libvirt instance disks (boolean):</p> <ul style="list-style-type: none"> false - The raw format is used. Without CoW, more space is used for common parts of the disk image true - The cqw2 format is used. With CoW, depending on the backing store and host caching, there may be better concurrency achieved by having each VM operate on its own copy. 	true

Section	Parameter	Description	Default
[DEFAULT]	preallocate_images	<p>Preallocation mode for libvirt instance disks. Value can be:</p> <ul style="list-style-type: none"> ● none - No storage is provisioned at instance start. ● space - Storage is fully allocated at instance start (using fallocate), which can help with both space guarantees and I/O performance. <p>Even when not using CoW instance disks, the copy each VM gets is sparse and so the VM may fail unexpectedly at run time with ENOSPC. By running fallocate(1) on the instance disk images, Compute immediately and efficiently allocates the space for them in the file system (if supported). Run time performance should also be improved because the file system does not have to dynamically allocate blocks at run time (reducing CPU overhead and more importantly file fragmentation).</p>	none
[DEFAULT]	resize_fs_using_block_device	<p>Whether to enable direct resizing of the base image by accessing the image over a block device (boolean). This is only necessary for images with older versions of cloud-init (that cannot resize themselves).</p> <p>Because this parameter enables the direct mounting of images which might otherwise be disabled for security reasons, it is not enabled by default.</p>	false
[DEFAULT]	default_ephemeral_format	<p>The default format that is used for a new ephemeral volume. Value can be: ext2, ext3, or ext4. The ext4 format provides much faster initialization times than ext3 for new, large disks. You can also override per instance using the guest_format configuration option.</p>	ext4
[DEFAULT]	image_cache_manager_interval	<p>Number of seconds to wait between runs of the image cache manager, which impacts base caching on libvirt compute nodes. This period is used in the auto removal of unused cached images (see remove_unused_base_images and remove_unused_original_minimum_age_seconds).</p>	2400

Section	Parameter	Description	Default
[DEFAULT]	remove_unused_base_images	Whether to enable the automatic removal of unused base images (checked every image_cache_manager_interval seconds). Images are defined as unused if they have not been accessed in remove_unused_original_minimum_age_seconds seconds.	true
[DEFAULT]	remove_unused_original_minimum_age_seconds	How old an unused base image must be before being removed from the libvirt cache (see remove_unused_base_images).	86400
[libvirt]	images_type	Image type to use for libvirt instance disks (deprecates use_cow_images). Value can be: raw , qcow2 , lvm , rbd , or default . If default is specified, the value used for the use_cow_images parameter is used.	default

CHAPTER 7. VIRTUAL MACHINE INSTANCES

OpenStack Compute is the central component that provides virtual machines on demand. Compute interacts with the Identity service for authentication, Image service for images (used to launch instances), and the dashboard service for the user and administrative interface.

Red Hat OpenStack Platform allows you to easily manage virtual machine instances in the cloud. The Compute service creates, schedules, and manages instances, and exposes this functionality to other OpenStack components. This chapter discusses these procedures along with procedures to add components like key pairs, security groups, host aggregates and flavors. The term *instance* is used by OpenStack to mean a virtual machine instance.

7.1. MANAGE INSTANCES

Before you can create an instance, you need to ensure certain other OpenStack components (for example, a network, key pair and an image or a volume as the boot source) are available for the instance.

This section discusses the procedures to add these components, create and manage an instance. Managing an instance refers to updating, and logging in to an instance, viewing how the instances are being used, resizing or deleting them.

7.1.1. Add Components

Use the following sections to create a network, key pair and upload an image or volume source. These components are used in the creation of an instance and are not available by default. You will also need to create a new security group to allow SSH access to the user.

1. In the dashboard, select **Project**.
2. Select **Network > Networks** and ensure there is a private network to which you can attach the new instance (to create a network, see [Create a Network](#) section in the *Networking Guide*).
3. Select **Compute > Access & Security > Key Pairs** and ensure there is a key pair (to create a key pair, see [Section 7.2.1.1, "Create a Key Pair"](#)).
4. Ensure that you have either an image or a volume that can be used as a boot source:
 - To view boot-source images, select the **Images** tab (to create an image, see [Section 1.2.1, "Creating an Image"](#)).
 - To view boot-source volumes, select the **Volumes** tab (to create a volume, see [Create a Volume](#) in the *Storage Guide*).
5. Select **Compute > Access & Security > Security Groups** and ensure you have created a security group rule (to create a security group, see [Project Security Management](#) in the *Users and Identity Management Guide*).

7.1.2. Launch an Instance

Launch one or more instances from the dashboard.

**NOTE**

By default, the Launch Instance form is used to launch instances. However, you can also enable a Launch Instance wizard that simplifies the steps required. For more information, see [Appendix B, *Enabling the Launch Instance Wizard*](#).

1. In the dashboard, select **Project > Compute > Instances**
2. Click **Launch Instance**.
3. Fill out the fields (those marked with '*' are required), and click **Launch**.


One or more instances are created, and launched based on the options provided.

7.1.2.1. Launch Instance Options

The following table outlines the options available when launching a new instance using the Launch Instance form. The same options are also available in the Launch instance wizard.

Table 7.1. Launch Instance Form Options

Tab	Field	Notes
Project and User	Project	Select the project from the dropdown list.
	User	Select the user from the dropdown list.
Details	Availability Zone	Zones are logical groupings of cloud resources in which your instance can be placed. If you are unsure, use the default zone.
	Instance Name	A name to identify your instance.
	Flavor	The flavor determines what resources the instance is given (for example, memory). For default flavor allocations and information on creating new flavors, see Section 7.3, "Manage Flavors" .
	Instance Count	The number of instances to create with these parameters. "1" is preselected.
	Instance Boot Source	Depending on the item selected, new fields are displayed allowing you to select the source: <ul style="list-style-type: none"> ● Image sources must be compatible with OpenStack (see Section 1.2, "Manage images"). ● If a volume or volume source is selected, the source must be formatted using an image (see Basic Volume Usage and Configuration in the <i>Storage Guide</i>).

Tab	Field	Notes
Access and Security	Key Pair	The specified key pair is injected into the instance and is used to remotely access the instance using SSH (if neither a direct login information or a static key pair is provided). Usually one key pair per project is created.
	Security Groups	Security groups contain firewall rules which filter the type and direction of the instance's network traffic (for more information on configuring groups, see Project Security Management in the <i>Users and Identity Management Guide</i>).
Networking	Selected Networks	You must select at least one network. Instances are typically assigned to a private network, and then later given a floating IP address to enable external access.
Post-Creation	Customization Script Source	<p>You can provide either a set of commands or a script file, which will run after the instance is booted (for example, to set the instance host name or a user password). If <i>Direct Input</i> is selected, write your commands in the Script Data field; otherwise, specify your script file.</p> <div>  <p>NOTE</p> <p>Any script that starts with <code>#cloud-config</code> is interpreted as using the cloud-config syntax (for information on the syntax, see http://cloudinit.readthedocs.org/en/latest/topics/examples.html).</p> </div>
Advanced Options	Disk Partition	By default, the instance is built as a single partition and dynamically resized as needed. However, you can choose to manually configure the partitions yourself.
	Configuration Drive	If selected, OpenStack writes metadata to a read-only configuration drive that is attached to the instance when it boots (instead of to Compute's metadata service). After the instance has booted, you can mount this drive to view its contents (enables you to provide files to the instance).

7.1.3. Update an Instance (Actions menu)

You can update an instance by selecting **Project > Compute > Instances** and selecting an action for that instance in the **Actions** column. Actions allow you to manipulate the instance in a number of ways:

Table 7.2. Update Instance Options

Action	Description
--------	-------------

Action	Description
Create Snapshot	Snapshots preserve the disk state of a running instance. You can create a snapshot to migrate the instance, as well as to preserve backup copies.
Associate/Disassociate Floating IP	You must associate an instance with a floating IP (external) address before it can communicate with external networks, or be reached by external users. Because there are a limited number of external addresses in your external subnets, it is recommended that you disassociate any unused addresses.
Edit Instance	Update the instance's name and associated security groups.
Edit Security Groups	Add and remove security groups to or from this instance using the list of available security groups (for more information on configuring groups, see Project Security Management in the <i>Users and Identity Management Guide</i>).
Console	View the instance's console in the browser (allows easy access to the instance).
View Log	View the most recent section of the instance's console log. Once opened, you can view the full log by clicking View Full Log.
Pause/Resume Instance	Immediately pause the instance (you are not asked for confirmation); the state of the instance is stored in memory (RAM).
Suspend/Resume Instance	Immediately suspend the instance (you are not asked for confirmation); like hibernation, the state of the instance is kept on disk.
Resize Instance	Bring up the Resize Instance window (see Section 7.1.4, "Resize an Instance").
Soft Reboot	Gracefully stop and restart the instance. A soft reboot attempts to gracefully shut down all processes before restarting the instance.
Hard Reboot	Stop and restart the instance. A hard reboot effectively just shuts down the instance's power and then turns it back on.
Shut Off Instance	Gracefully stop the instance.


Action	Description
Rebuild Instance	Use new image and disk-partition options to rebuild the image (shut down, re-image, and re-boot the instance). If encountering operating system issues, this option is easier to try than terminating the instance and starting over.
Terminate Instance	Permanently destroy the instance (you are asked for confirmation).

You can create and allocate an external IP address, see [Section 7.2.3, “Create, Assign, and Release Floating IP Addresses”](#)

7.1.4. Resize an Instance

To resize an instance (memory or CPU count), you must select a new flavor for the instance that has the right capacity. If you are increasing the size, remember to first ensure that the host has enough space.

1. Ensure communication between hosts by setting up each host with SSH key authentication so that Compute can use SSH to move disks to other hosts (for example, compute nodes can share the same SSH key).
2. Enable resizing on the original host by setting the **allow_resize_to_same_host** parameter to "True" in your Compute environment file.



NOTE

The **allow_resize_to_same_host** parameter does not resize the instance on the same host. Even if the parameter equals "True" on all Compute nodes, the scheduler does not force the instance to resize on the same host. This is the expected behavior.

3. In the dashboard, select **Project > Compute > Instances**
4. Click the instance's **Actions** arrow, and select **Resize Instance**.
5. Select a new flavor in the **New Flavor** field.
6. If you want to manually partition the instance when it launches (results in a faster build time):
 - a. Select **Advanced Options**.
 - b. In the **Disk Partition** field, select **Manual**.
7. Click **Resize**.

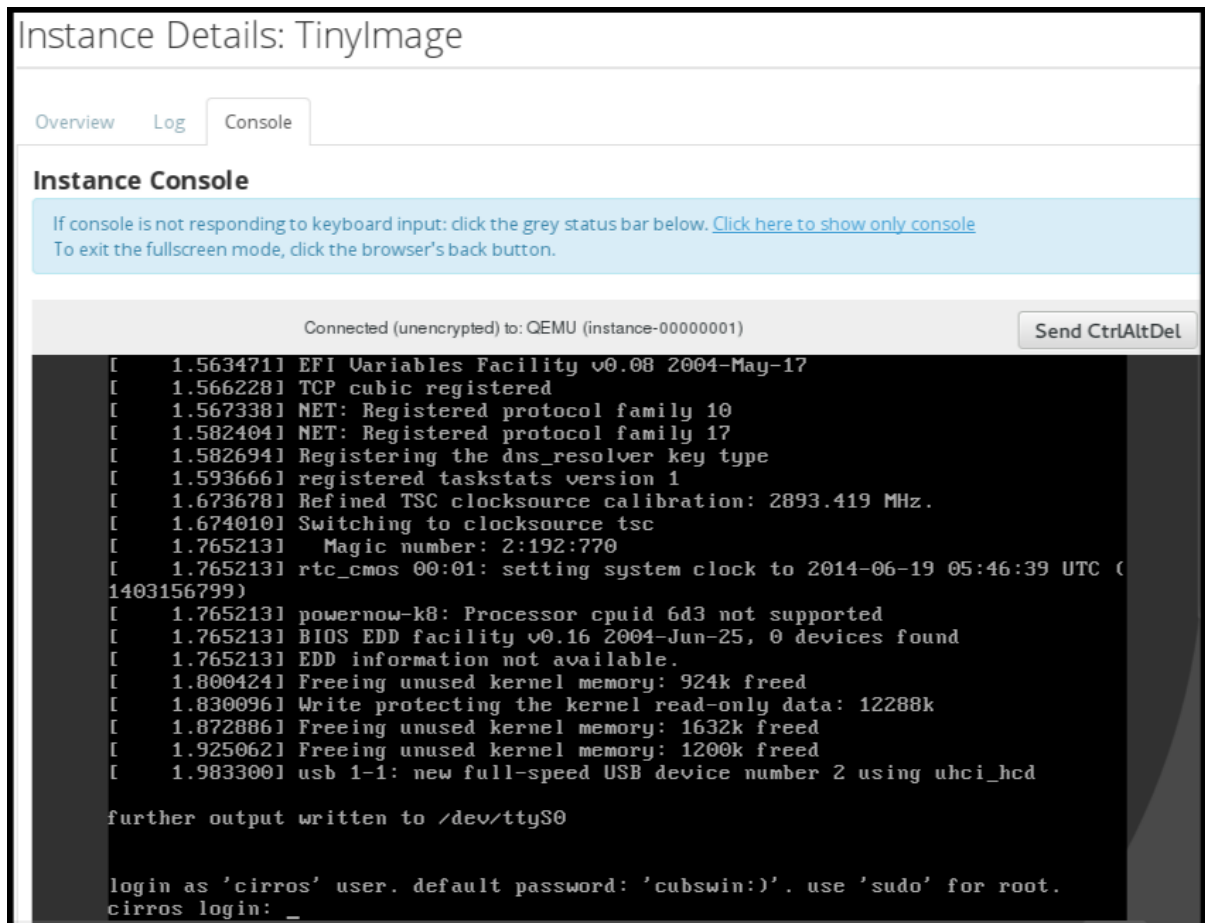
7.1.5. Connect to an Instance

This section discusses the different methods you can use to access an instance console using the dashboard or the command-line interface. You can also directly connect to an instance's serial port allowing you to debug even if the network connection fails.

7.1.5.1. Access an Instance Console using the Dashboard

The console allows you a way to directly access your instance within the dashboard.

1. In the dashboard, select **Compute > Instances**
2. Click the instance's **More** button and select **Console**.



3. Log in using the image's user name and password (for example, a CirrOS image uses *cirros/cubswin:)*).

7.1.5.2. Directly Connect to a VNC Console

You can directly access an instance's VNC console using a URL returned by **nova get-vnc-console** command.

Browser


To obtain a browser URL, use:

```
$ nova get-vnc-console INSTANCE_ID novnc
```

Java Client

To obtain a Java-client URL, use:

```
$ nova get-vnc-console INSTANCE_ID xvpvnc
```

**NOTE**

nova-xvpvncviewer provides a simple example of a Java client. To download the client, use:

```
# git clone https://github.com/cloudbuilders/nova-xvpvncviewer
# cd nova-xvpvncviewer/viewer
# make
```

Run the viewer with the instance's Java-client URL:

```
# java -jar VncViewer.jar URL
```

This tool is provided only for customer convenience, and is not officially supported by Red Hat.

7.1.6. View Instance Usage

The following usage statistics are available:

- Per Project

To view instance usage per project, select **Project > Compute > Overview**. A usage summary is immediately displayed for all project instances.

You can also view statistics for a specific period of time by specifying the date range and clicking **Submit**.

- Per Hypervisor

If logged in as an administrator, you can also view information for all projects. Click **Admin > System** and select one of the tabs. For example, the **Resource Usage** tab offers a way to view reports for a distinct time period. You might also click **Hypervisors** to view your current vCPU, memory, or disk statistics.

**NOTE**

The **vCPU Usage** value (**x of y**) reflects the number of total vCPUs of all virtual machines (x) and the total number of hypervisor cores (y).

7.1.7. Delete an Instance

1. In the dashboard, select **Project > Compute > Instances** and select your instance.
2. Click **Terminate Instance**.

**NOTE**

Deleting an instance does not delete its attached volumes; you must do this separately (see [Delete a Volume](#) in the *Storage Guide*).

7.1.8. Manage Multiple Instances at Once

If you need to start multiple instances at the same time (for example, those that were down for compute or controller maintenance) you can do so easily at **Project > Compute > Instances**

1. Click the check boxes in the first column for the instances that you want to start. If you want to select all of the instances, click the check box in the first row in the table.
2. Click **More Actions** above the table and select **Start Instances**.

Similarly, you can shut off or soft reboot multiple instances by selecting the respective actions.

7.2. MANAGE INSTANCE SECURITY

You can manage access to an instance by assigning it the correct security group (set of firewall rules) and key pair (enables SSH user access). Further, you can assign a floating IP address to an instance to enable external network access. The sections below outline how to create and manage key pairs, security groups, floating IP addresses and logging in to an instance using SSH. There is also a procedure for injecting an **admin** password in to an instance.

For information on managing security groups, see [Project Security Management](#) in the *Users and Identity Management Guide*.

7.2.1. Manage Key Pairs

Key pairs provide SSH access to the instances. Each time a key pair is generated, its certificate is downloaded to the local machine and can be distributed to users. Typically, one key pair is created for each project (and used for multiple instances).

You can also import an existing key pair into OpenStack.

7.2.1.1. Create a Key Pair

1. In the dashboard, select **Project > Compute > Access & Security**
2. On the **Key Pairs** tab, click **Create Key Pair**.
3. Specify a name in the **Key Pair Name** field, and click **Create Key Pair**.

When the key pair is created, a key pair file is automatically downloaded through the browser. Save this file for later connections from external machines. For command-line SSH connections, you can load this file into SSH by executing:

```
# ssh-add ~/.ssh/os-key.pem
```

7.2.1.2. Import a Key Pair

1. In the dashboard, select **Project > Compute > Access & Security**
2. On the **Key Pairs** tab, click **Import Key Pair**.
3. Specify a name in the **Key Pair Name** field, and copy and paste the contents of your public key into the **Public Key** field.
4. Click **Import Key Pair**.

7.2.1.3. Delete a Key Pair

1. In the dashboard, select **Project > Compute > Access & Security**

2. On the **Key Pairs** tab, click the key's **Delete Key Pair** button.

7.2.2. Create a Security Group

Security groups are sets of IP filter rules that can be assigned to project instances, and which define networking access to the instance. Security group are project specific; project members can edit the default rules for their security group and add new rule sets.

1. In the dashboard, select the **Project** tab, and click **Compute > Access & Security**
2. On the **Security Groups** tab, click **+ Create Security Group**.
3. Provide a name and description for the group, and click **Create Security Group**.

For more information on managing project security, see [Project Security Management](#) in the *Users and Identity Management Guide*.

7.2.3. Create, Assign, and Release Floating IP Addresses

By default, an instance is given an internal IP address when it is first created. However, you can enable access through the public network by creating and assigning a floating IP address (external address). You can change an instance's associated IP address regardless of the instance's state.


Projects have a limited range of floating IP address that can be used (by default, the limit is 50), so you should release these addresses for reuse when they are no longer needed. Floating IP addresses can only be allocated from an existing floating IP pool, see [Create Floating IP Pools](#) in the *Networking Guide*.

7.2.3.1. Allocate a Floating IP to the Project

1. In the dashboard, select **Project > Compute > Access & Security**
2. On the **Floating IPs** tab, click **Allocate IP to Project**
3. Select a network from which to allocate the IP address in the **Pool** field.
4. Click **Allocate IP**.

7.2.3.2. Assign a Floating IP

1. In the dashboard, select **Project > Compute > Access & Security**
2. On the **Floating IPs** tab, click the address' **Associate** button.
3. Select the address to be assigned in the IP address field.



NOTE

If no addresses are available, you can click the **+** button to create a new address.

4. Select the instance to be associated in the **Port** to be **Associated** field. An instance can only be associated with one floating IP address.
5. Click **Associate**.

7.2.3.3. Release a Floating IP

1. In the dashboard, select **Project > Compute > Access & Security**
2. On the **Floating IPs** tab, click the address' menu arrow (next to the **Associate/Disassociate** button).
3. Select **Release Floating IP**.

7.2.4. Log in to an Instance

Prerequisites:

- Ensure that the instance's security group has an SSH rule (see [Project Security Management](#) in the *Users and Identity Management Guide*).
- Ensure the instance has a floating IP address (external address) assigned to it (see [Section 7.2.3, "Create, Assign, and Release Floating IP Addresses"](#)).
- Obtain the instance's key-pair certificate. The certificate is downloaded when the key pair is created; if you did not create the key pair yourself, ask your administrator (see [Section 7.2.1, "Manage Key Pairs"](#)).

To first load the key pair file into SSH, and then use ssh without naming it

1. Change the permissions of the generated key-pair certificate.

```
$ chmod 600 os-key.pem
```

2. Check whether **ssh-agent** is already running:

```
# ps -ef | grep ssh-agent
```

3. If not already running, start it up with:

```
# eval `ssh-agent`
```

4. On your local machine, load the key-pair certificate into SSH. For example:

```
$ ssh-add ~/.ssh/os-key.pem
```

5. You can now SSH into the file with the user supplied by the image.

The following example command shows how to SSH into the Red Hat Enterprise Linux guest image with the user **cloud-user**:

```
$ ssh cloud-user@192.0.2.24
```



NOTE

You can also use the certificate directly. For example:

```
$ ssh -i /myDir/os-key.pem cloud-user@192.0.2.24
```


7.2.5. Inject an admin Password Into an Instance

You can inject an **admin (root)** password into an instance using the following procedure.

1. In the **/etc/openstack-dashboard/local_settings** file, set the **change_set_password** parameter value to **True**.

```
can_set_password: True
```

2. Set the **inject_password** parameter to "True" in your Compute environment file.

```
inject_password=true
```

3. Restart the Compute service.

```
# service nova-compute restart
```

When you use the **nova boot** command to launch a new instance, the output of the command displays an **adminPass** parameter. You can use this password to log into the instance as the **root** user.

The Compute service overwrites the password value in the **/etc/shadow** file for the **root** user. This procedure can also be used to activate the **root** account for the KVM guest images. For more information on how to use KVM guest images, see [Section 1.2.1.1, "Use a KVM Guest Image With Red Hat OpenStack Platform"](#)

You can also set a custom password from the dashboard. To enable this, run the following command after you have set **can_set_password** parameter to **true**.

```
# systemctl restart httpd.service
```

The newly added **admin** password fields are as follows:

Launch Instance

Project & User *

Details *

Access & Security

Networking *

Post-Creation

Advanced Options

Key Pair ?

+

Control access to your instance via key pairs, security groups, and other mechanisms.

Admin Password

Confirm Admin Password

Security Groups ?

☒ default

Cancel

Launch

These fields can be used when you launch or rebuild an instance.

7.3. MANAGE FLAVORS

Each created instance is given a flavor (resource template), which determines the instance's size and capacity. Flavors can also specify secondary ephemeral storage, swap disk, metadata to restrict usage, or special project access (none of the default flavors have these additional attributes defined).

Table 7.3. Default Flavors

Name	vCPUs	RAM	Root Disk Size
m1.tiny	1	512 MB	1 GB
m1.small	1	2048 MB	20 GB
m1.medium	2	4096 MB	40 GB
m1.large	4	8192 MB	80 GB
m1.xlarge	8	16384 MB	160 GB

The majority of end users will be able to use the default flavors. However, you can create and manage specialized flavors. For example, you can:

- Change default memory and capacity to suit the underlying hardware needs.
- Add metadata to force a specific I/O rate for the instance or to match a host aggregate.



NOTE

Behavior set using image properties overrides behavior set using flavors (for more information, see [Section 1.2, "Manage images"](#)).

7.3.1. Update Configuration Permissions

By default, only administrators can create flavors or view the complete flavor list (select **Admin > System > Flavors**). To allow all users to configure flavors, specify the following in the **/etc/nova/policy.json** file (nova-api server):

```
"compute_extension:flavormanage": "",
```

7.3.2. Create a Flavor

1. As an admin user in the dashboard, select **Admin > System > Flavors**
2. Click **Create Flavor**, and specify the following fields:

Table 7.4. Flavor Options

Tab	Field	Description
Flavor Information	Name	Unique name.
	ID	Unique ID. The default value, auto , generates a UUID4 value, but you can also manually specify an integer or UUID4 value.
	VCPUs	Number of virtual CPUs.
	RAM (MB)	Memory (in megabytes).
	Root Disk (GB)	Ephemeral disk size (in gigabytes); to use the native image size, specify 0 . This disk is not used if Instance Boot Source=Boot from Volume .
	Ephemeral Disk (GB)	<p>Secondary ephemeral disk size (in gigabytes) available to an instance. This disk is destroyed when an instance is deleted.</p> <p>The default value is 0, which implies that no ephemeral disk is created.</p>

Tab	Field	Description
	Swap Disk (MB)	Swap disk size (in megabytes).
Flavor Access	Selected Projects	Projects which can use the flavor. If no projects are selected, all projects have access (Public=Yes).

3. Click Create Flavor.

7.3.3. Update General Attributes

1. As an admin user in the dashboard, select **Admin > System > Flavors**
2. Click the flavor's **Edit Flavor** button.
3. Update the values, and click **Save**.

7.3.4. Update Flavor Metadata

In addition to editing general attributes, you can add metadata to a flavor (**extra_specs**), which can help fine-tune instance usage. For example, you might want to set the maximum-allowed bandwidth or disk writes.

- Pre-defined keys determine hardware support or quotas. Pre-defined keys are limited by the hypervisor you are using (for libvirt, see [Table 7.5, "Libvirt Metadata"](#)).
- Both pre-defined and user-defined keys can determine instance scheduling. For example, you might specify **SpecialComp=True**; any instance with this flavor can then only run in a host aggregate with the same key-value combination in its metadata.

7.3.4.1. View Metadata

1. As an admin user in the dashboard, select **Admin > System > Flavors**
2. Click the flavor's **Metadata** link (**Yes** or **No**). All current values are listed on the right-hand side under **Existing Metadata**.

7.3.4.2. Add Metadata

You specify a flavor's metadata using a **key/value** pair.

1. As an admin user in the dashboard, select **Admin > System > Flavors**
2. Click the flavor's **Metadata** link (**Yes** or **No**). All current values are listed on the right-hand side under **Existing Metadata**.
3. Under **Available Metadata**, click on the **Other** field, and specify the key you want to add (see [Table 7.5, "Libvirt Metadata"](#)).
4. Click the + button; you can now view the new key under **Existing Metadata**.


5. Fill in the key's value in its right-hand field.

The screenshot shows a web interface titled "Existing Metadata". At the top, there is a "Filter" search bar with a magnifying glass icon. Below this, there are two rows of metadata entries. Each row consists of a key field, a value field, and a minus sign button. The first row shows the key "quota:cpu_shares" with a value of "200". The second row shows the key "quota:disk_read_byte..." with a value of "10240000".

6. When finished with adding key-value pairs, click **Save**.

Table 7.5. Libvirt Metadata

Key	Description
hw:action	<p>Action that configures support limits per instance. Valid actions are:</p> <ul style="list-style-type: none"> • cpu_max_sockets - Maximum supported CPU sockets. • cpu_max_cores - Maximum supported CPU cores. • cpu_max_threads - Maximum supported CPU threads. • cpu_sockets - Preferred number of CPU sockets. • cpu_cores - Preferred number of CPU cores. • cpu_threads - Preferred number of CPU threads. • serial_port_count - Maximum serial ports per instance. <p>Example: hw:cpu_max_sockets=2</p>

Key	Description
hw:NUMA_def	<p>Definition of NUMA topology for the instance. For flavors whose RAM and vCPU allocations are larger than the size of NUMA nodes in the compute hosts, defining NUMA topology enables hosts to better utilize NUMA and improve performance of the guest OS. NUMA definitions defined through the flavor override image definitions. Valid definitions are:</p> <ul style="list-style-type: none"> • numa_nodes - Number of NUMA nodes to expose to the instance. Specify <i>1</i> to ensure image NUMA settings are overridden. • numa_cpus.0 - Mapping of vCPUs N-M to NUMA node 0 (comma-separated list). • numa_cpus.1 - Mapping of vCPUs N-M to NUMA node 1 (comma-separated list). • numa_mem.0 - Mapping N MB of RAM to NUMA node 0. • numa_mem.1 - Mapping N MB of RAM to NUMA node 1. • numa_cpu.N and numa_mem.N are only valid if numa_nodes is set. Additionally, they are only required if the instance's NUMA nodes have an asymmetrical allocation of CPUs and RAM (important for some NFV workloads). <div>  <p>NOTE</p> <p>If the values of numa_cpu or numa_mem.N specify more than that available, an exception is raised.</p> </div> <p>Example when the instance has 8 vCPUs and 4GB RAM:</p> <ul style="list-style-type: none"> • hw:numa_nodes=2 • hw:numa_cpus.0=0,1,2,3,4,5 • hw:numa_cpus.1=6,7 • hw:numa_mem.0=3072 • hw:numa_mem.1=1024 <p>The scheduler looks for a host with 2 NUMA nodes with the ability to run 6 CPUs + 3072 MB, or 3 GB, of RAM on one node, and 2 CPUs + 1024 MB, or 1 GB, of RAM on another node. If a host has a single NUMA node with capability to run 8 CPUs and 4 GB of RAM, it will not be considered a valid match.</p>

Key	Description
hw:watchdog_action	<p>An instance watchdog device can be used to trigger an action if the instance somehow fails (or hangs). Valid actions are:</p> <ul style="list-style-type: none"> ● disabled - The device is not attached (default value). ● pause - Pause the instance. ● poweroff - Forcefully shut down the instance. ● reset - Forcefully reset the instance. ● none - Enable the watchdog, but do nothing if the instance fails. <p>Example: hw:watchdog_action=poweroff</p>
hw:pci_numa_affinity_policy	<p>You can use this parameter to specify the NUMA affinity policy for PCI passthrough devices and SR-IOV interfaces. Set to one of the following valid values:</p> <ul style="list-style-type: none"> ● required: The Compute service only creates an instance that requests a PCI device when at least one of the NUMA nodes of the instance has affinity with the PCI device. This option provides the best performance. ● preferred: The Compute service attempts a best effort selection of PCI devices based on NUMA affinity. If this is not possible, then the Compute service schedules the instance on a NUMA node that has no affinity with the PCI device. ● legacy: (Default) The Compute service creates instances that request a PCI device when either: <ul style="list-style-type: none"> ○ The PCI device has affinity with at least one of the NUMA nodes; or ○ The PCI devices do not provide information on their NUMA affinities. <p>Example: hw:pci_numa_affinity_policy=required</p>
hw_rng:action	<p>A random-number generator device can be added to an instance using its image properties (see hw_rng_model in the "Command-Line Interface Reference" in Red Hat OpenStack Platform documentation).</p> <p>If the device has been added, valid actions are:</p> <ul style="list-style-type: none"> ● allowed - If True, the device is enabled; if False, disabled. By default, the device is disabled. ● rate_bytes - Maximum number of bytes the instance's kernel can read from the host to fill its entropy pool every rate_period (integer). ● rate_period - Duration of the read period in seconds (integer). <p>Example: hw_rng:allowed=True.</p>

Key	Description
hw_video:ram_max_mb	<p>Maximum permitted RAM to be allowed for video devices (in MB).</p> <p>Example: hw:ram_max_mb=64</p>
quota:option	<p>Enforcing limit for the instance. Valid options are:</p> <ul style="list-style-type: none"> • cpu_period - Time period for enforcing cpu_quota (in microseconds). Within the specified cpu_period, each vCPU cannot consume more than cpu_quota of runtime. The value must be in range [1000, 1000000]; 0 means <i>no value</i>. • cpu_quota - Maximum allowed bandwidth (in microseconds) for the vCPU in each `cpu_period`. The value must be in range [1000, 18446744073709551]. 0 means <i>no value</i>; a negative value means that the vCPU is not controlled. cpu_quota and cpu_period can be used to ensure that all vCPUs run at the same speed. • cpu_shares - Share of CPU time for the domain. The value only has meaning when weighted against other machine values in the same domain. That is, an instance with a flavor with 200 will get twice as much machine time as an instance with 100. • disk_read_bytes_sec - Maximum disk reads in bytes per second. • disk_read_iops_sec - Maximum read I/O operations per second. • disk_write_bytes_sec - Maximum disk writes in bytes per second. • disk_write_iops_sec - Maximum write I/O operations per second. • disk_total_bytes_sec - Maximum total throughput limit in bytes per second. • disk_total_iops_sec - Maximum total I/O operations per second. • vif_inbound_average - Desired average of incoming traffic. • vif_inbound_burst - Maximum amount of traffic that can be received at vif_inbound_peak speed. • vif_inbound_peak - Maximum rate at which incoming traffic can be received. • vif_outbound_average - Desired average of outgoing traffic. • vif_outbound_burst - Maximum amount of traffic that can be sent at vif_outbound_peak speed. • vif_outbound_peak - Maximum rate at which outgoing traffic can be sent. <p>Example: quota:vif_inbound_average=10240</p> <p>In addition, the VMware driver supports the following quota options, which control upper and lower limits for CPUs, RAM, disks, and networks, as well as <i>shares</i>, which can be used to control relative allocation of available resources among tenants:</p>

Key	<ul style="list-style-type: none"> • cpu_limit - Maximum CPU frequency available to a virtual machine (in MHz).
	<ul style="list-style-type: none"> • cpu_reservation - Guaranteed minimum amount of CPU resources available to a virtual machine (in MHz). • cpu_shares_level - CPU allocation level (shares) in the case of contention. Possible values are high, normal, low, and custom. • cpu_shares_share - The number of allocated CPU shares. Applicable when cpu_shares_level is set to custom. • memory_limit - Maximum amount of RAM available to a virtual machine (in MB). • memory_reservation - Guaranteed minimum amount of RAM available to a virtual machine (in MB). • memory_shares_level - RAM allocation level (shares) in the case of contention. Possible values are high, normal, low, and custom. • memory_shares_share - The number of allocated RAM shares. Applicable when memory_shares_level is set to custom. • disk_io_limit - Maximum I/O utilization by a virtual machine (in I/O operations per second). • disk_io_reservation - Guaranteed minimum amount of disk resources available to a virtual machine (in I/O operations per second). • disk_io_shares_level - I/O allocation level (shares) in the case of contention. Possible values are high, normal, low, and custom. • disk_io_shares_share - The number of allocated I/O shares. Applicable when disk_io_shares_level is set to custom. • vif_limit - Maximum network bandwidth available to a virtual network adapter (in Mbps). • vif_reservation - Guaranteed minimum network bandwidth available to a virtual network adapter (in Mbps). • vif_shares_level - Network bandwidth allocation level (shares) in the case of contention. Possible values are high, normal, low, and custom. • vif_shares_share - The number of allocated network bandwidth shares. Applicable when vif_shares_level is set to custom.

7.4. SCHEDULE HOSTS

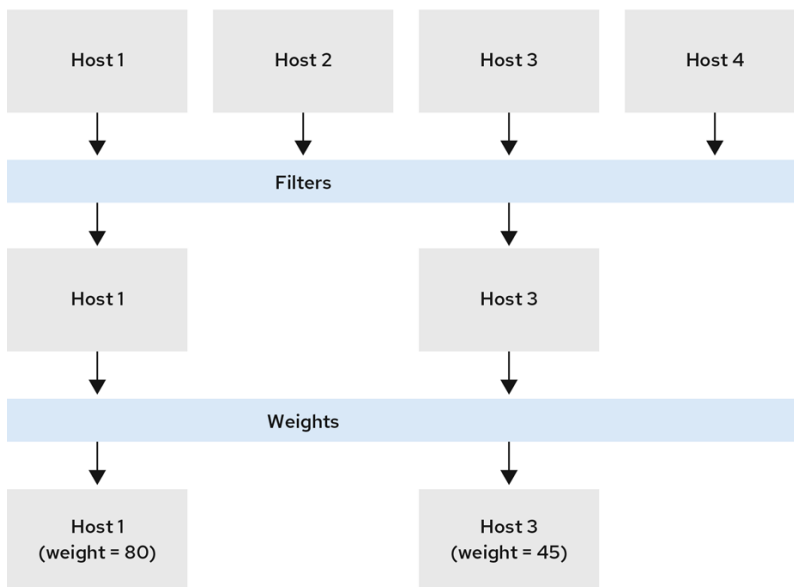
The Compute scheduling service determines on which host, or host aggregate, to place an instance. As an administrator, you can influence where the scheduler places an instance. For example, you might want to limit scheduling to hosts in a certain group or with the right RAM.

You can configure the following components:

- Filters - Determine the initial set of hosts on which an instance might be placed (see [Section 7.4.1, "Configure Scheduling Filters"](#)).

- **Weights** - When filtering is complete, the resulting set of hosts are prioritized using the weighting system. The highest weight has the highest priority (see [Section 7.4.2, "Configure Scheduling Weights"](#)).
- **Scheduler service** - There are a number of configuration options in the `/var/lib/config-data/puppet-generated/<nova_container>/etc/nova/nova.conf` file (on the scheduler host), which determine how the scheduler executes its tasks, and handles weights and filters.
- **Placement service** - Specify the traits an instance requires a host to have, such as the type of storage disk, or the Intel CPU instruction set extension (see [Section 7.4.3, "Configure Placement Service Traits"](#)).

In the following diagram, both host 1 and 3 are eligible after filtering. Host 1 has the highest weight and therefore has the highest priority for scheduling.



81_OpenStack_0520

7.4.1. Configure Scheduling Filters

You define the filters you want the scheduler to use using the **NovaSchedulerDefaultFilters** parameter in your Compute environment file. Filters can be added or removed.

The default configuration runs the following filters in the scheduler:

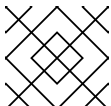
- `RetryFilter`
- `AvailabilityZoneFilter`
- `ComputeFilter`
- `ComputeCapabilitiesFilter`
- `ImagePropertiesFilter`
- `ServerGroupAntiAffinityFilter`
- `ServerGroupAffinityFilter`

Some filters use information in parameters passed to the instance in:

- The **nova boot** command.
- The instance's flavor (see [Section 7.3.4, "Update Flavor Metadata"](#))
- The instance's image (see [Appendix A, Image Configuration Parameters](#)).

All available filters are listed in the following table.

Table 7.6. Scheduling Filters

Filter	Description
AggregateImagePropertyIsolation	Only passes hosts in host aggregates whose metadata matches the instance's image metadata; only valid if a host aggregate is specified for the instance. For more information, see Section 1.2.1, "Creating an Image" .
AggregateInstanceExtraSpecsFilter	Metadata in the host aggregate must match the host's flavor metadata. For more information, see Section 7.3.4, "Update Flavor Metadata" .
	<p>To use this filter in the same NovaSchedulerDefaultFilters parameter as ComputeCapabilitiesFilter, you must scope your flavor extra_specs keys by prefixing them with the correct namespace:</p> <ul style="list-style-type: none"> • ComputeCapabilitiesFilter namespace = "capabilities:" • AggregateInstanceExtraSpecsFilter namespace = "aggregate_instance_extra_specs:"
AggregateMultiTenancyIsolation	<p>A host with the specified filter_tenant_id can only contain instances from that tenant (project).</p> <div>  <div> <p>NOTE</p> <p>The tenant can still place instances on other hosts.</p> </div> </div>
AllHostsFilter	Passes all available hosts (however, does not disable other filters).
AvailabilityZoneFilter	Filters using the instance's specified availability zone.
ComputeCapabilitiesFilter	Ensures Compute metadata is read correctly. Anything before the : is read as a namespace. For example, quota:cpu_period uses quota as the namespace and cpu_period as the key.
ComputeFilter	Passes only hosts that are operational and enabled.
DifferentHostFilter	Enables an instance to build on a host that is different from one or more specified hosts. Specify different hosts using the nova boot option --different_host option.
ImagePropertiesFilter	Only passes hosts that match the instance's image properties. For more information, see Section 1.2.1, "Creating an Image" .

Filter	Description
IsolatedHostsFilter	Passes only isolated hosts running isolated images that are specified using isolated_hosts and isolated_images (comma-separated values).
JsonFilter	<p>Recognises and uses an instance's custom JSON filters:</p> <ul style="list-style-type: none"> Valid operators are: =, <, >, in, ⇐, >=, not, or, and Recognised variables are: \$free_ram_mb, \$free_disk_mb, \$total_usable_ram_mb, \$vcpus_total, \$vcpus_used
	<p>The filter is specified as a query hint in the nova boot command. For example:</p> <p>--hint query='['>=', '\$free_disk_mb', 200 * 1024]'</p>
MetricsFilter	<p>Use this filter to limit scheduling to Compute nodes that report the metrics configured by using metrics/weight_setting.</p> <p>TIP</p> <p>By default, the Compute scheduling service updates the metrics every 60 seconds. To ensure the metrics are up-to-date, you can increase the frequency at which the metrics data is refreshed using the update_resources_interval configuration option. For example, use the following configuration to refresh the metrics data every 2 seconds:</p> <pre>parameter_defaults: ComputeExtraConfig: nova::config::nova_config: DEFAULT/update_resources_interval: value: '2'</pre>
NUMATopologyFilter	Filters out hosts based on its NUMA topology. If the instance has no topology defined, any host can be used. The filter tries to match the exact NUMA topology of the instance to those of the host (it does not attempt to pack the instance onto the host). The filter also looks at the standard over-subscription limits for each NUMA node, and provides limits to the compute host accordingly.
RetryFilter	Filters out hosts that have failed a scheduling attempt; valid if scheduler_max_attempts is greater than zero (defaults to "3").
SameHostFilter	Passes one or more specified hosts; specify hosts for the instance using the --hint same_host option for nova boot .

Filter	Description
ServerGroupAffinityFilter	<p>Only passes hosts for a specific server group:</p> <ul style="list-style-type: none"> • Give the server group the affinity policy (nova server-group-create --policy affinity groupName). • Build the instance with that group (nova boot option --hint group=UUID).
ServerGroupAntiAffinityFilter	<p>Only passes hosts in a server group that do not already host an instance:</p> <ul style="list-style-type: none"> • Give the server group the anti-affinity policy (nova server-group-create --policy anti-affinity groupName). • Build the instance with that group (nova boot option --hint group=UUID).
SimpleCIDRAffinityFilter	<p>Only passes hosts on the specified IP subnet range specified by the instance's cidr and build_new_host_ip hints. Example:</p> <p>--hint build_near_host_ip=192.0.2.0 --hint cidr=/24</p>

7.4.2. Configure Scheduling Weights

Hosts can be weighted for scheduling; the host with the largest weight (after filtering) is selected. All weighers are given a multiplier that is applied after normalising the node's weight. A node's weight is calculated as:

$$w1_multiplier * \text{norm}(w1) + w2_multiplier * \text{norm}(w2) + \dots$$

You can configure weight options in the Compute node configuration file.

Table 7.7. Configuration options for Scheduling service weights

Configuration option	Description
----------------------	-------------

Configuration option	Description
filter_scheduler/weight_classes	<p>Use this parameter to configure which of the following attributes to use for calculating the weight of each host:</p> <ul style="list-style-type: none"> ● nova.scheduler.weights.ram.RAMWeigher - Weighs the available RAM on the Compute node. ● nova.scheduler.weights.cpu.CPUWeigher - Weighs the available CPUs on the Compute node. ● nova.scheduler.weights.disk.DiskWeigher - Weighs the available disks on the Compute node. ● nova.scheduler.weights.metrics.MetricsWeigher - Weighs the metrics of the Compute node. ● nova.scheduler.weights.affinity.ServerGroupSoftAffinityWeigher - Weighs the proximity of the Compute node to other nodes in the given instance group. ● nova.scheduler.weights.affinity.ServerGroupSoftAntiAffinityWeigher - Weighs the proximity of the Compute node to other nodes in the given instance group. ● nova.scheduler.weights.compute.BuildFailureWeigher - Weighs Compute nodes by the number of recent failed boot attempts. ● nova.scheduler.weights.io_ops.IoOpsWeigher - Weighs Compute nodes by their workload. ● nova.scheduler.weights.pci.PCIWeigher - Weighs Compute nodes by their PCI availability. ● nova.scheduler.weights.cross_cell.CrossCellWeigher - Weighs Compute nodes based on which cell they are in, giving preference to Compute nodes in the source cell when moving an instance. ● nova.scheduler.weights.all_weighters - (Default) Uses all the above weighers. <p>Type: String</p>

Configuration option	Description
filter_scheduler/ram_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts based on the available RAM.</p> <p>Set to a positive value to prefer hosts with more available RAM, which spreads instances across many hosts.</p> <p>Set to a negative value to prefer hosts with less available RAM, which fills up (stacks) hosts as much as possible before scheduling to a less-used host.</p> <p>The absolute value, whether positive or negative, controls how strong the RAM weigher is relative to other weighers.</p> <p>By default, the scheduler spreads instances across all hosts evenly (ram_weight_multiplier=1.0).</p> <p>Type: Floating point</p>
filter_scheduler/disk_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts based on the available disk space.</p> <p>Set to a positive value to prefer hosts with more available disk space, which spreads instances across many hosts.</p> <p>Set to a negative value to prefer hosts with less available disk space, which fills up (stacks) hosts as much as possible before scheduling to a less-used host.</p> <p>The absolute value, whether positive or negative, controls how strong the disk weigher is relative to other weighers.</p> <p>By default, the scheduler spreads instances across all hosts evenly (disk_weight_multiplier=1.0).</p> <p>Type: Floating point</p>
filter_scheduler/cpu_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts based on the available vCPUs.</p> <p>Set to a positive value to prefer hosts with more available vCPUs, which spreads instances across many hosts.</p> <p>Set to a negative value to prefer hosts with less available vCPUs, which fills up (stacks) hosts as much as possible before scheduling to a less-used host.</p> <p>The absolute value, whether positive or negative, controls how strong the vCPU weigher is relative to other weighers.</p> <p>By default, the scheduler spreads instances across all hosts evenly (cpu_weight_multiplier=1.0).</p> <p>Type: Floating point</p>

Configuration option	Description
filter_scheduler/io_ops_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts based on the host workload.</p> <p>Set to a negative value to prefer hosts with lighter workloads, which distributes the workload across more hosts.</p> <p>Set to a positive value to prefer hosts with heavier workloads, which schedules instances onto hosts that are already busy.</p> <p>The absolute value, whether positive or negative, controls how strong the I/O operations weigher is relative to other weighers.</p> <p>By default, the scheduler distributes the workload across more hosts (io_ops_weight_multiplier=-1.0).</p> <p>Type: Floating point</p>
filter_scheduler/build_failure_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts based on recent build failures.</p> <p>Set to a positive value to increase the significance of build failures recently reported by the host. Hosts with recent build failures are then less likely to be chosen.</p> <p>Set to 0 to disable weighing compute hosts by the number of recent failures.</p> <p>Default: 1000000.0</p> <p>Type: Floating point</p>
filter_scheduler/cross_cell_move_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts during a cross-cell move. This option determines how much weight is placed on a host which is within the same source cell when moving an instance. By default, the scheduler prefers hosts within the same source cell when migrating an instance.</p> <p>Set to a positive value to prefer hosts within the same cell the instance is currently running. Set to a negative value to prefer hosts located in a different cell from that where the instance is currently running.</p> <p>Default: 1000000.0</p> <p>Type: Floating point</p>

Configuration option	Description
filter_scheduler/pci_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts based on the number of PCI devices on the host and the number of PCI devices requested by an instance. If an instance requests PCI devices, then the more PCI devices a Compute node has the higher the weight allocated to the Compute node.</p> <p>For example, if there are three hosts available, one with a single PCI device, one with multiple PCI devices and one without any PCI devices, then the Compute scheduler prioritizes these hosts based on the demands of the instance. The first host should be preferred if the instance requests one PCI device, the second host if the instance requires multiple PCI devices and the third host if the instance does not request a PCI device.</p> <p>Configure this option to prevent non-PCI instances from occupying resources on hosts with PCI devices.</p> <p>Default: 1.0</p> <p>Type: Positive floating point</p>
filter_scheduler/host_subset_size	<p>Use this parameter to specify the size of the subset of filtered hosts from which to select the host. Must be set to at least 1. A value of 1 selects the first host returned by the weighing functions. Any value less than 1 is ignored and 1 is used instead.</p> <p>Set to a value greater than 1 to prevent multiple scheduler processes handling similar requests selecting the same host, creating a potential race condition. By selecting a host randomly from the N hosts that best fit the request, the chance of a conflict is reduced. However, the higher you set this value, the less optimal the chosen host may be for a given request.</p> <p>Default: 1</p> <p>Type: Integer</p>
filter_scheduler/soft_affinity_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts for group soft-affinity.</p> <p>Default: 1.0</p> <p>Type: Positive floating point</p>

Configuration option	Description
filter_scheduler/soft_anti_affinity_weight_multiplier	<p>Use this parameter to specify the multiplier to use to weigh hosts for group soft-anti-affinity.</p> <p>Default: 1.0</p> <p>Type: Positive floating point</p>
metrics/weight_multiplier	<p>Use this parameter to specify the multiplier to use for weighting metrics. By default, weight_multiplier=1.0, which spreads instances across possible hosts.</p> <p>Set to a number greater than 1.0 to increase the effect of the metric on the overall weight.</p> <p>Set to a number between 0.0 and 1.0 to reduce the effect of the metric on the overall weight.</p> <p>Set to 0.0 to ignore the metric value and return the value of the 'weight_of_unavailable' option.</p> <p>Set to a negative number to prioritize the host with lower metrics, and stack instances in hosts.</p> <p>Default: 1.0</p> <p>Type: Floating point</p>
metrics/weight_setting	<p>Use this parameter to specify the metrics to use for weighting, and the ratio to use to calculate the weight of each metric. Valid metric names:</p> <ul style="list-style-type: none"> ● cpu.frequency - CPU frequency ● cpu.user.time - CPU user mode time ● cpu.kernel.time - CPU kernel time ● cpu.idle.time - CPU idle time ● cpu.iowait.time - CPU I/O wait time ● cpu.user.percent - CPU user mode percentage ● cpu.kernel.percent - CPU kernel percentage ● cpu.idle.percent - CPU idle percentage ● cpu.iowait.percent - CPU I/O wait percentage ● cpu.percent - Generic CPU utilization <p>Example: weight_setting=cpu.user.time=1.0</p> <p>Type: Comma-separated list of metric=ratio pairs.</p>

Configuration option	Description
metrics/required	<p>Use this parameter to specify how to handle configured metrics/weight_setting metrics that are unavailable:</p> <ul style="list-style-type: none"> • True - Metrics are required. If the metric is unavailable, an exception is raised. To avoid the exception, use the MetricsFilter filter in NovaSchedulerDefaultFilters. • False - The unavailable metric is treated as a negative factor in the weighing process. Set the returned value by using the weight_of_unavailable configuration option. <p>Type: Boolean</p>
metrics/weight_of_unavailable	<p>Use this parameter to specify the weight to use if any metrics/weight_setting metric is unavailable, and metrics/required=False.</p> <p>Default: -10000.0</p> <p>Type: Floating point</p>

7.4.3. Configure Placement Service Traits

The placement service tracks the inventory and usage of resource providers, which can be a compute node, a shared storage pool, or an IP allocation pool. Any service that needs to manage the selection and consumption of resources can use the placement service.

To query the placement service, install the **python3-osc-placement** package on the undercloud.

Each resource provider has a set of traits. Traits are the qualitative aspects of a resource provider, for example, the type of storage disk, or the Intel CPU instruction set extension. An instance can specify which of these traits it requires.

The Compute (nova) service interacts with the placement service when it creates instances, with the **nova-compute** and **nova-scheduler** processes.

nova-compute

- Creates the resource provider record.
- Sets the inventory that describes the available quantitative resources, such as the available vCPUs.
- Sets the traits that describe qualitative aspects of the resource provider. The **libvirt** virtualization driver reports these traits to the placement service (see [Section 7.4.3.1, "libvirt virtualization driver capabilities as placement service traits"](#) for details).

nova-scheduler

- Sends a request to the placement service for a list of allocation candidates.

- Decides which destination host to build a server on, based on the traits required by the instance.

7.4.3.1. libvirt virtualization driver capabilities as placement service traits

You can use the capabilities of **libvirt** virtualization drivers as placement service traits. The traits that you can specify are defined in the **os-traits** library, for example:

- **COMPUTE_TRUSTED_CERTS**
- **COMPUTE_NET_ATTACH_INTERFACE_WITH_TAG**
- **COMPUTE_IMAGE_TYPE_RAW**
- **HW_CPU_X86_AVX**
- **HW_CPU_X86_AVX512VL**
- **HW_CPU_X86_AVX512CD**

See the [os-traits library](#) for a catalog of the standardized constants that an instance can request for a particular hardware, virtualization, storage, network, or device trait.

The following **libvirt** virtualization drivers automatically report the features that a host CPU provides, such as the type of instruction set, for example, SSE4, AVX, or AVX-512, to the placement service:

- Libvirt QEMU (x86)
- Libvirt KVM (x86)
- Libvirt KVM (ppc64)

If you are using one of these drivers, you can configure the flavor extra specs or image metadata for an instance to request a resource provider with specific CPU features.

7.4.3.2. Using placement service traits to specify resource provider requirements

You can use one of the following methods to specify the required resource provider traits for an instance:

- [Requesting a trait using image metadata](#)
- [Requesting a trait using flavor extra specs](#)

In the following example procedures, the instance requires a particular type of CPU.

Prerequisites

- The placement service package, **python3-osc-placement**, is installed on the undercloud.
- Your deployment uses one of the following **libvirt** virtualization drivers:
 - Libvirt QEMU (x86)
 - Libvirt KVM (x86)

- Libvirt KVM (ppc64)

Procedure: Requesting a trait using image metadata

1. Create a new image or modify an existing one to set the required trait:

```
$ openstack image create ... $IMAGE
$ openstack image set --property trait:HW_CPU_X86_AVX512BW=required $IMAGE
```

2. Boot an instance using the image:

```
$ openstack server create --image=$IMAGE ... $SERVER_NAME
```

Result: The instance is created on a host that supports AVX-512.

Procedure: Requesting a trait using flavor extra specs

1. Create a new flavor or modify an existing one to set the required trait:

```
$ openstack flavor create ... $FLAVOR
$ openstack flavor set --property trait:HW_CPU_X86_AVX512BW=required $FLAVOR
```

2. Boot an instance using the flavor:

```
$ openstack server create --flavor=$FLAVOR ... $SERVER_NAME
```

Result: The instance is created on a host that supports AVX-512.

7.4.4. Configuring a guaranteed minimum bandwidth QoS

You can create instances that request a guaranteed minimum bandwidth by using a Quality of Service (QoS) policy.

QoS policies with a guaranteed minimum bandwidth rule are assigned to ports on a specific physical network. When you create an instance that uses the configured port, the Compute scheduling service selects a host for the instance that satisfies this request. The Compute scheduling service checks the Placement service for the amount of bandwidth reserved by other instances on each physical interface, before selecting a host to deploy an instance on.

Limitations/Restrictions

- You can only assign a guaranteed minimum bandwidth QoS policy when creating a new instance. You cannot assign a guaranteed minimum bandwidth QoS policy to instances that are already running, as the Compute service only updates resource usage for an instance in placement during creation or move operations, which means the minimum bandwidth available to the instance cannot be guaranteed.
- You cannot live migrate an instance that uses a port that has resource requests, such as a guaranteed minimum bandwidth QoS policy. Run the following command to check if a port has resource requests:

```
$ openstack port show <port_name/port_id>
```

Prerequisites

Field	Value
description	
id	b80acc64-4fc2-41f2-a346-520d7cfe0e2b
is_default	False
location	cloud='', project.domain_id=, project.domain_name='Default', project.id=ba4de51bf7694228a350dd22b7a3dc24, project.name=admin, region_name=regionOne, zone=
name	policy0
project_id	ba4de51bf7694228a350dd22b7a3dc24
rules	[[{min_kbps: 100000, direction: egress, id: d46218fe-9218-4e96-952b-9f45a5cb3b3c, qos_policy_id: b80acc64-4fc2-41f2-a346-520d7cfe0e2b, type: minimum_bandwidth}, {min_kbps: 100000, direction: ingress, id: 1202c4e3-a03a-464c-80d5-0bf90bb74c9d, qos_policy_id: b80acc64-4fc2-41f2-a346-520d7cfe0e2b, type: minimum_bandwidth}]]
shared	True
tags	[]

3. Create a port from the appropriate policy:

```
(overcloud) $ openstack port create port-normal-qos --network net0 --qos-policy policy0
```

4. Create an instance, specifying the NIC port to use:

```
$ openstack server create --flavor cirros256 --image cirros-0.3.5-x86_64-disk --nic port-id=port-normal-qos --wait qos_instance
```

An "ACTIVE" status in the output indicates that you have successfully created the instance on a host that can provide the requested guaranteed minimum bandwidth.

7.4.4.1. Removing a guaranteed minimum bandwidth QoS from an instance

If you want to lift the guaranteed minimum bandwidth QoS policy restriction from an instance, you can detach the interface.

1. To detach the interface, enter the following command:

```
$ openstack server remove port <vm_name|vm_id> <port_name|port_id>
```

7.4.5. Reserve NUMA Nodes with PCI Devices

Compute uses the filter scheduler to prioritize hosts with PCI devices for instances requesting PCI. The hosts are weighted using the **PCIWeigher** option, based on the number of PCI devices available on the host and the number of PCI devices requested by an instance. If an instance requests PCI devices, then the hosts with more PCI devices are allocated a higher weight than the others. If an instance is not requesting PCI devices, then prioritization does not take place.

This feature is especially useful in the following cases:

- As an operator, if you want to reserve nodes with PCI devices (typically expensive and with limited resources) for guest instances that request them.
- As a user launching instances, you want to ensure that PCI devices are available when required.



NOTE

For this value to be considered, one of the following values must be added to the **NovaSchedulerDefaultFilters** parameter in your Compute environment file: **PciPassthroughFilter** or **NUMATopologyFilter**.

The **pci_weight_multiplier** configuration option must be a positive value.

7.4.6. Configure Emulator Threads to run on Dedicated Physical CPU

The Compute scheduler determines the CPU resource utilization and places instances based on the number of virtual CPUs (vCPUs) in the flavor. There are a number of hypervisor operations that are performed on the host, on behalf of the guest instance, for example, with QEMU, there are threads used for the QEMU main event loop, asynchronous I/O operations and so on and these operations need to be accounted and scheduled separately.

The **libvirt** driver implements a generic placement policy for KVM which allows QEMU emulator threads

to float across the same physical CPUs (pCPUs) that the vCPUs are running on. This leads to the emulator threads using time borrowed from the vCPUs operations. When you need a guest to have dedicated vCPU allocation, it is necessary to allocate one or more pCPUs for emulator threads. It is therefore necessary to describe to the scheduler any other CPU usage that might be associated with a guest and account for that during placement.



NOTE

In an NFV deployment, to avoid packet loss, you have to make sure that the vCPUs are never preempted.

Before you enable the emulator threads placement policy on a flavor, check that the following heat parameters are defined as follows:

- **NovaComputeCpuSharedSet:** Set this parameter to a list of CPUs defined to run emulator threads.
- **NovaSchedulerDefaultFilters:** Include **NUMATopologyFilter** in the list of defined filters.



NOTE

You can define or change heat parameter values on an active cluster, and then redeploy for those changes to take effect.

To isolate emulator threads, you must use a flavor configured as follows:

```
# openstack flavor set FLAVOR-NAME \
--property hw:cpu_policy=dedicated \
--property hw:emulator_threads_policy=share
```

7.5. MANAGE INSTANCE SNAPSHOTS

An instance snapshot allows you to create a new image from an instance. This is very convenient for upgrading base images or for taking a published image and customizing it for local use.

The difference between an image that you upload directly to the Image Service and an image that you create by snapshot is that an image created by snapshot has additional properties in the Image Service database. These properties are found in the **image_properties** table and include the following parameters:

Table 7.8. Snapshot Options

Name	Value
image_type	snapshot
instance_uuid	<uuid of instance that was snapshotted>
base_image_ref	<uuid of original image of instance that was snapshotted>
image_location	snapshot

Snapshots allow you to create new instances based on that snapshot, and potentially restore an instance to that state. Moreover, this can be performed while the instance is running.

By default, a snapshot is accessible to the users and projects that were selected while launching an instance that the snapshot is based on.

7.5.1. Create an Instance Snapshot

NOTE

If you intend to use an instance snapshot as a template to create new instances, you must ensure that the disk state is consistent. Before you create a snapshot, set the snapshot image metadata property **os_require_quiesce=yes**. For example,

```
$ glance image-update IMAGE_ID --property os_require_quiesce=yes
```

For this to work, the guest should have the **qemu-guest-agent** package installed, and the image should be created with the metadata property parameter **hw_qemu_guest_agent=yes** set. For example,

```
$ glance image-create --name NAME \
  --disk-format raw \
  --container-format bare \
  --file FILE_NAME \
  --is-public True \
  --property hw_qemu_guest_agent=yes \
  --progress
```

If you unconditionally enable the **hw_qemu_guest_agent=yes** parameter, then you are adding another device to the guest. This consumes a PCI slot, and will limit the number of other devices you can allocate to the guest. It also causes Windows guests to display a warning message about an unknown hardware device.

For these reasons, setting the **hw_qemu_guest_agent=yes** parameter is optional, and the parameter should be used for only those images that require the QEMU guest agent.

1. In the dashboard, select **Project > Compute > Instances**
2. Select the instance from which you want to create a snapshot.
3. In the **Actions** column, click **Create Snapshot**.
4. In the **Create Snapshot** dialog, enter a name for the snapshot and click **Create Snapshot**. The **Images** category now shows the instance snapshot.

To launch an instance from a snapshot, select the snapshot and click **Launch**.

7.5.2. Manage a Snapshot

1. In the dashboard, select **Project > Images**.
2. All snapshots you created, appear under the **Project** option.
3. For every snapshot you create, you can perform the following functions, using the dropdown list:

- a. Use the **Create Volume** option to create a volume and entering the values for volume name, description, image source, volume type, size and availability zone. For more information, see [Create a Volume](#) in the *Storage Guide*.
- b. Use the **Edit Image** option to update the snapshot image by updating the values for name, description, Kernel ID, Ramdisk ID, Architecture, Format, Minimum Disk (GB), Minimum RAM (MB), public or private. For more information, see [Section 1.2.3, “Update an image”](#).
- c. Use the **Delete Image** option to delete the snapshot.

7.5.3. Rebuild an Instance to a State in a Snapshot

In an event that you delete an instance on which a snapshot is based, the snapshot still stores the instance ID. You can check this information using the **nova image-list** command and use the snapshot to restore the instance.

1. In the dashboard, select **Project > Compute > Images**
2. Select the snapshot from which you want to restore the instance.
3. In the **Actions** column, click **Launch Instance**.
4. In the **Launch Instance** dialog, enter a name and the other details for the instance and click **Launch**.

For more information on launching an instance, see [Section 7.1.2, “Launch an Instance”](#).

7.5.4. Consistent Snapshots

Previously, file systems had to be quiesced manually (fsfreeze) before taking a snapshot of active instances for consistent backups.

Compute’s **libvirt** driver automatically requests the *QEMU Guest Agent* to freeze the file systems (and applications if **fsfreeze-hook** is installed) during an image snapshot. Support for quiescing file systems enables scheduled, automatic snapshots at the block device level.

This feature is only valid if the QEMU Guest Agent is installed (**qemu-ga**) and the image metadata enables the agent (**hw_qemu_guest_agent=yes**)



NOTE

Snapshots should not be considered a substitute for an actual system backup.


7.6. USE RESCUE MODE FOR INSTANCES

Compute has a method to reboot a virtual machine in rescue mode. Rescue mode provides a mechanism for access when the virtual machine image renders the instance inaccessible. A rescue virtual machine allows a user to fix their virtual machine by accessing the instance with a new root password. This feature is useful if an instance’s filesystem is corrupted. By default, rescue mode starts an instance from the initial image attaching the current boot disk as a secondary one.

7.6.1. Preparing an Image for a Rescue Mode Instance

Due to the fact that both the boot disk and the disk for rescue mode have same UUID, sometimes the virtual machine can be booted from the boot disk instead of the disk for rescue mode.

To avoid this issue, you should create a new image as rescue image based on the procedure in [Section 1.2.1, "Creating an Image"](#):



NOTE

The **rescue** image is stored in **glance** and configured in the **nova.conf** as a default, or you can select when you do the rescue.

7.6.1.1. Rescue Image if Using **ext4** Filesystem

When the base image uses **ext4** filesystem, you can create a rescue image from it using the following procedure:

1. Change the *UUID* to a random value using the **tune2fs** command:

```
# tune2fs -U random /dev/DEVICE_NODE
```

Here *DEVICE_NODE* is the root device node (for example, **sda**, **vda**, and so on).

2. Verify the details of the filesystem, including the new *UUID*:

```
# tune2fs -l
```

3. Update the **/etc/fstab** to use the new *UUID*. You may need to repeat this for any additional partitions you have, that are mounted in the **fstab** by *UUID*.
4. Update the **/boot/grub2/grub.conf** file and update the *UUID* parameter with the new *UUID* of the root disk.
5. Shut down and use this image as your rescue image. This will cause the rescue image to have a new random *UUID* that will not conflict with the instance that you are rescuing.



NOTE

The XFS filesystem cannot change the *UUID* of the root device on the running virtual machine. Reboot the virtual machine until the virtual machine is launched from the disk for rescue mode.

7.6.2. Adding the Rescue Image to the OpenStack Image Service

When you have completed modifying the *UUID* of your image, use the following commands to add the generated rescue image to the OpenStack Image service:

1. Add the rescue image to the Image service:

```
# glance image-create --name IMAGE_NAME --disk-format qcow2 \
  --container-format bare --is-public True --file IMAGE_PATH
```

Here *IMAGE_NAME* is the name of the image, *IMAGE_PATH* is the location of the image.

2. Use the **image-list** command to obtain the *IMAGE_ID* required for launching an instance in the rescue mode.

```
# glance image-list
```

You can also upload an image using the OpenStack Dashboard, see [Section 1.2.2, “Upload an image”](#).

7.6.3. Launching an Instance in Rescue Mode

1. Since you need to rescue an instance with a specific image, rather than the default one, use the **--image** parameter:

```
# nova rescue --image IMAGE_ID VIRTUAL_MACHINE_ID
```

Here *IMAGE_ID* is the ID of the image you want to use and *VIRTUAL_MACHINE_ID* is ID of a virtual machine that you want to rescue.



NOTE

The **nova rescue** command allows an instance to perform a soft shut down. This allows the guest operating system to perform a controlled shutdown before the instance is powered off. The shut down behavior is configured using **shutdown_timeout** in your Compute configuration file. The value stands for the overall period (in seconds) a guest operation system is allowed to complete the shutdown. The default timeout is 60 seconds.

The timeout value can be overridden on a per image basis by means of **os_shutdown_timeout** that is an image metadata setting allowing different types of operating systems to specify how much time they need to shut down cleanly.

2. Reboot the virtual machine.
3. Confirm the status of the virtual machine is *RESCUE* on the controller node by using **nova list** command or by using dashboard.
4. Log in to the new virtual machine dashboard by using the password for rescue mode.

You can now make the necessary changes to your instance to fix any issues.

7.6.4. Unrescuing an Instance

You can **unrescue** the fixed instance to restart it from the boot disk.

1. Execute the following commands on the controller node.

```
# nova unrescue VIRTUAL_MACHINE_ID
```

Here *VIRTUAL_MACHINE_ID* is ID of a virtual machine that you want to unrescue.

The status of your instance returns to *ACTIVE* once the unrescue operation has completed successfully.

7.7. SET A CONFIGURATION DRIVE FOR INSTANCES

You can use the **config-drive** parameter to present a read-only drive to your instances. This drive can contain selected files that are then accessible to the instance. The configuration drive is attached to the instance at boot, and is presented to the instance as a partition. Configuration drives are useful when combined with *cloud-init* (for server bootstrapping), and when you want to pass large files to your instances.

7.7.1. Configuration Drive Options

Use your Compute environment file to set the following configuration drive parameters:


- **config_drive_format** - sets the format of the drive, and accepts the options **iso9660** and **vfat**. By default, it uses **iso9660**.
- **force_config_drive** - this forces the configuration drive to be presented to all instances. Set to "True".
- **mkisofs_cmd** - specifies the command to use for ISO file creation. This value must not be changed, as only *genisoimage* is supported.

7.7.2. Use a Configuration Drive

An instance attaches its configuration drive at boot time. This is enabled by the **--config-drive** option. For example, this command creates a new instance named *test-instance01* and attaches a drive containing a file named */root/user-data.txt*:

```
# nova boot --flavor m1.tiny --config-drive true --file /root/user-data.txt=/root/user-data.txt --image
cirros test-instance01
```

Once the instance has booted, you can log in to it and see a file named */root/user-data.txt*.



NOTE

You can use the configuration drive as a source for *cloud-init* information. During the initial instance boot, *cloud-init* can automatically mount the configuration drive and run the setup scripts.

CHAPTER 8. CONFIGURING COMPUTE NODES FOR PERFORMANCE

You can configure the scheduling and placement of instances for optimal performance by creating customized flavors to target specialized workloads, including NFV and High Performance Computing (HPC).

Use the following features to tune your instances for optimal performance:

- CPU pinning: Pin virtual CPUs to physical CPUs.
- Emulator threads: Pin emulator threads associated with the instance to physical CPUs.
- Huge pages: Tune instance memory allocation policies both for normal memory (4k pages) and huge pages (2 MB or 1 GB pages).



NOTE

Configuring any of these features creates an implicit NUMA topology on the instance if there is no NUMA topology already present.

8.1. CONFIGURING CPU PINNING ON THE COMPUTE NODE

You can configure instances to run on dedicated host CPUs. Enabling CPU pinning implicitly configures a guest NUMA topology. Each NUMA node of this NUMA topology maps to a separate host NUMA node. For more information about NUMA, see [CPUs and NUMA nodes](#) in the *Network Functions Virtualization Product Guide*.

Configure CPU pinning on your Compute node based on the NUMA topology of your host system. Reserve some CPU cores across all the NUMA nodes for the host processes for efficiency. Assign the remaining CPU cores to managing your instances.

The following example illustrates eight CPU cores spread across two NUMA nodes.

Table 8.1. Example of NUMA Topology

NUMA Node 0		NUMA Node 1	
Core 0	Core 1	Core 2	Core 3
Core 4	Core 5	Core 6	Core 7

You can schedule dedicated (pinned) and shared (unpinned) instances on the same Compute node. The following procedure reserves cores 0 and 4 for host processes, cores 1, 3, 5 and 7 for instances that require CPU pinning, and cores 2 and 6 for floating instances that do not require CPU pinning.



NOTE

If the host supports simultaneous multithreading (SMT), group thread siblings together in either the dedicated or the shared set. Thread siblings share some common hardware which means it is possible for a process running on one thread sibling to impact the performance of the other thread sibling.

For example, the host identifies four CPUs in a dual core CPU with SMT: 0, 1, 2, and 3. Of these four, there are two pairs of thread siblings:

- Thread sibling 1: CPUs 0 and 2
- Thread sibling 2: CPUs 1 and 3

In this scenario, you should not assign CPUs 0 and 1 as dedicated and 2 and 3 as shared. Instead, you should assign 0 and 2 as dedicated and 1 and 3 as shared.

Prerequisite

- You know the NUMA topology of your Compute node. For more information, see [Discovering your NUMA node topology](#) in the *Network Functions Virtualization Planning and Configuration Guide*.

Procedure

1. Reserve physical CPU cores for the dedicated instances by setting the **NovaComputeCpuDedicatedSet** configuration in the Compute environment file for each Compute node:

```
NovaComputeCpuDedicatedSet: 1,3,5,7
```

2. Reserve physical CPU cores for the shared instances by setting the **NovaComputeCpuSharedSet** configuration in the Compute environment file for each Compute node:

```
NovaComputeCpuSharedSet: 2,6
```

3. Set the **NovaReservedHostMemory** option in the same files to the amount of RAM to reserve for host processes. For example, if you want to reserve 512 MB, use:

```
NovaReservedHostMemory: 512
```

4. To ensure that host processes do not run on the CPU cores reserved for instances, set the parameter **IsolCpusList** in each Compute environment file to the CPU cores you have reserved for instances. Specify the value of the **IsolCpusList** parameter using a list, or ranges, of CPU indices separated by a whitespace.

```
IsolCpusList: 1 2 3 5 6 7
```

5. To filter out hosts based on its NUMA topology, add **NUMATopologyFilter** to the **NovaSchedulerDefaultFilters** parameter in each Compute environment file.
6. To apply this configuration, add the environment file(s) to your deployment command and deploy the overcloud:

■

```
(undercloud) $ openstack overcloud deploy --templates \
-e [your environment files]
-e /home/stack/templates/<compute_environment_file>.yaml
```

8.1.1. Upgrading CPU pinning configuration

From Red Hat OpenStack Platform (RHOSP) 16+ it is not necessary to use host aggregates to ensure dedicated (pinned) and shared (unpinned) instance types run on separate hosts. Also, the **[DEFAULT] reserved_host_cpus** config option is no longer necessary and can be unset.

To upgrade your CPU pinning configuration from earlier versions of RHOSP:

- Migrate the value of **NovaVcpuPinSet** to **NovaComputeCpuDedicatedSet** for hosts that were previously used for pinned instances.
- Migrate the value of **NovaVcpuPinSet** to **NovaComputeCpuSharedSet** for hosts that were previously used for unpinned instances.
- If there is no value set for **NovaVcpuPinSet**, then all host cores should be assigned to either **NovaComputeCpuDedicatedSet** or **NovaComputeCpuSharedSet**, depending on the type of instance running there.

Once the upgrade is complete, it is possible to start setting both options on the same host. However, to do this, all the instances should be migrated from the host, as the Compute service cannot start when cores for an unpinned instance are not listed in **NovaComputeCpuSharedSet**, or when cores for a pinned instance are not listed in **NovaComputeCpuDedicatedSet**.

8.1.2. Launching an instance with CPU pinning

You can launch an instance that uses CPU pinning by specifying a flavor with a dedicated CPU policy.

Prerequisites

- Simultaneous multithreading (SMT) is enabled on the host.
- The Compute node is configured to allow CPU pinning. For more information, see [Configuring CPU pinning on the Compute node](#).

Procedure

1. Create a flavor for instances that require CPU pinning:


```
(overcloud) $ openstack flavor create --ram <size-mb> --disk <size-gb> --vcpus
<no_reserved_vcpus> pinned_cpus
```

2. To request pinned CPUs, set the **hw:cpu_policy** property of the flavor to **dedicated**:

```
(overcloud) $ openstack flavor set --property hw:cpu_policy=dedicated pinned_cpus
```

3. To place each vCPU on thread siblings, set the **hw:cpu_thread_policy** property of the flavor to **require**:

```
(overcloud) $ openstack flavor set --property hw:cpu_thread_policy=require pinned_cpus
```

NOTE

- If the host does not have an SMT architecture or enough CPU cores with available thread siblings, scheduling will fail. To prevent this, set **hw:cpu_thread_policy** to **prefer** instead of **require**. The (default) **prefer** policy ensures that thread siblings are used when available.
- If you use **cpu_thread_policy=isolate**, you must have SMT disabled or use a platform that does not support SMT.

4. Create an instance using the new flavor:

```
(overcloud) $ openstack server create --flavor pinned_cpus --image <image>
pinned_cpu_instance
```

5. To verify correct placement of the new instance, run the following command and check for **OS-EXT-SRV-ATTR:hypervisor_hostname** in the output:

```
(overcloud) $ openstack server show pinned_cpu_instance
```

8.1.3. Launching a floating instance

You can launch an instance that is placed on a floating CPU by specifying a flavor with a shared CPU policy.

Prerequisites

- The Compute node is configured to reserve physical CPU cores for the floating instances. For more information, see [Configuring CPU pinning on the Compute node](#).

Procedure

1. Create a flavor for instances that do not require CPU pinning:

```
(overcloud) $ openstack flavor create --ram <size-mb> --disk <size-gb> --vcpus
<no_reserved_vcpus> floating_cpus
```

2. To request floating CPUs, set the **hw:cpu_policy** property of the flavor to **shared**:

```
(overcloud) $ openstack flavor set --property hw:cpu_policy=shared floating_cpus
```

3. Create an instance using the new flavor:

```
(overcloud) $ openstack server create --flavor floating_cpus --image <image>
floating_cpu_instance
```

4. To verify correct placement of the new instance, run the following command and check for **OS-EXT-SRV-ATTR:hypervisor_hostname** in the output:

```
(overcloud) $ openstack server show floating_cpu_instance
```

8.2. CONFIGURING HUGE PAGES ON THE COMPUTE NODE

Configure the Compute node to enable instances to request huge pages.

Procedure

1. Configure the amount of huge page memory to reserve on each NUMA node for processes that are not instances:

```
parameter_defaults:
  NovaReservedHugePages: ["node:0,size:2048,count:64","node:1,size:1GB,count:1"]
```

Where:

Attribute	Description
size	The size of the allocated huge page. Valid values: * 2048 (for 2MB) * 1GB
count	The number of huge pages used by OVS per NUMA node. For example, for 4096 of socket memory used by Open vSwitch, set this to 2.

2. (Optional) To allow instances to allocate 1GB huge pages, configure the CPU feature flags, **cpu_model_extra_flags**, to include "pdpe1gb":

```
parameter_defaults:
  ComputeExtraConfig:
    nova::compute::libvirt::libvirt_cpu_mode: 'custom'
    nova::compute::libvirt::libvirt_cpu_model: 'Haswell-noTSX'
    nova::compute::libvirt::libvirt_cpu_model_extra_flags: 'vmx, pdpe1gb'
```

NOTE

- CPU feature flags do not need to be configured to allow instances to only request 2 MB huge pages.
- You can only allocate 1G huge pages to an instance if the host supports 1G huge page allocation.
- You only need to set **cpu_model_extra_flags** to **pdpe1gb** when **cpu_mode** is set to **host-model** or **custom**.
- If the host supports **pdpe1gb**, and **host-passthrough** is used as the **cpu_mode**, then you do not need to set **pdpe1gb** as a **cpu_model_extra_flags**. The **pdpe1gb** flag is only included in Opteron_G4 and Opteron_G5 CPU models, it is not included in any of the Intel CPU models supported by QEMU.
- To mitigate for CPU hardware issues, such as Microarchitectural Data Sampling (MDS), you might need to configure other CPU flags. For more information, see [RHOS Mitigation for MDS \("Microarchitectural Data Sampling"\) Security Flaws](#).

3. To avoid loss of performance after applying Meltdown protection, configure the CPU feature flags, **cpu_model_extra_flags**, to include "+pcid":

```
parameter_defaults:
  ComputeExtraConfig:
    nova::compute::libvirt::libvirt_cpu_mode: 'custom'
    nova::compute::libvirt::libvirt_cpu_model: 'Haswell-noTSX'
    nova::compute::libvirt::libvirt_cpu_model_extra_flags: 'vmx, pdpe1gb, +pcid'
```

TIP

For more information, see [Reducing the performance impact of Meltdown CVE fixes for OpenStack guests with "PCID" CPU feature flag](#).

4. Add **NUMATopologyFilter** to the **NovaSchedulerDefaultFilters** parameter in each Compute environment file, if not already present.
5. Apply this huge page configuration by adding the environment file(s) to your deployment command and deploying the overcloud:

```
(undercloud) $ openstack overcloud deploy --templates \
-e [your environment files]
-e /home/stack/templates/<compute_environment_file>.yaml
```

8.2.1. Allocating huge pages to instances

Create a flavor with the **hw:mem_page_size** extra specification key to specify that the instance should use huge pages.

Prerequisites

- The Compute node is configured for huge pages. For more information, see [Configuring huge pages on the Compute node](#).

Procedure

1. Create a flavor for instances that require huge pages:

```
$ openstack flavor create --ram <size-mb> --disk <size-gb> --vcpus <no_reserved_vcpus>
huge_pages
```

2. Set the flavor for huge pages:

```
$ openstack flavor set huge_pages --property hw:mem_page_size=1GB
```

Valid values for **hw:mem_page_size**:

- **large** - Selects the largest page size supported on the host, which may be 2 MB or 1 GB on x86_64 systems.
- **small** - (Default) Selects the smallest page size supported on the host. On x86_64 systems this is 4 kB (normal pages).

- **any** - Selects the largest available huge page size, as determined by the libvirt driver.
- `<pagesize>`: (string) Set an explicit page size if the workload has specific requirements. Use an integer value for the page size in KB, or any standard suffix. For example: 4KB, 2MB, 2048, 1GB.

3. Create an instance using the new flavor:

```
$ openstack server create --flavor huge_pages --image <image> huge_pages_instance
```

Validation

The scheduler identifies a host with enough free huge pages of the required size to back the memory of the instance. If the scheduler is unable to find a host and NUMA node with enough pages, then the request will fail with a `NoValidHost` error.

8.3. CONFIGURING COMPUTE NODES TO USE FILE-BACKED MEMORY FOR INSTANCES

You can use file-backed memory to expand your Compute node memory capacity, by allocating files within the libvirt memory backing directory as instance memory. You can configure the amount of host disk that is available for instance memory, and the location on the disk of the instance memory files.

The OpenStack Compute service reports the capacity configured for file-backed memory to the Placement service as the total system memory capacity. This allows the Compute node to run more instances than would normally fit within the system memory.

To use file-backed memory for instances, you must enable file-backed memory on the Compute node.

Limitations

- You cannot live migrate instances between Compute nodes that have file-backed memory enabled and Compute nodes that do not have file-backed memory enabled.
- File-backed memory is not compatible with huge pages. Instances that use huge pages cannot start on a Compute node with file-backed memory enabled. Use host aggregates to ensure that instances that use huge pages are not placed on Compute nodes with file-backed memory enabled.
- File-backed memory is not compatible with memory overcommit.
- You cannot reserve memory for host processes using **NovaReservedHostMemory**. When file-backed memory is in use, reserved memory corresponds to disk space not set aside for file-backed memory. File-backed memory is reported to the Placement service as the total system memory, with RAM used as cache memory.

Prerequisites

- **NovaRAMAllocationRatio** must be set to "1.0" on the node and any host aggregate the node is added to.
- **NovaReservedHostMemory** must be set to "0".

Procedure

1. Open your Compute environment file.
2. Configure the amount of host disk space, in MiB, to make available for instance RAM, by adding the following parameter to your Compute environment file:

```
parameter_defaults:
  NovaLibvirtFileBackedMemory: 102400
```

3. Optional: To configure the directory to store the memory backing files, set the **QemuMemoryBackingDir** parameter in your Compute environment file. If not set, the memory backing directory defaults to **/var/lib/libvirt/qemu/ram/**.



NOTE

You must locate your backing store in a directory at or above the default directory location, **/var/lib/libvirt/qemu/ram/**.

You can also change the host disk for the backing store. For more information, see [Section 8.3.1, “Changing the memory backing directory host disk”](#).

4. Save the updates to your Compute environment file.
5. To apply this configuration, add your Compute environment file to the stack with your other environment files and deploy the overcloud:

```
(undercloud) $ openstack overcloud deploy --templates \
  -e [your environment files] \
  -e /home/stack/templates/<compute_environment_file>.yaml
```

8.3.1. Changing the memory backing directory host disk

You can move the memory backing directory from the default primary disk location to an alternative disk.

Procedure

1. Create a file system on the alternative backing device. For example, enter the following command to create an **ext4** filesystem on **/dev/sdb**:

```
# mkfs.ext4 /dev/sdb
```

2. Mount the backing device. For example, enter the following command to mount **/dev/sdb** on the default libvirt memory backing directory:

```
# mount /dev/sdb /var/lib/libvirt/qemu/ram
```



NOTE

The mount point must match the value of the **QemuMemoryBackingDir** parameter.

CHAPTER 9. CONFIGURING SEV-CAPABLE COMPUTE NODES TO PROVIDE MEMORY ENCRYPTION FOR INSTANCES



IMPORTANT

This feature is available in this release as a *Technology Preview*, and therefore is not fully supported by Red Hat. It should only be used for testing, and should not be deployed in a production environment. For more information about Technology Preview features, see [Scope of Coverage Details](#).

As a cloud administrator, you can provide cloud users the ability to create instances that run on SEV-capable Compute nodes with memory encryption enabled.

To enable your cloud users to create instances that use memory encryption, you must complete the following procedures:

- Configure the Compute nodes that have the SEV-capable hardware.
- Create a SEV-enabled flavor or image for launching instances.

9.1. SECURE ENCRYPTED VIRTUALIZATION (SEV)

Secure Encrypted Virtualization (SEV), provided by AMD, protects the data in DRAM that a running virtual machine instance is using. SEV encrypts the memory of each instance with a unique key.

SEV increases security when you use non-volatile memory technology (NVDIMM), because an NVDIMM chip can be physically removed from a system with the data intact, similar to a hard drive. Without encryption, any stored information such as sensitive data, passwords, or secret keys can be compromised.

For more information, see the [AMD Secure Encrypted Virtualization \(SEV\)](#) documentation.

Limitations of SEV-encrypted instances

- You cannot live migrate, or suspend and resume SEV-encrypted instances.
- You cannot use PCI passthrough on SEV-encrypted instances to directly access devices.
- You cannot use virtio-blk as the boot disk of SEV-encrypted instances.



NOTE

You can use virtio-scsi or SATA as the boot disk, or virtio-blk for non-boot disks.

- The operating system running in an encrypted instance must contain SEV support.
- Machines that support SEV have a limited number of slots in their memory controller for storing encryption keys. Each running instance with encrypted memory consumes one of these slots. Therefore, the number of SEV instances that can run concurrently is limited to the number of slots in the memory controller. For example, on AMD EPYC Zen 1 the limit is 16, and on AMD EPYC Zen 2, the limit is 255.

- Memory-encrypted instances pin pages in RAM. The Compute service cannot swap these pages, therefore you cannot safely overcommit a Compute node that hosts memory-encrypted instances.

9.2. CONFIGURING A SEV-CAPABLE COMPUTE NODE

To enable your cloud users to create instances that use memory encryption, you must configure the Compute nodes that have the SEV-capable hardware.

Prerequisites

- Your deployment must include a Compute node that runs on AMD hardware capable of supporting SEV, such as an AMD EPYC CPU. You can use the following command to determine if your deployment is SEV-capable:


```
$ lscpu | grep sev
```

- Your deployment must include libvirt 4.5 or later, which includes support for SEV.

Procedure

1. Open your Compute environment file.
2. Optional: Add the following configuration to your Compute environment file to specify the maximum number of memory-encrypted instances the SEV-capable Compute node can host concurrently:

```
parameter_defaults:
  ComputeExtraConfig:
    nova::config::nova_config:
      libvirt/num_memory_encrypted_guests:
        value: 15
```



NOTE

If not set, **libvirt/num_memory_encrypted_guests** defaults to **none**, which means the SEV-capable Compute node does not impose a limit on the number of memory-encrypted instances that can be hosted concurrently. Instead, the hardware determines the maximum number of memory-encrypted instances the SEV-capable Compute node can host concurrently, which might cause some memory-encrypted instances to fail to launch.

3. Optional: To specify that all x86_64 images use the q35 machine type by default, add the **NovaHWMachineType** parameter to the Compute environment file, and set it to **x86_64=q35**. This configuration removes the need to set the **hw_machine_type** property to **q35** on every SEV-enabled instance image.
4. To prevent memory overcommit, set the **NovaRAMAllocationRatio** parameter to **1.0** in the Compute environment file.
5. To ensure that the SEV-capable Compute nodes reserve enough memory for host-level services to function, add 16MB for each potential SEV instance (the maximum number of concurrent SEV instances), to your value for **NovaReservedHostMemory** in the Compute environment file.

6. Add the following configuration to your Compute environment file to schedule memory-encrypted instances on a SEV-capable Compute node aggregate:

```
parameter_defaults:
  ControllerExtraConfig:
    nova::config::nova_config:
      scheduler/enable_isolated_aggregate_filtering:
        value: 'True'
```

7. Save the updates to your Compute environment file.
8. To apply this configuration, add your Compute environment file to the stack with your other environment files and deploy the overcloud:

```
(undercloud) $ openstack overcloud deploy --templates \
-e [your environment files]
-e /home/stack/templates/<compute_environment_file>.yaml
```

9. Create a host aggregate for SEV Compute nodes to ensure that instances that do not request memory encryption are not created on SEV-capable hosts:

```
(undercloud) $ source ~/overcloudrc
(overcloud) $ openstack aggregate create sev_agg
(overcloud) $ openstack aggregate add host sev_agg hostA
(overcloud) $ openstack aggregate add host sev_agg hostB
(overcloud) $ openstack --os-compute-api-version 2.53 aggregate set --property
trait:HW_CPU_X86_AMD_SEV=required sev_agg
```

9.3. CREATING THE IMAGE AND FLAVOR FOR MEMORY ENCRYPTION

To enable your cloud users to create instances that use memory encryption, you can define a SEV-enabled flavor, and you can create a SEV-enabled image.

9.3.1. Creating a SEV-enabled image for instances

When the overcloud contains SEV-capable Compute nodes, you can create a SEV-enabled instance image that your cloud users can use to launch instances that have memory encryption.

Procedure

1. Create a new image for SEV:

```
(overcloud) $ openstack image create ... --property hw_firmware_type=uefi sev-image
```



NOTE

If you use an existing image, the image must have the **hw_firmware_type** property set to **uefi**.

2. Optional: Add the property **hw_mem_encryption=True** to the image to enable SEV memory encryption on the image:


```
(overcloud) $ openstack image set --property hw_mem_encryption=True sev-image
```

TIP

You can enable SEV memory encryption on the flavor. For more information, see [Creating a SEV-enabled flavor for instances](#).

- Optional: Set the machine type to **q35**, if not already set in the Compute node configuration:

```
(overcloud) $ openstack image set --property hw_machine_type=q35 sev-image
```

- Optional: To schedule memory-encrypted instances on a SEV-capable host aggregate, add the following trait to the image extra specs:

```
(overcloud) $ openstack image set --property trait:HW_CPU_X86_AMD_SEV=required sev-image
```

TIP

You can also specify this trait on the flavor. For more information, see [Creating a SEV-enabled flavor for instances](#).

9.3.2. Creating a SEV-enabled flavor for instances

When the overcloud contains SEV-capable Compute nodes, you can create one or more SEV-enabled flavors that your cloud users can use to launch instances that have memory encryption.

Procedure

- Create a flavor for SEV:

```
(overcloud) $ openstack flavor create --vcpus 1 --ram 512 --disk 2 --property hw:mem_encryption=True m1.small-sev
```

- To schedule memory-encrypted instances on a SEV-capable host aggregate, add the following trait to the flavor extra specs:

```
(overcloud) $ openstack flavor set --property trait:HW_CPU_X86_AMD_SEV=required m1.small-sev
```

CHAPTER 10. MIGRATING VIRTUAL MACHINE INSTANCES BETWEEN COMPUTE NODES

You sometimes need to migrate instances from one Compute node to another Compute node in the overcloud, to perform maintenance, rebalance the workload, or replace a failed or failing node.

Compute node maintenance

If you need to temporarily take a Compute node out of service, for instance, to perform hardware maintenance or repair, kernel upgrades and software updates, you can migrate instances running on the Compute node to another Compute node.

Failing Compute node

If a Compute node is about to fail and you need to service it or replace it, you can migrate instances from the failing Compute node to a healthy Compute node.

Failed Compute nodes

If a Compute node has already failed, you can evacuate the instances. You can rebuild instances from the original image on another Compute node, using the same name, UUID, network addresses, and any other allocated resources the instance had before the Compute node failed.

Workload rebalancing

You can migrate one or more instances to another Compute node to rebalance the workload. For example, you can consolidate instances on a Compute node to conserve power, migrate instances to a Compute node that is physically closer to other networked resources to reduce latency, or distribute instances across Compute nodes to avoid hot spots and increase resiliency.

Director configures all Compute nodes to provide secure migration. All Compute nodes also require a shared SSH key to provide the users of each host with access to other Compute nodes during the migration process. Director creates this key using the **OS::TripleO::Services::NovaCompute** composable service. This composable service is one of the main services included on all Compute roles by default. For more information, see [Composable Services and Custom Roles](#) in the *Advanced Overcloud Customization* guide.



NOTE

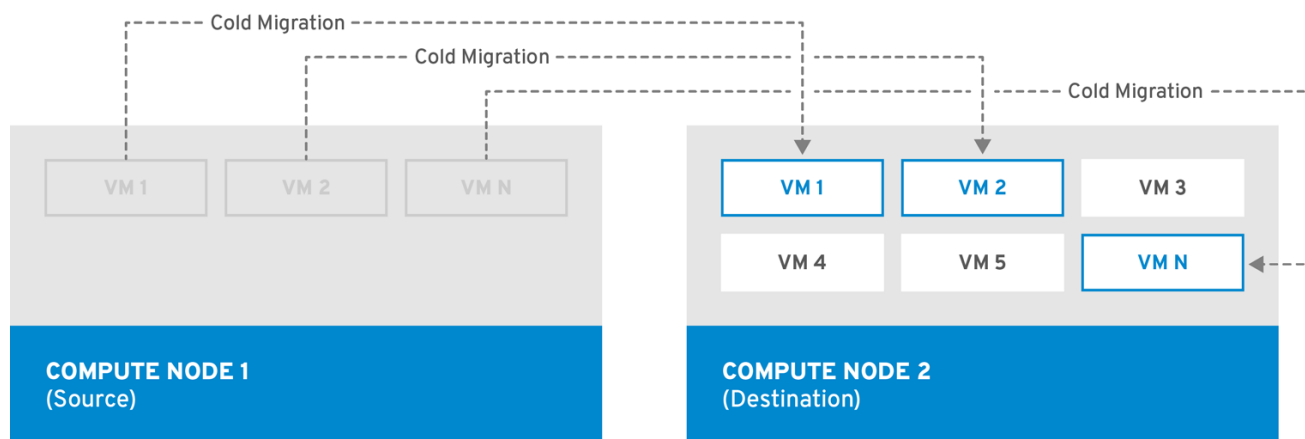
If you have a functioning Compute node, and you want to make a copy of an instance for backup purposes, or to copy the instance to a different environment, follow the procedure in [Importing virtual machines into the overcloud](#) in the *Director Installation and Usage* guide.

10.1. MIGRATION TYPES

Red Hat OpenStack Platform (RHOSP) supports the following types of migration.

Cold migration

Cold migration, or non-live migration, involves shutting down a running instance before migrating it from the source Compute node to the destination Compute node.



OPENSTACK_11_0419

Cold migration involves some downtime for the instance. The migrated instance maintains access to the same volumes and IP addresses.




NOTE

Cold migration requires that both the source and destination Compute nodes are running.

Live migration

Live migration involves moving the instance from the source Compute node to the destination Compute node without shutting it down, and while maintaining state consistency.



OPENSTACK_11_0419

Live migrating an instance involves little or no perceptible downtime. However, live migration does impact performance for the duration of the migration operation. Therefore, instances should be taken out of the critical path while being migrated.



NOTE

Live migration requires that both the source and destination Compute nodes are running.

In some cases, instances cannot use live migration. For more information, see [Migration Constraints](#).

Evacuation

If you need to migrate instances because the source Compute node has already failed, you can evacuate the instances.

10.2. MIGRATION CONSTRAINTS

Migration constraints typically arise with block migration, configuration disks, or when one or more instances access physical hardware on the Compute node.

CPU constraints

The source and destination Compute nodes must have the same CPU architecture. For example, Red Hat does not support migrating an instance from an **x86_64** CPU to a **ppc64le** CPU. In some cases, the CPU of the source and destination Compute node must match exactly, such as instances that use CPU host passthrough. In all cases, the CPU features of the destination node must be a superset of the CPU features on the source node.

Memory constraints

The destination Compute node must have sufficient available RAM. Memory oversubscription can cause migration to fail.

Block migration constraints

Migrating instances that use disks that are stored locally on a Compute node takes significantly longer than migrating volume-backed instances that use shared storage, such as Red Hat Ceph Storage. This latency arises because OpenStack Compute (nova) migrates local disks block-by-block between the Compute nodes over the control plane network by default. By contrast, volume-backed instances that use shared storage, such as Red Hat Ceph Storage, do not have to migrate the volumes, because each Compute node already has access to the shared storage.



NOTE

Network congestion in the control plane network caused by migrating local disks or instances that consume large amounts of RAM might impact the performance of other systems that use the control plane network, such as RabbitMQ.

Read-only drive migration constraints

Migrating a drive is supported only if the drive has both read and write capabilities. For example, OpenStack Compute (nova) cannot migrate a CD-ROM drive or a read-only config drive. However, OpenStack Compute (nova) can migrate a drive with both read and write capabilities, including a config drive with a drive format such as **vfat**.

Live migration constraints

In some cases, live migrating instances involves additional constraints.

No new operations during migration

To achieve state consistency between the copies of the instance on the source and destination nodes, RHOSP must prevent new operations during live migration. Otherwise, live migration might take a long time or potentially never end if writes to memory occur faster than live migration can replicate the state of the memory.

CPU pinning with NUMA

NovaSchedulerDefaultFilters parameter in the Compute configuration must include the values **AggregateInstanceExtraSpecsFilter** and **NUMATopologyFilter**.

Multi-cell clouds

In a multi-cell cloud, instances can be live migrated to a different host in the same cell, but not across cells.

Floating instances

When live migrating floating instances, if the configuration of **NovaComputeCpuSharedSet** on the destination Compute node is different from the configuration of **NovaComputeCpuSharedSet** on the source Compute node, the instances will not be allocated to the CPUs configured for shared (unpinned) instances on the destination Compute node. Therefore, if you need to live migrate floating instances, you must configure all the Compute nodes with the same CPU mappings for dedicated (pinned) and shared (unpinned) instances, or use a host aggregate for the shared instances.

Destination Compute node capacity

The destination Compute node must have sufficient capacity to host the instance that you want to migrate.

SR-IOV live migration

Instances with SR-IOV-based network interfaces can be live migrated. Live migrating instances with direct mode SR-IOV network interfaces attached incurs network downtime while the direct mode interfaces are being detached and re-attached.

Constraints that preclude live migration

You cannot live migrate an instance that uses the following features.

PCI passthrough

QEMU/KVM hypervisors support attaching PCI devices on the Compute node to an instance. Use PCI passthrough to give an instance exclusive access to PCI devices, which appear and behave as if they are physically attached to the operating system of the instance. However, because PCI passthrough involves physical addresses, OpenStack Compute does not support live migration of instances using PCI passthrough.

Port resource requests

You cannot live migrate an instance that uses a port that has resource requests, such as a guaranteed minimum bandwidth QoS policy. Use the following command to check if a port has resource requests:

```
$ openstack port show <port_name/port_id>
```

10.3. PREPARING TO MIGRATE

Before you migrate one or more instances, you need to determine the Compute node names and the IDs of the instances to migrate.

Procedure

1. Identify the source Compute node host name and the destination Compute node host name:

```
(undercloud) $ source ~/overcloudrc
(overcloud) $ openstack compute service list
```

2. List the instances on the source Compute node and locate the ID of the instance or instances that you want to migrate:

■

```
(overcloud) $ openstack server list --host <source> --all-projects
```

Replace **<source>** with the name or ID of the source Compute node.

- Optional: If you are migrating instances from a source Compute node to perform maintenance on the node, you must disable the node to prevent the scheduler from assigning new instances to the node during maintenance:

```
(overcloud) $ source ~/stackrc
(undercloud) $ openstack compute service set <source> nova-compute --disable
```

Replace **<source>** with the name or ID of the source Compute node.

You are now ready to perform the migration. Follow the required procedure detailed in [Cold migrating an instance](#) or [Live migrating an instance](#).

10.4. COLD MIGRATING AN INSTANCE

Cold migrating an instance involves stopping the instance and moving it to another Compute node. Cold migration facilitates migration scenarios that live migrating cannot facilitate, such as migrating instances that use PCI passthrough. The scheduler automatically selects the destination Compute node. For more information, see [Migration Constraints](#).

Procedure

- To cold migrate an instance, enter the following command to power off and move the instance:

```
(overcloud) $ openstack server migrate <vm> --wait
```

- Replace **<vm>** with the name or ID of the instance to migrate.
 - Specify the **--block-migration** flag if migrating a locally stored volume.
- Wait for migration to complete. While you wait for the instance migration to complete, you can check the migration status. For more information, see [Checking migration status](#).
 - Check the status of the instance:

```
(overcloud) $ openstack server list --all-projects
```

A status of "VERIFY_RESIZE" indicates you need to confirm or revert the migration:

- If the migration worked as expected, confirm it:

```
(overcloud) $ openstack server resize --confirm <vm>
```

Replace **<vm>** with the name or ID of the instance to migrate. A status of "ACTIVE" indicates that the instance is ready to use.

- If the migration did not work as expected, revert it:

```
(overcloud) $ openstack server resize --revert <vm>
```

Replace **<vm>** with the name or ID of the instance.

- Restart the instance:

```
(overcloud) $ openstack server start <vm>
```

Replace **<vm>** with the name or ID of the instance.

- Optional: If you disabled the source Compute node for maintenance, you must re-enable the node so that new instances can be assigned to it:

```
(overcloud) $ source ~/stackrc
(undercloud) $ openstack compute service set <source> nova-compute --enable
```

Replace **<source>** with the host name of the source Compute node.

10.5. LIVE MIGRATING AN INSTANCE

Live migration moves an instance from a source Compute node to a destination Compute node with a minimal amount of downtime. Live migration might not be appropriate for all instances. For more information, see [Migration Constraints](#).

Procedure

- To live migrate an instance, specify the instance and the destination Compute node:

```
(overcloud) $ openstack server migrate <vm> --live-migration [--host <dest>] --wait
```

- Replace **<vm>** with the name or ID of the instance.
- Replace **<dest>** with the name or ID of the destination Compute node.



NOTE

The **openstack server migrate** command covers migrating instances with shared storage, which is the default. Specify the **--block-migration** flag to migrate a locally stored volume:

```
(overcloud) $ openstack server migrate <vm> --live-migration [--host <dest>] --wait --block-migration
```

- Confirm that the instance is migrating:

```
(overcloud) $ openstack server show <vm>
```

```
+-----+-----+
| Field          | Value                |
+-----+-----+
| ...            | ...                  |
| status         | MIGRATING            |
| ...            | ...                  |
+-----+-----+
```

- Wait for migration to complete. While you wait for the instance migration to complete, you can check the migration status. For more information, see [Checking migration status](#).

4. Check the status of the instance to confirm if the migration was successful:

```
(overcloud) $ openstack server list --host <dest> --all-projects
```

Replace **<dest>** with the name or ID of the destination Compute node.

5. Optional: If you disabled the source Compute node for maintenance, you must re-enable the node so that new instances can be assigned to it:

```
(overcloud) $ source ~/stackrc
(undercloud) $ openstack compute service set <source> nova-compute --enable
```

Replace **<source>** with the host name of the source Compute node.

10.6. CHECKING MIGRATION STATUS

Migration involves several state transitions before migration is complete. During a healthy migration, the migration state typically transitions as follows:

1. **Queued:** The Compute service has accepted the request to migrate an instance, and migration is pending.
2. **Preparing:** The Compute service is preparing to migrate the instance.
3. **Running:** The Compute service is migrating the instance.
4. **Post-migrating:** The Compute service has built the instance on the destination Compute node and is releasing resources on the source Compute node.
5. **Completed:** The Compute service has completed migrating the instance and finished releasing resources on the source Compute node.

Procedure

1. Retrieve the list of migration IDs for the instance:

```
$ nova server-migration-list <vm>

+----+-----+-----+ (...)
| Id | Source Node | Dest Node | (...)
+----+-----+-----+ (...)
| 2 | -          | -          | (...)
+----+-----+-----+ (...)
```

Replace **<vm>** with the name or ID of the instance.

2. Show the status of the migration:

```
$ <vm> <migration-id>
```

- Replace **<vm>** with the name or ID of the instance.
- Replace **<migration-id>** with the ID of the migration.
Running the **nova server-migration-show** command returns the following example output:

```
■
```


Property	Value
created_at	2017-03-08T02:53:06.000000
dest_compute	controller
dest_host	-
dest_node	-
disk_processed_bytes	0
disk_remaining_bytes	0
disk_total_bytes	0
id	2
memory_processed_bytes	65502513
memory_remaining_bytes	786427904
memory_total_bytes	1091379200
server_uuid	d1df1b5a-70c4-4fed-98b7-423362f2c47c
source_compute	compute2
source_node	-
status	running
updated_at	2017-03-08T02:53:47.000000

TIP

The OpenStack Compute service measures progress of the migration by the number of remaining memory bytes to copy. If this number does not decrease over time, the migration might be unable to complete, and the Compute service might abort it.


Sometimes instance migration can take a long time or encounter errors. For more information, see [Troubleshooting migration](#).

10.7. EVACUATING AN INSTANCE

If you want to move an instance from a dead or shut-down Compute node to a new host in the same environment, you can evacuate it.

The evacuate process destroys the original instance and rebuilds it on another Compute node using the original image, instance name, UUID, network addresses, and any other resources the original instance had allocated to it.

If the instance uses shared storage, the instance root disk is not rebuilt during the evacuate process, as the disk remains accessible by the destination Compute node. If the instance does not use shared storage, then the instance root disk is also rebuilt on the destination Compute node.

**NOTE**

- You can only perform an evacuation when the Compute node is fenced, and the API reports that the state of the Compute node is "down" or "forced-down". If the Compute node is not reported as "down" or "forced-down", the **evacuate** command fails.
- To perform an evacuation, you must be a cloud administrator.

10.7.1. Evacuating one instance

You can evacuate instances one at a time.

Procedure

1. Log onto the failed Compute node as an administrator.
2. Disable the Compute node:

```
(overcloud) [stack@director ~]$ openstack compute service set \
<host> <service> --disable
```

- Replace **<host>** with the name of the Compute node to evacuate the instance from.
- Replace **<service>** with the name of the service to disable, for example **nova-compute**.

3. To evacuate an instance, enter the following command:

```
(overcloud) [stack@director ~]$ nova evacuate [--password <pass>] <vm> [<dest>]
```

- Replace **<pass>** with the admin password to set for the evacuated instance. If a password is not specified, a random password is generated and output when the evacuation is complete.
- Replace **<vm>** with the name or ID of the instance to evacuate.
- Replace **<dest>** with the name of the Compute node to evacuate the instance to. If you do not specify the destination Compute node, the Compute scheduler selects one for you. You can find possible Compute nodes by using the following command:

```
(overcloud) [stack@director ~]$ openstack hypervisor list
```

10.7.2. Evacuating all instances on a host

You can evacuate all instances on a specified Compute node.

Procedure

1. Log onto the failed Compute node as an administrator.
2. Disable the Compute node:

```
(overcloud) [stack@director ~]$ openstack compute service set \
<host> <service> --disable
```

- Replace **<host>** with the name of the Compute node to evacuate the instances from.
- Replace **<service>** with the name of the service to disable, for example **nova-compute**.

3. Evacuate all instances on a specified Compute node:

```
(overcloud) [stack@director ~]$ nova host-evacuate [--target_host <dest>] [--force] <host>
```

- Replace **<dest>** with the name of the destination Compute node to evacuate the instances to. If you do not specify the destination, the Compute scheduler selects one for you. You can find possible Compute nodes by using the following command:

```
(overcloud) [stack@director ~]$ openstack hypervisor list
```

- Replace **<host>** with the name of the Compute node to evacuate the instances from.

10.7.3. Configuring shared storage

If you are using shared storage, export the instance directory for the Compute service to the two nodes, and ensure that the nodes have access. The directory path is set in the **state_path** and **instances_path** parameters in your Compute environment file. This procedure uses the default value, which is **/var/lib/nova/instances**. Only users with root access can set up shared storage. The Compute service user in the following procedure must be the same across Controller and Compute nodes.

Procedure

1. Perform the following steps on the Controller node:

- a. Ensure that the **/var/lib/nova/instances** directory has read-write access by the Compute service user, as shown in the following example:

```
drwxr-xr-x. 9 nova nova 4096 Nov  5 20:37 instances
```

- b. Add the following lines to the **/etc/exports** file:

```
/var/lib/nova/instances node1_IP(rw,sync,fsid=0,no_root_squash)
/var/lib/nova/instances node2_IP(rw,sync,fsid=0,no_root_squash)
```

Replace **node1_IP** and **node2_IP** for the IP addresses of the two Compute nodes, for example:

```
/var/lib/nova/instances 192.168.24.9(rw,sync,fsid=0,no_root_squash)
/var/lib/nova/instances 192.168.24.21(rw,sync,fsid=0,no_root_squash)
```

- c. Export the **/var/lib/nova/instances** directory to the Compute nodes:

```
# exportfs -avr
```

- d. Restart the NFS server:

```
# systemctl restart nfs-server
```

2. Perform the following steps on each Compute node:

- a. Ensure that the **/var/lib/nova/instances** directory exists locally.
- b. Add the following line to the **/etc/fstab** file:

```
NFS_SHARE_PATH:/var/lib/nova/instances /var/lib/nova/instances nfs4 defaults 0 0
```

- c. Mount the controller's instance directory to mount all the devices listed in **/etc/fstab**:

```
# mount -a -v
```

- d. Ensure that QEMU can access the directory's images:

```
# ls -ld /var/lib/nova/instances
drwxr-xr-x. 9 nova nova 4096 Nov  5 20:37 /var/lib/nova/instances
```

- e. Ensure that the node can see the instances directory with:

```
drwxr-xr-x. 9 nova nova 4096 Nov  5 20:37 /var/lib/nova/instances
```



NOTE

You can also run the following to view all mounted devices:

```
# df -k
```

10.8. TROUBLESHOOTING MIGRATION

The following issues can arise during instance migration:

- The migration process encounters errors.
- The migration process never ends.
- Performance of the instance degrades after migration.

10.8.1. Errors during migration

The following issues can send the migration operation into an **error** state:

- Running a cluster with different versions of Red Hat OpenStack Platform (RHOSP).
- Specifying an instance ID that cannot be found.
- The instance you are trying to migrate is in an **error** state.
- The Compute service is shutting down.
- A race condition occurs.
- Live migration enters a **failed** state.

When live migration enters a **failed** state, it is typically followed by an **error** state. The following common issues can cause a **failed** state:

- A destination Compute host is not available.
- A scheduler exception occurs.
- The rebuild process fails due to insufficient computing resources.
- A server group check fails.
- The instance on the source Compute node gets deleted before migration to the destination Compute node is complete.

10.8.2. Never-ending live migration

Live migration can fail to complete, which leaves migration in a perpetual **running** state. A common reason for a live migration that never completes is that client requests to the instance running on the source Compute node create changes that occur faster than the Compute service can replicate them to the destination Compute node.

Use one of the following methods to address this situation:

- Abort the live migration.
- Force the live migration to complete.

Aborting live migration

If the instance state changes faster than the migration procedure can copy it to the destination node, and you do not want to temporarily suspend the instance operations, you can abort the live migration.

Procedure

1. Retrieve the list of migrations for the instance:

```
$ nova server-migration-list <vm>
```

Replace **<vm>** with the name or ID of the instance.


2. Abort the live migration:

```
$ nova live-migration-abort <vm> <migration-id>
```

- Replace **<vm>** with the name or ID of the instance.
- Replace **<migration-id>** with the ID of the migration.

Forcing live migration to complete

If the instance state changes faster than the migration procedure can copy it to the destination node, and you want to temporarily suspend the instance operations to force migration to complete, you can force the live migration procedure to complete.



IMPORTANT

Forcing live migration to complete might lead to perceptible downtime.

Procedure

1. Retrieve the list of migrations for the instance:

```
$ nova server-migration-list <vm>
```

Replace **<vm>** with the name or ID of the instance.

2. Force the live migration to complete:

```
$ nova live-migration-force-complete <vm> <migration-id>
```

- Replace **<vm>** with the name or ID of the instance.
- Replace **<migration-id>** with the ID of the migration.

10.8.3. Instance performance degrades after migration

For instances that use a NUMA topology, the source and destination Compute nodes must have the same NUMA topology and configuration. The NUMA topology of the destination Compute node must have sufficient resources available. If the NUMA configuration between the source and destination Compute nodes is not the same, it is possible that live migration succeeds while the instance performance degrades. For example, if the source Compute node maps NIC 1 to NUMA node 0, but the destination Compute node maps NIC 1 to NUMA node 5, after migration the instance might route network traffic from a first CPU across the bus to a second CPU with NUMA node 5 to route traffic to NIC 1. This can result in expected behavior, but degraded performance. Similarly, if NUMA node 0 on the source Compute node has sufficient available CPU and RAM, but NUMA node 0 on the destination Compute node already has instances using some of the resources, the instance might run correctly but suffer performance degradation. For more information, see [Migration constraints](#).

CHAPTER 11. CONFIGURING VIRTUAL GPUS FOR INSTANCES

To support GPU-based rendering on your instances, you can define and manage virtual GPU (vGPU) resources according to your available physical GPU devices and your hypervisor type. You can use this configuration to divide the rendering workloads between all your physical GPU devices more effectively, and to have more control over scheduling your vGPU-enabled instances.

To enable vGPU in OpenStack Compute, create flavors that your cloud users can use to create Red Hat Enterprise Linux (RHEL) instances with vGPU devices. Each instance can then support GPU workloads with virtual GPU devices that correspond to the physical GPU devices.

The OpenStack Compute service tracks the number of vGPU devices that are available for each GPU profile you define on each host. The Compute service schedules instances to these hosts based on the flavor, attaches the devices, and monitors usage on an ongoing basis. When an instance is deleted, the Compute service adds the vGPU devices back to the available pool.

11.1. SUPPORTED CONFIGURATIONS AND LIMITATIONS

Supported GPU cards

For a list of supported NVIDIA GPU cards, see [Virtual GPU Software Supported Products](#) on the NVIDIA website.

Limitations when using vGPU devices

- You can enable only one vGPU type on each Compute node.
- Each instance can use only one vGPU resource.
- Live migration of vGPU between hosts is not supported.
- Suspend operations on a vGPU-enabled instance is not supported due to a libvirt limitation. Instead, you can snapshot or shelve the instance.
- Resize and cold migration operations on an instance with a vGPU flavor does not automatically re-allocate the vGPU resources to the instance. After you resize or migrate the instance, you must rebuild it manually to re-allocate the vGPU resources.
- By default, vGPU types on Compute hosts are not exposed to API users. To grant access, add the hosts to a host aggregate. For more information, see [Creating and managing host aggregates](#).
- If you use NVIDIA accelerator hardware, you must comply with the NVIDIA licensing requirements. For example, NVIDIA vGPU GRID requires a licensing server. For more information about the NVIDIA licensing requirements, see [NVIDIA License Server Release Notes](#) on the NVIDIA website.

11.2. CONFIGURING VGPU ON THE COMPUTE NODES

To enable your cloud users to create instances that use a virtual GPU (vGPU), you must configure the Compute nodes that have the physical GPUs:

1. Build a custom GPU-enabled overcloud image.
2. Prepare the GPU role, profile, and flavor for designating Compute nodes for vGPU.

3. Configure the Compute node for vGPU.
4. Deploy the overcloud.

**NOTE**

To use an NVIDIA GRID vGPU, you must comply with the NVIDIA GRID licensing requirements and you must have the URL of your self-hosted license server. For more information, see the [NVIDIA License Server Release Notes](#) web page.

11.2.1. Building a custom GPU overcloud image

Perform the following steps on the director node to install the NVIDIA GRID host driver on an overcloud Compute image and upload the image to the OpenStack Image Service (glance).

Procedure

1. Copy the overcloud image and add the **gpu** suffix to the copied image.

```
$ cp overcloud-full.qcow2 overcloud-full-gpu.qcow2
```

2. Install an ISO image generator tool from YUM.

```
$ sudo yum install genisoimage -y
```

3. Download the NVIDIA GRID host driver RPM package that corresponds to your GPU device from the NVIDIA website. To determine which driver you need, see the [NVIDIA Driver Downloads Portal](#).

**NOTE**

You must be a registered NVIDIA customer to download the drivers from the portal.

4. Create an ISO image from the driver RPM package and save the image in the **nvidia-host** directory.

```
$ genisoimage -o nvidia-host.iso -R -J -V NVIDIA nvidia-host/
l: -input-charset not specified, using utf-8 (detected in locale settings)
 9.06% done, estimate finish Wed Oct 31 11:24:46 2018
18.08% done, estimate finish Wed Oct 31 11:24:46 2018
27.14% done, estimate finish Wed Oct 31 11:24:46 2018
36.17% done, estimate finish Wed Oct 31 11:24:46 2018
45.22% done, estimate finish Wed Oct 31 11:24:46 2018
54.25% done, estimate finish Wed Oct 31 11:24:46 2018
63.31% done, estimate finish Wed Oct 31 11:24:46 2018
72.34% done, estimate finish Wed Oct 31 11:24:46 2018
81.39% done, estimate finish Wed Oct 31 11:24:46 2018
90.42% done, estimate finish Wed Oct 31 11:24:46 2018
99.48% done, estimate finish Wed Oct 31 11:24:46 2018
Total translation table size: 0
Total rockridge attributes bytes: 358
Total directory bytes: 0
```



```
Path table size(bytes): 10
Max brk space used 0
55297 extents written (108 MB)
```

5. Create a driver installation script for your Compute nodes. This script installs the NVIDIA GRID host driver on each Compute node that you run it on. The following example creates a script named **install_nvidia.sh**:

```
#!/bin/bash

# NVIDIA GRID package
mkdir /tmp/mount
mount LABEL=NVIDIA /tmp/mount
rpm -ivh /tmp/mount/NVIDIA-vGPU-rhel-8.1-430.27.x86_64.rpm
```

6. Customize the overcloud image by attaching the ISO image that you generated in Step 4, and running the driver installation script that you created in Step 5:

```
$ virt-customize --attach nvidia-packages.iso -a overcloud-full-gpu.qcow2 -v --run
install_nvidia.sh
[ 0.0] Examining the guest ...
libguestfs: launch: program=virt-customize
libguestfs: launch: version=1.36.10rhel=8,release=6.el8_5.2,libvirt
libguestfs: launch: backend registered: unix
libguestfs: launch: backend registered: uml
libguestfs: launch: backend registered: libvirt
```

7. Relabel the customized image with SELinux:

```
$ virt-customize -a overcloud-full-gpu.qcow2 --selinux-relabel
[ 0.0] Examining the guest ...
[ 2.2] Setting a random seed
[ 2.2] SELinux relabelling
[ 27.4] Finishing off
```

8. Prepare the custom image files for upload to the OpenStack Image Service:

```
$ mkdir /var/image/x86_64/image
$ guestmount -a overcloud-full-gpu.qcow2 -i --ro image
$ cp image/boot/vmlinuz-3.10.0-862.14.4.el8.x86_64 ./overcloud-full-gpu.vmlinuz
$ cp image/boot/initramfs-3.10.0-862.14.4.el8.x86_64.img ./overcloud-full-gpu.initrd
```

9. From the undercloud, upload the custom image to the OpenStack Image Service:

```
(undercloud) $ openstack overcloud image upload --update-existing --os-image-name
overcloud-full-gpu.qcow2
```

11.2.2. Designating Compute nodes for vGPU

To designate Compute nodes for vGPU workloads, you must create a new role file to configure the vGPU role, and configure a new flavor to use to tag the GPU-enabled Compute nodes.

Procedure

1. To create the new **ComputeGPU** role file, copy the file `/usr/share/openstack-tripleo-heat-templates/roles/Compute.yaml` to `/usr/share/openstack-tripleo-heat-templates/roles/ComputeGPU.yaml` and edit the following file sections:

Table 11.1. ComputeGPU role file edits

Section/Parameter	Current value	New value
Role comment	Role: Compute	Role: ComputeGpu
Role name	name: Compute	name: ComputeGpu
description	Basic Compute Node role	GPU Compute Node role
ImageDefault	overcloud-full	overcloud-full-gpu
HostnameFormatDefault	-compute-	-computegpu-
deprecated_nic_config_name	compute.yaml	compute-gpu.yaml

2. Generate a new roles data file named **gpu_roles_data.yaml** that includes the **Controller**, **Compute**, and **ComputeGpu** roles.

```
(undercloud) [stack@director templates]$ openstack overcloud roles generate -o
/home/stack/templates/gpu_roles_data.yaml Controller Compute ComputeGpu
```

The following example shows the **ComputeGpu** role details:

```
#####
# Role: ComputeGpu                                     #
#####
- name: ComputeGpu
  description: |
    GPU Compute Node role
  CountDefault: 1
  ImageDefault: overcloud-full-gpu
  networks:
    - InternalApi
    - Tenant
    - Storage
  HostnameFormatDefault: '%stackname%-computegpu-%index%'
  RoleParametersDefault:
    TunedProfileName: "virtual-host"
  # Deprecated & backward-compatible values (FIXME: Make parameters consistent)
  # Set uses_deprecated_params to True if any deprecated params are used.
  uses_deprecated_params: True
  deprecated_param_image: 'NovalImage'
  deprecated_param_extraconfig: 'NovaComputeExtraConfig'
  deprecated_param_metadata: 'NovaComputeServerMetadata'
  deprecated_param_scheduler_hints: 'NovaComputeSchedulerHints'
  deprecated_param_ips: 'NovaComputeIPs'
```

```

deprecated_server_resource_name: 'NovaCompute'
deprecated_nic_config_name: 'compute-gpu.yaml'
ServicesDefault:
- OS::TripleO::Services::Aide
- OS::TripleO::Services::AuditD
- OS::TripleO::Services::BootParams
- OS::TripleO::Services::CACerts
- OS::TripleO::Services::CephClient
- OS::TripleO::Services::CephExternal
- OS::TripleO::Services::CertmongerUser
- OS::TripleO::Services::Collectd
- OS::TripleO::Services::ComputeCeilometerAgent
- OS::TripleO::Services::ComputeNeutronCorePlugin
- OS::TripleO::Services::ComputeNeutronL3Agent
- OS::TripleO::Services::ComputeNeutronMetadataAgent
- OS::TripleO::Services::ComputeNeutronOvsAgent
- OS::TripleO::Services::Docker
- OS::TripleO::Services::Fluentd
- OS::TripleO::Services::IpaClient
- OS::TripleO::Services::Ipsec
- OS::TripleO::Services::Iscsid
- OS::TripleO::Services::Kernel
- OS::TripleO::Services::LoginDefs
- OS::TripleO::Services::MetricsQdr
- OS::TripleO::Services::MySQLClient
- OS::TripleO::Services::NeutronBgpVpnBagpipe
- OS::TripleO::Services::NeutronLinuxbridgeAgent
- OS::TripleO::Services::NeutronVppAgent
- OS::TripleO::Services::NovaCompute
- OS::TripleO::Services::NovaLibvirt
- OS::TripleO::Services::NovaLibvirtGuests
- OS::TripleO::Services::NovaMigrationTarget
- OS::TripleO::Services::ContainersLogrotateCron
- OS::TripleO::Services::OpenDaylightOvs
- OS::TripleO::Services::Podman
- OS::TripleO::Services::Rhsm
- OS::TripleO::Services::RsyslogSidecar
- OS::TripleO::Services::Securetty
- OS::TripleO::Services::SensuClient
- OS::TripleO::Services::Snmp
- OS::TripleO::Services::Sshd
- OS::TripleO::Services::Timesync
- OS::TripleO::Services::Timezone
- OS::TripleO::Services::TripleoFirewall
- OS::TripleO::Services::TripleoPackages
- OS::TripleO::Services::Tuned
- OS::TripleO::Services::Vpp
- OS::TripleO::Services::OVNController
- OS::TripleO::Services::OVNMetadataAgent

```

3. Register the node for the overcloud. For more information, see [Registering nodes for the overcloud](#) in the *Director Installation and Usage* guide.
4. Inspect the node hardware. For more information, see [Inspecting the hardware of nodes](#) in the *Director Installation and Usage* guide.

5. Create the **compute-vgpu-nvidia** flavor to use to tag nodes that you want to designate for vGPU workloads:

```
(undercloud) [stack@director templates]$ openstack flavor create --id auto --ram 6144 --disk 40 --vcpus 4 compute-vgpu-nvidia
```

Field	Value
OS-FLV-DISABLED:disabled	False
OS-FLV-EXT-DATA:ephemeral	0
disk	40
id	9cb47954-be00-47c6-a57f-44db35be3e69
name	compute-vgpu-nvidia
os-flavor-access:is_public	True
properties	
ram	6144
rxtx_factor	1.0
swap	
vcpus	4

6. Tag each node that you want to designate for GPU workloads with the **compute-vgpu-nvidia** profile.

```
(undercloud) [stack@director templates]$ openstack baremetal node set --property capabilities='profile:compute-vgpu-nvidia,boot_option:local' <node>
```

Replace **<node>** with the ID of the baremetal node.

11.2.3. Configuring the Compute node for vGPU and deploying the overcloud

You need to retrieve and assign the vGPU type that corresponds to the physical GPU device in your environment, and prepare the environment files to configure the Compute node for vGPU.

Procedure

1. Install Red Hat Enterprise Linux and the NVIDIA GRID driver on a temporary Compute node and launch the node. For more information about installing the NVIDIA GRID driver, see [Section 11.2.1, "Building a custom GPU overcloud image"](#).
2. On the Compute node, locate the vGPU type of the physical GPU device that you want to enable. For libvirt, virtual GPUs are mediated devices, or **mdev** type devices. To discover the supported **mdev** devices, enter the following command:

```
[root@overcloud-computegpu-0 ~]# ls
/sys/class/mdev_bus/0000\:06\:00.0/mdev_supported_types/
nvidia-11 nvidia-12 nvidia-13 nvidia-14 nvidia-15 nvidia-16 nvidia-17 nvidia-18 nvidia-19
nvidia-20 nvidia-21 nvidia-210 nvidia-22

[root@overcloud-computegpu-0 ~]# cat
/sys/class/mdev_bus/0000\:06\:00.0/mdev_supported_types/nvidia-18/description
num_heads=4, frl_config=60, framebuffer=2048M, max_resolution=4096x2160,
max_instance=4
```

3. Add the **compute-gpu.yaml** file to the **network-environment.yaml** file:

```
resource_registry:
  OS::TripleO::Compute::Net::SoftwareConfig: /home/stack/templates/nic-
  configs/compute.yaml
  OS::TripleO::ComputeGpu::Net::SoftwareConfig: /home/stack/templates/nic-
  configs/compute-gpu.yaml
  OS::TripleO::Controller::Net::SoftwareConfig: /home/stack/templates/nic-
  configs/controller.yaml
  #OS::TripleO::AllNodes::Validation: OS::Heat::None
```

4. Add the following parameters to the **node-info.yaml** file to specify the number of GPU Compute nodes, and the flavor to use for the GPU-designated Compute nodes:

```
parameter_defaults:
  OvercloudControllerFlavor: control
  OvercloudComputeFlavor: compute
  OvercloudComputeGpuFlavor: compute-vgpu-nvidia
  ControllerCount: 1
  ComputeCount: 0
  ComputeGpuCount: 1
```

5. Create a **gpu.yaml** file to specify the vGPU type of your GPU device:

```
parameter_defaults:
  ComputeGpuExtraConfig:
    nova::compute::vgpu::enabled_vgpu_types:
      - nvidia-18
```



NOTE

Each physical GPU supports only one virtual GPU type. If you specify multiple vGPU types in this property, only the first type is used.

6. Deploy the overcloud, adding your new role and environment files to the stack along with your other environment files:

```
(undercloud) $ openstack overcloud deploy --templates \
  -r /home/stack/templates/nvidia/gpu_roles_data.yaml
  -e /home/stack/templates/node-info.yaml
  -e /home/stack/templates/network-environment.yaml
  -e [your environment files]
  -e /home/stack/templates/gpu.yaml
```

11.3. CREATING THE VGPU IMAGE AND FLAVOR

To enable your cloud users to create instances that use a virtual GPU (vGPU), you can define a custom vGPU-enabled image, and you can create a vGPU flavor.

11.3.1. Creating a custom GPU instance image

After you deploy the overcloud with GPU-enabled Compute nodes, you can create a custom vGPU-enabled instance image with the NVIDIA GRID guest driver and license file.

Procedure

1. Create an instance with the hardware and software profile that your vGPU instances require:

```
(overcloud) [stack@director ~]$ openstack server create --flavor <flavor> --image <image>
temp_vgpu_instance
```

- Replace **<flavor>** with the name or ID of the flavor that has the hardware profile that your vGPU instances require. For information on default flavors, see [Manage flavors](#).
 - Replace **<image>** with the name or ID of the image that has the software profile that your vGPU instances require. For information on downloading RHEL cloud images, see [Image service](#).
2. Log in to the instance as a cloud-user. For more information, see [Log in to an Instance](#).
 3. Create the **gridd.conf** NVIDIA GRID license file on the instance, following the NVIDIA guidance: [Licensing an NVIDIA vGPU on Linux by Using a Configuration File](#).
 4. Install the GPU driver on the instance. For more information about installing an NVIDIA driver, see [Installing the NVIDIA vGPU Software Graphics Driver on Linux](#).



NOTE

Use the **hw_video_model** image property to define the GPU driver type. You can choose **none** if you want to disable the emulated GPUs for your vGPU instances. For more information about supported drivers, see [Appendix A, Image Configuration Parameters](#).

5. Create an image snapshot of the instance:

```
(overcloud) [stack@director ~]$ openstack server image create --name vgpu_image
temp_vgpu_instance
```

6. Optional: Delete the instance.

11.3.2. Creating a vGPU flavor for instances

After you deploy the overcloud with GPU-enabled Compute nodes, you can create a custom flavor that your cloud users can use to launch instances for GPU workloads.

Procedure

1. Create an NVIDIA GPU flavor. For example:

```
(overcloud) [stack@virtlab-director2 ~]$ openstack flavor create --vcpus 6 --ram 8192 --disk
100 m1.small-gpu
```

Field	Value
OS-FLV-DISABLED:disabled	False
OS-FLV-EXT-DATA:ephemeral	0
disk	100
id	a27b14dd-c42d-4084-9b6a-225555876f68

```

| name           | m1.small-gpu           |
| os-flavor-access:is_public | True                   |
| properties     |                         |
| ram            | 8192                   |
| rxtx_factor    | 1.0                    |
| swap           |                         |
| vcpus          | 6                      |
+-----+-----+

```

2. Assign a vGPU resource to the flavor that you created. You can assign only one vGPU for each instance.

```

(overcloud) [stack@virtlab-director2 ~]$ openstack flavor set m1.small-gpu --property
"resources:VGPU=1"

```

```

(overcloud) [stack@virtlab-director2 ~]$ openstack flavor show m1.small-gpu

```

```

+-----+-----+
| Field           | Value                   |
+-----+-----+
| OS-FLV-DISABLED:disabled | False                   |
| OS-FLV-EXT-DATA:ephemeral | 0                       |
| access_project_ids      | None                    |
| disk                   | 100                     |
| id                     | a27b14dd-c42d-4084-9b6a-225555876f68 |
| name                   | m1.small-gpu           |
| os-flavor-access:is_public | True                   |
| properties             | resources:VGPU='1'     |
| ram                     | 8192                   |
| rxtx_factor            | 1.0                    |
| swap                    |                         |
| vcpus                   | 6                      |
+-----+-----+

```

11.3.3. Launching a vGPU instance

You can create a GPU-enabled instance for GPU workloads.

Procedure

1. Create an instance using a GPU flavor and image. For example:

```

(overcloud) [stack@virtlab-director2 ~]$ openstack server create --flavor m1.small-gpu --
image vgpu_image --security-group web --nic net-id=internal0 --key-name lambda vgpu-
instance

```

2. Log in to the instance as a cloud-user. For more information, see [Log in to an Instance](#) .
3. To verify that the GPU is accessible from the instance, run the following command from the instance:

```

$ lspci -nn | grep <gpu_name>

```

11.4. ENABLING PCI PASSTHROUGH FOR A GPU DEVICE

You can use PCI passthrough to attach a physical PCI device, such as a graphics card, to an instance. If you use PCI passthrough for a device, the instance reserves exclusive access to the device for performing tasks, and the device is not available to the host.

Prerequisites

- The **pciutils** package is installed on the physical servers that have the PCI cards.
- The GPU driver is available to install on the GPU instances. For more information, see [Section 11.2.1, “Building a custom GPU overcloud image”](#).

Procedure

1. To determine the vendor ID and product ID for each passthrough device type, run the following command on the physical server that has the PCI cards:

```
# lspci -nn | grep -i <gpu_name>
```

For example, to determine the vendor and product ID for an NVIDIA GPU, run the following command:

```
# lspci -nn | grep -i nvidia
3b:00.0 3D controller [0302]: NVIDIA Corporation TU104GL [Tesla T4] [10de:1eb8] (rev a1)
d8:00.0 3D controller [0302]: NVIDIA Corporation TU104GL [Tesla T4] [10de:1db4] (rev a1)
```

2. To configure the Controller node on the overcloud for PCI passthrough, create an environment file, for example, **pci_passthru_controller.yaml**.
3. Add **PciPassthroughFilter** to the **NovaSchedulerDefaultFilters** parameter in **pci_passthru_controller.yaml**:

```
parameter_defaults:
  NovaSchedulerDefaultFilters:
    ['RetryFilter','AvailabilityZoneFilter','ComputeFilter','ComputeCapabilitiesFilter','ImageProperties
    Filter','ServerGroupAntiAffinityFilter','ServerGroupAffinityFilter','PciPassthroughFilter','NUMATo
    pologyFilter']
```

4. To specify the PCI alias for the devices on the Controller node, add the following to **pci_passthru_controller.yaml**:

```
ControllerExtraConfig:
  nova::pci::aliases:
    - name: "t4"
      product_id: "1eb8"
      vendor_id: "10de"
    - name: "v100"
      product_id: "1db4"
      vendor_id: "10de"
```



NOTE

If the **nova-api** service is running in a role other than the Controller, then replace **ControllerExtraConfig** with the user role, in the format **<Role>ExtraConfig**.

5. To configure the Compute node on the overcloud for PCI passthrough, create an environment file, for example, **pci_passthru_compute.yaml**.
6. To specify the available PCI IDs for the devices on the Compute node, add the following to **pci_passthru_compute.yaml**:

```
parameter_defaults:
  NovaPCIPassthrough:
    - vendor_id: "10de"
      product_id: "1eb8"
```

7. To enable IOMMU in the server BIOS of the Compute nodes to support PCI passthrough, add the **KernelArgs** parameter to **pci_passthru_compute.yaml**:

```
parameter_defaults:
  ...
  ComputeParameters:
    KernelArgs: "intel_iommu=on iommu=pt"
```

8. Deploy the overcloud, adding your custom environment files to the stack along with your other environment files:

```
(undercloud) $ openstack overcloud deploy --templates \
  -e [your environment files]
  -e /home/stack/templates/pci_passthru_controller.yaml
  -e /home/stack/templates/pci_passthru_compute.yaml
```

9. Configure a flavor to request the PCI devices. The following example requests two devices, each with a vendor ID of **10de** and a product ID of **13f2**:

```
# openstack flavor set m1.large --property "pci_passthrough:alias"="t4:2"
```

10. Create an instance with a PCI passthrough device:

```
# openstack server create --flavor m1.large --image rhelgpu --wait test-pci
```

11. Log in to the instance as a cloud-user. For more information, see [Log in to an Instance](#).
12. Install the GPU driver on the instance. For example, run the following script to install an NVIDIA driver:

```
$ sh NVIDIA-Linux-x86_64-430.24-grid.run
```

Verification

1. To verify that the GPU is accessible from the instance, run the following command from the instance:

```
$ lspci -nn | grep <gpu_name>
```

2. To check the NVIDIA System Management Interface status, run the following command from the instance:

```
$ nvidia-smi
```

Example output:

```
-----
| NVIDIA-SMI 440.33.01    Driver Version: 440.33.01    CUDA Version: 10.2    |
|-----+-----|
| GPU Name      Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|=====+=====|
| 0 Tesla T4          Off | 00000000:01:00.0 Off |             0 |
| N/A   43C   P0   20W / 70W |    0MiB / 15109MiB |      0%    Default |
|-----+-----|

Processes:                               GPU Memory |
| GPU    PID  Type  Process name                        Usage      |
|=====+=====|
| No running processes found                        |
|-----+-----|
```

CHAPTER 12. CONFIGURING REAL-TIME COMPUTE

In some use-cases, you might need instances on your Compute nodes to adhere to low-latency policies and perform real-time processing. Real-time Compute nodes include a real-time capable kernel, specific virtualization modules, and optimized deployment parameters, to facilitate real-time processing requirements and minimize latency.

The process to enable Real-time Compute includes:

- configuring the BIOS settings of the Compute nodes
- building a real-time image with real-time kernel and Real-Time KVM (RT-KVM) kernel module
- assigning the **ComputeRealTime** role to the Compute nodes

For a use-case example of Real-time Compute deployment for NFV workloads, see the [Example: Configuring OVS-DPDK with ODL and VXLAN tunnelling](#) section in the *Network Functions Virtualization Planning and Configuration Guide*.

12.1. PREPARING YOUR COMPUTE NODES FOR REAL-TIME



NOTE

Real-time Compute nodes are supported only with Red Hat Enterprise Linux version 7.5 or later.

Before you can deploy Real-time Compute in your overcloud, you must enable Red Hat Enterprise Linux Real-Time KVM (RT-KVM), configure your BIOS to support real-time, and build the real-time image.

Prerequisites

- You must use Red Hat certified servers for your RT-KVM Compute nodes. See [Red Hat Enterprise Linux for Real Time 7 certified servers](#) for details.
- You must enable the **rhel-8-for-x86_64-nfv-rpms** repository for RT-KVM to build the real-time image.



NOTE

You need a separate subscription to *Red Hat OpenStack Platform for Real Time* before you can access this repository. For details on managing repositories and subscriptions for your undercloud, see the [Registering and updating your undercloud](#) section in the *Director Installation and Usage* guide.

To check which packages will be installed from the repository, run the following command:

```
$ dnf repo-pkgs rhel-8-for-x86_64-nfv-rpms list
Loaded plugins: product-id, search-disabled-repos, subscription-manager
Available Packages
kernel-rt.x86_64          4.18.0-80.7.1.rt9.153.el8_0      rhel-8-for-x86_64-nfv-rpms
kernel-rt-debug.x86_64    4.18.0-80.7.1.rt9.153.el8_0      rhel-8-for-x86_64-nfv-
rpms
kernel-rt-debug-devel.x86_64 4.18.0-80.7.1.rt9.153.el8_0      rhel-8-for-x86_64-nfv-
```

```

rpms
kernel-rt-debug-kvm.x86_64      4.18.0-80.7.1.rt9.153.el8_0      rhel-8-for-x86_64-nfv-
rpms
kernel-rt-devel.x86_64         4.18.0-80.7.1.rt9.153.el8_0      rhel-8-for-x86_64-nfv-
rpms
kernel-rt-doc.noarch           4.18.0-80.7.1.rt9.153.el8_0      rhel-8-for-x86_64-nfv-rpms
kernel-rt-kvm.x86_64           4.18.0-80.7.1.rt9.153.el8_0      rhel-8-for-x86_64-nfv-
rpms
[ output omitted...]

```

Building the real-time image

To build the overcloud image for Real-time Compute nodes:

1. Install the **libguestfs-tools** package on the undercloud to get the **virt-customize** tool:

```
(undercloud) [stack@undercloud-0 ~]$ sudo dnf install libguestfs-tools
```

2. Extract the images:

```

(undercloud) [stack@undercloud-0 ~]$ tar -xf /usr/share/rhosp-director-images/overcloud-
full.tar
(undercloud) [stack@undercloud-0 ~]$ tar -xf /usr/share/rhosp-director-images/ironic-python-
agent.tar

```

3. Copy the default image:

```
(undercloud) [stack@undercloud-0 ~]$ cp overcloud-full.qcow2 overcloud-realtime-
compute.qcow2
```

4. Register the image and configure the required subscriptions:

```

(undercloud) [stack@undercloud-0 ~]$ virt-customize -a overcloud-realtime-compute.qcow2
--run-command 'subscription-manager register --username=[username] --password=
[password]'
[ 0.0] Examining the guest ...
[ 10.0] Setting a random seed
[ 10.0] Running: subscription-manager register --username=[username] --password=
[password]
[ 24.0] Finishing off

```

Replace the **username** and **password** values with your Red Hat customer account details. For general information about building a Real-time overcloud image, see the [Modifying the Red Hat Enterprise Linux OpenStack Platform Overcloud Image with virt-customize](#) knowledgebase article.

5. Find the SKU of the *Red Hat OpenStack Platform for Real Time* subscription. The SKU might be located on a system that is already registered to the Red Hat Subscription Manager with the same account and credentials. For example:

```
$ sudo subscription-manager list
```

6. Attach the *Red Hat OpenStack Platform for Real Time* subscription to the image:

```
(undercloud) [stack@undercloud-0 ~]$ virt-customize -a overcloud-realtime-compute.qcow2
--run-command 'subscription-manager attach --pool [subscription-pool]'
```

7. Create a script to configure **rt** on the image:

```
(undercloud) [stack@undercloud-0 ~]$ cat rt.sh
#!/bin/bash

set -eux

subscription-manager repos --enable=[REPO_ID]
dnf -v -y --setopt=protected_packages= erase kernel.$(uname -m)
dnf -v -y install kernel-rt kernel-rt-kvm tuned-profiles-nfv-host

# END OF SCRIPT
```

8. Run the script to configure the real-time image:

```
(undercloud) [stack@undercloud-0 ~]$ virt-customize -a overcloud-realtime-compute.qcow2 -
v --run rt.sh 2>&1 | tee virt-customize.log
```

9. Re-label SELinux:

```
(undercloud) [stack@undercloud-0 ~]$ virt-customize -a overcloud-realtime-compute.qcow2 -
-selinux-relabel
```

10. Extract **vmlinuz** and **initrd**. For example:

```
(undercloud) [stack@undercloud-0 ~]$ mkdir image
(undercloud) [stack@undercloud-0 ~]$ guestmount -a overcloud-realtime-compute.qcow2 -i -
-ro image
(undercloud) [stack@undercloud-0 ~]$ cp image/boot/vmlinuz-4.18.0-
80.7.1.rt9.153.el8_0.x86_64 ./overcloud-realtime-compute.vmlinuz
(undercloud) [stack@undercloud-0 ~]$ cp image/boot/initramfs-4.18.0-
80.7.1.rt9.153.el8_0.x86_64.img ./overcloud-realtime-compute.initrd
(undercloud) [stack@undercloud-0 ~]$ guestunmount image
```



NOTE

The software version in the **vmlinuz** and **initramfs** filenames vary with the kernel version.

11. Upload the image:

```
(undercloud) [stack@undercloud-0 ~]$ openstack overcloud image upload --update-existing -
-os-image-name overcloud-realtime-compute.qcow2
```

You now have a real-time image you can use with the **ComputeRealTime** composable role on select Compute nodes.

Modifying BIOS settings on Real-time Compute nodes

To reduce latency on your Real-time Compute nodes, you must modify the BIOS settings in the Compute nodes. You should disable all options for the following components in your Compute node BIOS settings:

- Power Management
- Hyper-Threading
- CPU sleep states
- Logical processors

See [Setting BIOS parameters](#) for descriptions of these settings and the impact of disabling them. See your hardware manufacturer documentation for complete details on how to change BIOS settings.

12.2. DEPLOYING THE REAL-TIME COMPUTE ROLE

Red Hat OpenStack Platform director provides the template for the **ComputeRealTime** role, which you can use to deploy real-time Compute nodes. You must perform additional steps to designate Compute nodes for real-time.

Procedure

1. Based on the `/usr/share/openstack-tripleo-heat-templates/environments/compute-real-time-example.yaml` file, create a `compute-real-time.yaml` environment file that sets the parameters for the **ComputeRealTime** role.

```
cp /usr/share/openstack-tripleo-heat-templates/environments/compute-real-time-example.yaml /home/stack/templates/compute-real-time.yaml
```

The file must include values for the following parameters:

- **IsolCpusList** and **NovaComputeCpuDedicatedSet**: List of isolated CPU cores and virtual CPU pins to reserve for real-time workloads. This value depends on the CPU hardware of your real-time Compute nodes.
 - **NovaComputeCpuSharedSet**: List of host CPUs to reserve for emulator threads.
 - **KernelArgs**: Arguments to pass to the kernel of the Real-time Compute nodes. For example, you can use `default_hugepagesz=1G hugepagesz=1G hugepages=<number_of_1G_pages_to_reserve> hugepagesz=2M hugepages=<number_of_2M_pages>` to define the memory requirements of guests that have huge pages with multiple sizes. In this example, the default size is 1GB but you can also reserve 2M huge pages.
2. Add the **ComputeRealTime** role to your roles data file and regenerate the file. For example:

```
$ openstack overcloud roles generate -o /home/stack/templates/rt_roles_data.yaml Controller Compute ComputeRealTime
```

This command generates a **ComputeRealTime** role with contents similar to the following example, and also sets the **ImageDefault** option to **overcloud-realtime-compute**.

```
- name: ComputeRealTime
  description: |
```

Compute role that is optimized for real-time behaviour. When using this role it is mandatory that an overcloud-realtime-compute image is available and the role specific parameters `IsolCpusList`, `NovaComputeCpuDedicatedSet` and `NovaComputeCpuSharedSet` are set accordingly to the hardware of the real-time compute nodes.

`CountDefault: 1`

`networks:`

`InternalApi:`

`subnet: internal_api_subnet`

`Tenant:`

`subnet: tenant_subnet`

`Storage:`

`subnet: storage_subnet`

`HostnameFormatDefault: '%stackname%-computerealttime-%index%'`

`ImageDefault: overcloud-realtime-compute`

`RoleParametersDefault:`

`TunedProfileName: "realtime-virtual-host"`

`KernelArgs: "" # these must be set in an environment file`

`IsolCpusList: "" # or similar according to the hardware`

`NovaComputeCpuDedicatedSet: "" # of real-time nodes`

`NovaComputeCpuSharedSet: "" #`

`NovaLibvirtMemStatsPeriodSeconds: 0`

`ServicesDefault:`

- `OS::TripleO::Services::Aide`
- `OS::TripleO::Services::AuditD`
- `OS::TripleO::Services::BootParams`
- `OS::TripleO::Services::CACerts`
- `OS::TripleO::Services::CephClient`
- `OS::TripleO::Services::CephExternal`
- `OS::TripleO::Services::CertmongerUser`
- `OS::TripleO::Services::Collectd`
- `OS::TripleO::Services::ComputeCeilometerAgent`
- `OS::TripleO::Services::ComputeNeutronCorePlugin`
- `OS::TripleO::Services::ComputeNeutronL3Agent`
- `OS::TripleO::Services::ComputeNeutronMetadataAgent`
- `OS::TripleO::Services::ComputeNeutronOvsAgent`
- `OS::TripleO::Services::Docker`
- `OS::TripleO::Services::Fluentd`
- `OS::TripleO::Services::IpaClient`
- `OS::TripleO::Services::Ipsec`
- `OS::TripleO::Services::Iscsid`
- `OS::TripleO::Services::Kernel`
- `OS::TripleO::Services::LoginDefs`
- `OS::TripleO::Services::MetricsQdr`
- `OS::TripleO::Services::MySQLClient`
- `OS::TripleO::Services::NeutronBgpVpnBagpipe`
- `OS::TripleO::Services::NeutronLinuxbridgeAgent`
- `OS::TripleO::Services::NeutronVppAgent`
- `OS::TripleO::Services::NovaCompute`
- `OS::TripleO::Services::NovaLibvirt`
- `OS::TripleO::Services::NovaLibvirtGuests`
- `OS::TripleO::Services::NovaMigrationTarget`
- `OS::TripleO::Services::ContainersLogrotateCron`
- `OS::TripleO::Services::OpenDaylightOvs`
- `OS::TripleO::Services::Podman`
- `OS::TripleO::Services::Rhsm`

```
- OS::TripleO::Services::RsyslogSidecar
- OS::TripleO::Services::Securetty
- OS::TripleO::Services::SensuClient
- OS::TripleO::Services::SkydiveAgent
- OS::TripleO::Services::Snmp
- OS::TripleO::Services::Sshd
- OS::TripleO::Services::Timesync
- OS::TripleO::Services::Timezone
- OS::TripleO::Services::TripleoFirewall
- OS::TripleO::Services::TripleoPackages
- OS::TripleO::Services::Vpp
- OS::TripleO::Services::OVNController
- OS::TripleO::Services::OVNMetadataAgent
```

For general information about custom roles and about the *roles-data.yaml*, see the [Roles](#) section.

3. Create the **compute-realtime** flavor to tag nodes that you want to designate for real-time workloads. For example:

```
$ source ~/stackrc
$ openstack flavor create --id auto --ram 6144 --disk 40 --vcpus 4 compute-realtime
$ openstack flavor set --property "cpu_arch"="x86_64" --property
"capabilities:boot_option"="local" --property "capabilities:profile"="compute-realtime"
compute-realtime
```

4. Tag each node that you want to designate for real-time workloads with the **compute-realtime** profile.

```
$ openstack baremetal node set --property capabilities='profile:compute-
realtime,boot_option:local' <NODE UUID>
```

5. Map the **ComputeRealTime** role to the **compute-realtime** flavor by creating an environment file with the following content:

```
parameter_defaults:
  OvercloudComputeRealTimeFlavor: compute-realtime
```

6. Run the **openstack overcloud deploy** command with the **-e** option and specify all the environment files that you created, as well as the new roles file. For example:

```
$ openstack overcloud deploy -r /home/stack/templates/rt~/my_roles_data.yaml -e
home/stack/templates/compute-real-time.yaml <FLAVOR_ENV_FILE>
```

12.3. SAMPLE DEPLOYMENT AND TESTING SCENARIO

The following example procedure uses a simple single-node deployment to test that the environment variables and other supporting configuration is set up correctly. Actual performance results might vary, depending on the number of nodes and guests that you deploy in your cloud.

1. Create the **compute-real-time.yaml** file with the following parameters:

```
parameter_defaults:
  ComputeRealTimeParameters:
```



```
IsolCpusList: "1"
NovaComputeCpuDedicatedSet: "1"
NovaComputeCpuSharedSet: "0"
KernelArgs: "default_hugepagesz=1G hugepagesz=1G hugepages=16"
```

2. Create a new **rt_roles_data.yaml** file with the **ComputeRealTime** role:

```
$ openstack overcloud roles generate -o ~/rt_roles_data.yaml Controller ComputeRealTime
```

3. Deploy the overcloud, adding both your new real-time roles data file and your real-time environment file to the stack along with your other environment files:

```
(undercloud) $ openstack overcloud deploy --templates \
-r /home/stack/rt_roles_data.yaml
-e [your environment files]
-e /home/stack/templates/compute-real-time.yaml
```

This command deploys one Controller node and one Real-time Compute node.

4. Log into the Real-time Compute node and check the following parameters. Replace <...> with the values of the relevant parameters from the **compute-real-time.yaml**.

```
[root@overcloud-computerealttime-0 ~]# uname -a
Linux overcloud-computerealttime-0 4.18.0-80.7.1.rt9.153.el8_0.x86_64 #1 SMP PREEMPT
RT Wed Dec 13 13:37:53 UTC 2017 x86_64 x86_64 x86_64 GNU/Linux
[root@overcloud-computerealttime-0 ~]# cat /proc/cmdline
BOOT_IMAGE=/boot/vmlinuz-4.18.0-80.7.1.rt9.153.el8_0.x86_64 root=UUID=45ae42d0-
58e7-44fe-b5b1-993fe97b760f ro console=tty0 crashkernel=auto console=ttyS0,115200
default_hugepagesz=1G hugepagesz=1G hugepages=16
[root@overcloud-computerealttime-0 ~]# tuned-adm active
Current active profile: realtime-virtual-host
[root@overcloud-computerealttime-0 ~]# grep ^isolated_cores /etc/tuned/realtime-virtual-host-
variables.conf
isolated_cores=<IsolCpusList>
[root@overcloud-computerealttime-0 ~]# cat /usr/lib/tuned/realtime-virtual-
host/lapic_timer_adv_ns
X (X != 0)
[root@overcloud-computerealttime-0 ~]# cat
/sys/module/kvm/parameters/lapic_timer_advance_ns
X (X != 0)
[root@overcloud-computerealttime-0 ~]# cat
/sys/devices/system/node/node0/hugepages/hugepages-1048576kB/nr_hugepages
X (X != 0)
[root@overcloud-computerealttime-0 ~]# crudini --get /var/lib/config-data/puppet-
generated/nova_libvirt/etc/nova/nova.conf compute cpu_dedicated_set
<NovaComputeCpuDedicatedSet>
[root@overcloud-computerealttime-0 ~]# crudini --get /var/lib/config-data/puppet-
generated/nova_libvirt/etc/nova/nova.conf compute cpu_shared_set
<NovaComputeCpuSharedSet>
```

12.4. LAUNCHING AND TUNING REAL-TIME INSTANCES

After you deploy and configure Real-time Compute nodes, you can launch real-time instances on those nodes. You can further configure these real-time instances with CPU pinning, NUMA topology filters, and huge pages.

Launching a real-time instance

1. Make sure that the **compute-realttime** flavor exists on the overcloud, as described in the *Deploying the Real-time Compute Role* section.
2. Launch the real-time instance.

```
# openstack server create --image <rhel> --flavor r1.small --nic net-id=<dpdk-net> test-rt
```

3. Optionally, verify that the instance uses the assigned emulator threads.

```
# virsh dumpxml <instance-id> | grep vcpu -A1
<vcpu placement='static'>4</vcpu>
<cputune>
  <vcpupin vcpu='0' cpuset='1'>
  <vcpupin vcpu='1' cpuset='3'>
  <vcpupin vcpu='2' cpuset='5'>
  <vcpupin vcpu='3' cpuset='7'>
  <emulatorpin cpuset='0-1'>
  <vcpusched vcpus='2-3' scheduler='fifo'
  priority='1'>
</cputune>
```

Pinning CPUs and setting emulator thread policy

To ensure that there are enough CPUs on each Real-time Compute node for real-time workloads, you need to pin at least one virtual CPU (vCPU) for an instance to a physical CPU (pCPUs) on the host. The emulator threads for that vCPU then remain dedicated to that pCPU.

Configure your flavor to use a dedicated CPU policy. To do so, set the **hw:cpu_policy** parameter to **dedicated** on the flavor. For example:

```
# openstack flavor set --property hw:cpu_policy=dedicated 99
```



NOTE

Make sure that your resources quota has enough pCPUs for the Real-time Compute nodes to consume.

Optimizing your network configuration

Depending on the needs of your deployment, you might need to set parameters in the **network-environment.yaml** file to tune your network for certain real-time workloads.

To review an example configuration optimized for OVS-DPDK, see the [Configuring the OVS-DPDK parameters](#) section of the *Network Functions Virtualization Planning and Configuration Guide*.

Configuring huge pages

It is recommended to set the default huge pages size to 1GB. Otherwise, TLB flushes might create jitter in the vCPU execution. For general information about using huge pages, see the [Running DPDK applications](#) web page.

Disabling Performance Monitoring Unit (PMU) emulation

Instances can provide PMU metrics by specifying an image or flavor with a vPMU. Providing PMU metrics introduces latency.



NOTE

The vPMU defaults to enabled when **cpu_mode=host-passthrough**.

If you do not need PMU metrics, then disable the vPMU to reduce latency by setting the PMU property to "False" in the image or flavor used to create the instance:

- Image: **hw_pmu=False**
- Flavor: **hw:pmu=False**

APPENDIX A. IMAGE CONFIGURATION PARAMETERS

The following keys can be used with the **property** option for both the **glance image-update** and **glance image-create** commands.

```
$ glance image-update IMG-UUID --property architecture=x86_64
```



NOTE

Behavior set using image properties overrides behavior set using flavors. For more information, see [Section 7.3, “Manage Flavors”](#).

Table A.1. Property Keys

Specific to	Key	Description	Supported values
All	architecture	The CPU architecture that must be supported by the hypervisor. For example, x86_64 , arm , or ppc64 . Run uname -m to get the architecture of a machine.	<ul style="list-style-type: none"> • alpha - DEC 64-bit RISC • armv7l - ARM Cortex-A7 MPCore • cris - Ethernet, Token Ring, AXis-Code Reduced Instruction Set • i686 - Intel sixth-generation x86 (P6 micro architecture) • ia64 - Itanium • lm32 - Lattice Micro32 • m68k - Motorola 68000 • microblaze - Xilinx 32-bit FPGA (Big Endian) • microblazeel - Xilinx 32-bit FPGA (Little Endian) • mips - MIPS 32-bit RISC (Big Endian) • mipsel - MIPS 32-bit RISC (Little Endian) • mips64 - MIPS 64-bit RISC (Big Endian) • mips64el - MIPS 64-bit RISC (Little Endian) • openrisc - OpenCores RISC • parisc - HP Precision Architecture RISC • parisc64 - HP Precision Architecture 64-bit RISC • ppc - PowerPC 32-bit

Specific to	Key	Description	Supported values
			<ul style="list-style-type: none"> ● ppc64 - PowerPC 64-bit ● ppcemb - PowerPC (Embedded 32-bit) ● s390 - IBM Enterprise Systems Architecture/390 ● s390x - S/390 64-bit ● sh4 - SuperH SH-4 (Little Endian) ● sh4eb - SuperH SH-4 (Big Endian) ● sparc - Scalable Processor Architecture, 32-bit ● sparc64 - Scalable Processor Architecture, 64-bit ● unicore32 - Microprocessor Research and Development Center RISC Unicores2 ● x86_64 - 64-bit extension of IA-32 ● xtensa - Tensilica Xtensa configurable microprocessor core ● xtensaeb - Tensilica Xtensa configurable microprocessor core (Big Endian)
All	hypervisor_type	The hypervisor type.	kvm, vmware
All	instance_uuid	For snapshot images, this is the UUID of the server used to create this image.	Valid server UUID
All	kernel_id	The ID of an image stored in the Image Service that should be used as the kernel when booting an AMI-style image.	Valid image ID

Specific to	Key	Description	Supported values
All	os_distro	The common name of the operating system distribution in lowercase.	<ul style="list-style-type: none"> ● arch - Arch Linux. Do not use archlinux or org.archlinux. ● centos - Community Enterprise Operating System. Do not use org.centos or CentOS. ● debian - Debian. Do not use Debian or org.debian. ● fedora - Fedora. Do not use Fedora, org.fedora, or org.fedoraproject. ● freebsd - FreeBSD. Do not use org.freebsd, freeBSD, or FreeBSD. ● gentoo - Gentoo Linux. Do not use Gentoo or org.gentoo. ● mandrake - Mandrakelinux (MandrakeSoft) distribution. Do not use mandrakelinux or MandrakeLinux. ● mandriva - Mandriva Linux. Do not use mandrivalinux. ● mes - Mandriva Enterprise Server. Do not use mandrivaent or mandrivaES. ● msdos - Microsoft Disc Operating System. Do not use ms-dos. ● netbsd - NetBSD. Do not use NetBSD or org.netbsd. ● netware - Novell NetWare. Do not use novell or NetWare. ● openbsd - OpenBSD. Do not use OpenBSD or org.openbsd. ● opensolaris - OpenSolaris. Do not use OpenSolaris or org.opensolaris. ● opensuse - openSUSE. Do not use suse, SuSE, or org.opensuse. ● rhel - Red Hat Enterprise Linux. Do not use redhat, RedHat, or com.redhat. ● sled - SUSE Linux Enterprise Desktop. Do not use com.suse. ● ubuntu - Ubuntu. Do not use Ubuntu, com.ubuntu, org.ubuntu, or canonical.

Specific to	Key	Description	<ul style="list-style-type: none"> windows - Microsoft Windows. Do not Supported values microsoft.server .
All	os_version	The operating system version as specified by the distributor.	Version number (for example, "11.10")
All	ramdisk_id	The ID of image stored in the Image Service that should be used as the ramdisk when booting an AMI-style image.	Valid image ID
All	vm_mode	The virtual machine mode. This represents the host/guest ABI (application binary interface) used for the virtual machine.	hvm -Fully virtualized. This is the mode used by QEMU and KVM.
libvirt API driver	hw_disk_bus	Specifies the type of disk controller to attach disk devices to.	scsi , virtio , ide , or usb . Note that if using iscsi , the hw_scsi_model needs to be set to virtio-scsi .
libvirt API driver	hw_cdrom_buses	Specifies the type of disk controller to attach CD-ROM devices to.	scsi , virtio , ide , or usb . If you specify iscsi , you must set the hw_scsi_model parameter to virtio-scsi .
libvirt API driver	hw_numa_nodes	Number of NUMA nodes to expose to the instance (does not override flavor definition).	Integer. For a detailed example of NUMA-topology definition, see the <code>hw:NUMA_def</code> key in Add Metadata .
libvirt API driver	hw_numa_cpus.0	Mapping of vCPUs N-M to NUMA node 0 (does not override flavor definition).	Comma-separated list of integers.

Specific to	Key	Description	Supported values
libvirt API driver	hw_numa_cpus.1	Mapping of vCPUs N-M to NUMA node 1 (does not override flavor definition).	Comma-separated list of integers.
libvirt API driver	hw_numa_mem.0	Mapping N MB of RAM to NUMA node 0 (does not override flavor definition).	Integer
libvirt API driver	hw_numa_mem.1	Mapping N MB of RAM to NUMA node 1 (does not override flavor definition).	Integer
libvirt API driver	hw_qemu_guest_agent	Guest agent support. If set to yes , and if qemu-ga is also installed, file systems can be quiesced (frozen) and snapshots created automatically.	yes / no

Specific to	Key	Description	Supported values
libvirt API driver	hw_rng_model	<p>Adds a random-number generator device to the image's instances. The cloud administrator can enable and control device behavior by configuring the instance's flavor. By default:</p> <ul style="list-style-type: none"> • The generator device is disabled. • <code>/dev/random</code> is used as the default entropy source. To specify a physical HW RNG device, set rng_dev_path to <code>"/dev/hwrng"</code> in your Compute environment file. 	virtio , or other supported device.

Specific to	Key	Description	Supported values
libvirt API driver	hw_scsi_model	Enables the use of VirtIO SCSI (virtio-scsi) to provide block device access for compute instances; by default, instances use VirtIO Block (virtio-blk). VirtIO SCSI is a para-virtualized SCSI controller device that provides improved scalability and performance, and supports advanced SCSI hardware.	virtio-scsi

Specific to	Key	Description	Supported values
libvirt API driver	hw_video_model	The video device driver to use in virtual machine instances.	<p>List of supported drivers, in order of precedence:</p> <ul style="list-style-type: none"> ● virtio. (Recommended) Virtual GPU with the Gallium GPU specification that uses the VIRGL renderer to render OpenGL. This GPU model is supported in all architectures, and can leverage hardware acceleration if the host has a dedicated GPU. For more information, see https://virgil3d.github.io/. ● qxl. High-performance driver for Spice or noVNC environments. ● cirrus. Legacy driver, use if the QXL driver is not available. ● vga. Use this driver for IBM Power environments. ● gop. Not supported for QEMU/KVM environments. ● xen. Not supported for KVM environments. ● vmvga. Legacy driver, do not use. ● none. Use this value to disable emulated graphics or video in virtual GPU (vGPU) instances where the driver is configured separately. For more information, see Chapter 11, Configuring virtual GPUs for instances.
libvirt API driver	hw_video_ram	Maximum RAM for the video image. Used only if a hw_video:ram_max_mb value has been set in the flavor's extra_specs and that value is higher than the value set in hw_video_ram .	Integer in MB (for example, 64)

Specific to	Key	Description	Supported values
libvirt API driver	hw_watchdog_action	Enables a virtual hardware watchdog device that carries out the specified action if the server hangs. The watchdog uses the i6300esb device (emulating a PCI Intel 6300ESB). If hw_watchdog_action is not specified, the watchdog is disabled.	<ul style="list-style-type: none"> ● disabled-The device is not attached. Allows the user to disable the watchdog for the image, even if it has been enabled using the image's flavor. The default value for this parameter is disabled. ● reset-Forcefully reset the guest. ● poweroff-Forcefully power off the guest. ● pause-Pause the guest. ● none-Only enable the watchdog; do nothing if the server hangs.
libvirt API driver	os_command_line	The kernel command line to be used by the libvirt driver, instead of the default. For Linux Containers (LXC), the value is used as arguments for initialization. This key is valid only for Amazon kernel, ramdisk, or machine images (aki, ari, or ami).	
libvirt API driver and VMware API driver	hw_vif_model	Specifies the model of virtual network interface device to use.	<p>The valid options depend on the configured hypervisor.</p> <ul style="list-style-type: none"> ● KVM and QEMU: e1000, ne2k_pci, pcnet, rtl8139, and virtio. ● VMware: e1000, e1000e, VirtualE1000, VirtualE1000e, VirtualPCNet32, VirtualSriovEthernetCard, and VirtualVmxnet. ● Xen: e1000, netfront, ne2k_pci, pcnet, and rtl8139.

Specific to	Key	Description	Supported values
VMware API driver	vmware_adaptertype	The virtual SCSI or IDE controller used by the hypervisor.	lsiLogic , busLogic , or ide
VMware API driver	vmware_ostype	A VMware GuestID which describes the operating system installed in the image. This value is passed to the hypervisor when creating a virtual machine. If not specified, the key defaults to otherGuest .	For more information, see Images with VMware vSphere .
VMware API driver	vmware_image_version	Currently unused.	1
XenAPI driver	auto_disk_config	If true, the root partition on the disk is automatically resized before the instance boots. This value is only taken into account by the Compute service when using a Xen-based hypervisor with the XenAPI driver. The Compute service will only attempt to resize if there is a single partition on the image, and only if the partition is in ext3 or ext4 format.	true / false

Specific to	Key	Description	Supported values
libvirt API driver and XenAPI driver	os_type	The operating system installed on the image. The XenAPI driver contains logic that takes different actions depending on the value of the os_type parameter of the image. For example, for os_type=windows images, it creates a FAT32-based swap partition instead of a Linux swap partition, and it limits the injected host name to less than 16 characters.	linux or windows

APPENDIX B. ENABLING THE LAUNCH INSTANCE WIZARD

There are two methods that you can use to launch instances from the dashboard:

- The Launch Instance form
- The Launch Instance wizard

The Launch Instance form is enabled by default, but you can enable the Launch Instance wizard at any time. You can also enable both the Launch Instance form and the Launch Instance wizard at the same time. The Launch Instance wizard simplifies the steps required to create instances.

1. Edit **/etc/openstack-dashboard/local_settings** file, and add the following values:

```
LAUNCH_INSTANCE_LEGACY_ENABLED = False
LAUNCH_INSTANCE_NG_ENABLED = True
```

2. Restart the *httpd* service:

```
# systemctl restart httpd
```

The preferences for the Launch Instance form and Launch Instance wizard are updated.

If you enabled only one of these options, the **Launch Instance** button in the dashboard opens that option by default. If you enabled both options, two **Launch Instance** buttons are displayed in the dashboard, with the button on the left opening the Launch Instance wizard and the button on the right opening the Launch Instance form.